AUTO-RUBRIC: LEARNING TO EXTRACT GENERALIZ-ABLE CRITERIA FOR REWARD MODELING

¹Alibaba Group, ²Ant Group

{xielipeng.xlp, huangsen.huang, zz297429, zouanni.zan,
zhaiyunpeng.zyp, liuboyin.lby, congling.chr,
jingmu.lzy, bolin.ding}@alibaba-inc.com
{rendingchao.rdc, kezun.zkz, haoyuan.huhy}@antgroup.com

• Code: https://github.com/modelscope/RM-Gallery

Dataset: https://huggingface.co/datasets/agentscope-ai/Auto-Rubric

ABSTRACT

Reward models are essential for aligning Large Language Models (LLMs) with human values, yet their development is hampered by costly preference datasets and poor interpretability. While recent rubric-based approaches offer transparency, they often lack systematic quality control and optimization, creating a trade-off between scalability and reliability. We address these limitations with a novel, training-free framework built on a key assumption: evaluation rubrics underlying human preferences exhibit significant generalization ability across diverse queries, a property that enables remarkable data efficiency. Our two-stage approach first infers high-quality, query-specific rubrics using a validation-guided Propose-Evaluate-Revise pipeline. Second, it generalizes these granular rubrics into a compact, non-redundant core set by maximizing an **information-theoretic coding rate**. The final output is an interpretable, hierarchical "Theme-Tips" rubric set. Extensive experiments demonstrate the framework's exceptional data efficiency and performance. Critically, using just 70 preference pairs (1.5% of the source data), our method also empowers smaller models like Owen3-8B to outperform specialized, fully-trained counterparts. This work pioneers a scalable, interpretable, and data-efficient path for reward modeling.

1 Introduction

Reinforcement Learning from Human Feedback (RLHF) is a powerful paradigm for aligning Large Language Models (LLMs) with human values (Ouyang et al., 2022). As shown in Figure 1, the core of RLHF is a reward model (RM) trained on vast datasets of human preferences to serve as a proxy for human judgment (Gao et al., 2023; Guo et al., 2025). However, this approach is fundamentally limited by the prohibitive cost of data acquisition and the "black-box" nature of the reward models (Liu et al., 2025a). This lack of interpretability not only hinders our ability to diagnose failures but also elevates the risk of "reward hacking" (DeepSeek-AI et al., 2025), where models exploit the proxy in unintended ways.

To address these shortcomings, rubric-based evaluation using explicit criteria has gained traction as a more transparent alternative. The rubric is a set of explicit human-readable criteria, such as factual accuracy and well-organized content, which can be effectively integrated as part of the prompt for the "LLM-as-a-Judge" paradigm. Early approaches relied on expert-defined rubrics (Hashemi et al., 2024) or large-scale crowd annotations(Bai et al., 2022), but their limited scalability prompted a shift

^{*}Corresponding author

towards automated rubric generation (Wang & Xiong, 2025; Gupta et al., 2025). These methods often produce rubrics that suffer from noise, redundancy, and misalignment with human preferences due to the lack of a verification mechanism. Consequently, a fundamental tension arises between scalability and fidelity, which poses the primary bottleneck for the broader adoption of rubric-based evaluation.

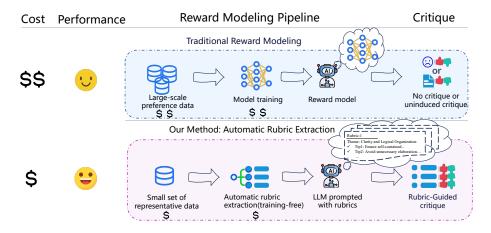


Figure 1: Comparison between traditional reward model training and our automatic rubric extraction method. Traditional methods require large-scale preference data and costly training to produce reward models. In contrast, our method uses a small set of representative data to automatically extract interpretable "Theme-Tips" rubrics for reward modeling, achieving low cost and high interpretability.

To resolve this tension, we propose a new framework for the automated generation and refinement of high-quality evaluation rubrics with a small corpus of preference data. Our work is built on a key assumption: evaluation rubrics underlying human preferences exhibit significant generalization ability across diverse queries. For example, humans generally prefer more logical, well-organized, and factual responses to different queries. Our goal is not to learn an opaque reward function but to explicitly infer the underlying principles—the rubrics—that govern human choices. This represents a fundamental shift from typical **reward model learning** to **rubric learning**, a contrast visually summarized in Figure 1.

To operationalize this new paradigm, our approach operates through two stages: **Query-Specific Rubric Generation** and **Query-Agnostic Rubric Aggregation**. First, Query-Specific Rubric Generation employs an iterative **Propose-Evaluate-Revise** loop that treats rubric generation as a constrained optimization problem, ensuring each rubric is validated for its discriminative power. Second, Query-Agnostic Rubric Aggregation uses an **information-theoretic selection** algorithm to distill the large pool of validated, granular rubrics into a compact, hierarchically structured rubric we term the "Theme-Tips" rubric. This rubric comprises high-level themes (e.g., Prioritize clarity) and corresponding actionable tips (e.g., Ensure narrative fidelity).

Our primary contributions are as follows:

- A data-efficient, training-free framework for automated rubric extraction. Our two-stage Propose-Evaluate-Revise and information-theoretic selection mechanism achieves state-of-the-art performance using only a fraction of typical preference data.
- Open-source rubric datasets. We release public datasets of query-agnostic rubrics inferred from preference data to facilitate research into interpretable alignment.
- A novel rubric analysis framework. We introduce a quantitative method to dissect rubric utility via Coverage, Precision, and Contribution metrics, offering deeper insight into the evaluation process.
- State-of-the-art performance on reward modeling benchmarks. Our method consistently improves base LLMs across four benchmarks. Notably, our performance on RewardBench2 sets a new state-of-the-art for training-free methods, with our rubric-enhanced Qwen3-235B (86.46%) and Qwen3-8B (80.91%) outperforming many specialized, fully-trained reward models.

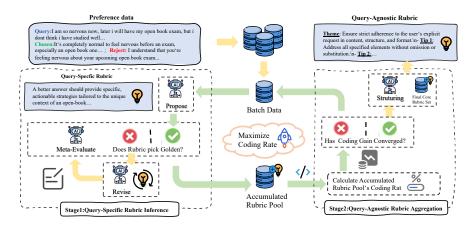


Figure 2: Overview of our two-stage rubric extraction framework. **Stage 1 (Query-Specific Rubric Inference):** An iterative Propose-Evaluate-Revise pipeline generates high-quality, query-specific rubrics from batches of preference data. **Stage 2 (Query-Agnostic Rubric Aggregation):** An information-theoretic selection algorithm distills the accumulated rubric pool into a compact, non-redundant core set by maximizing coding rate. The final output is structured into an interpretable "Theme-Tips" hierarchy for enhanced usability and generalization.

2 Related Work

LLM-as-a-Judge Evaluation. The paradigm of using LLMs as automatic evaluators is promising but undermined by significant reliability challenges. Early work identified surface-level biases such as positional and verbosity effects (Zheng et al., 2023), while recent studies have revealed deeper misalignments: LLM judges can systematically prioritize stylistic qualities over factual accuracy and safety (Feuer et al., 2025). Subsequent work has aimed to mitigate these issues through calibration techniques like CHARM (Zhu et al., 2025) or by developing specialized judge models such as Auto-J(Li et al., 2023) and JudgeLM(Zhu et al., 2023). However, these approaches often address the symptoms of bias rather than its root cause: an opaque and implicit judgment process. Our framework provides a more fundamental solution by replacing this implicit judgment with an explicit, verifiable rubric structure. This makes latent rubrics transparent, enabling the direct mitigation of such biases.

Rubric-Based Reward Modeling. The evolution of rubric-based methods reveals a persistent gap between rubric generation and effective rubric Optimization. Early approaches(Hashemi et al., 2024) relied on static, expert-authored rubrics that, while interpretable, were fundamentally unscalable. To overcome this limitation, recent work has automated rubric extraction using methods like chain-of-thought reasoning (Wang & Xiong, 2025) and templated prompts (Anugraha et al., 2025). However, these automatic methods typically produce a disorganized corpus of unrefined, often conflicting rules, and many approaches remain tethered to expensive parametric training (Viswanathan et al., 2025; Sun et al., 2024). Our work addresses this comprehensive lifecycle within a training-free paradigm, introducing a systematic framework to propose, refine, select, and structure rubrics into coherent, generalizable hierarchies from minimal data.

3 METHODOLOGY

Overview. Our framework systematically infers a general-purpose, interpretable set of evaluation rubrics from a small sample of human preferences. The methodology is structured into several stages, beginning at a granular level to maximize data efficiency. We first formulate rubric learning as an alternative to traditional reward modeling. Next, in the initial generation stage, a small seed batch is processed to infer high-fidelity, query-specific rubrics for each preference pair through a validation-centric loop, as illustrated in Figure 2. These granular rubrics are then aggregated into a compact, query-agnostic set using an information-theoretic approach. Finally, we introduce a quantitative framework to analyze the utility and contribution of each rubric in the final set.

3.1 FORMULATION

The conventional approach to learning from human preferences is to train a parametric reward model. Given a preference dataset $\mathcal{D} = \{(x_i, y_i^+, y_i^-)\}_{i=1}^N$, the goal is to learn a scalar reward function $r_{\theta}(x,y)$ that assigns a higher score to the preferred response. This is often optimized using a Bradley-Terry model(Bradley & Terry, 1952), where the probability of a preference is modeled as:

$$P(y_i^+ \succ y_i^- | x_i) = \sigma(r_\theta(x_i, y_i^+) - r_\theta(x_i, y_i^-)). \tag{1}$$

The objective is to find the optimal parameters θ by maximizing the log-likelihood over the dataset. While effective, this process yields an opaque reward function r_{θ} —a "black box" that offers limited insight into why one response is preferred over another. This lack of interpretability hinders failure diagnosis and trust.

To overcome these challenges, our work attempts a paradigm shift from **Reward Model Learning** to **Rubric Learning**. Instead of optimizing the parameters θ of an inscrutable function, our objective is to directly infer the explicit, human-readable rubric set R that best explains the preferences in \mathcal{D} . Our optimization problem remains:

$$R_{\text{task}}^* = \arg\max_{R} \sum_{i=1}^{N} \mathbb{I}[\text{eval}_{R}(x_i, y_i^+, y_i^-) = \text{correct}]. \tag{2}$$

However, the evaluation function, $\operatorname{eval}_R(\cdot)$, is no longer a parametric model but a transparent reasoning process guided by the natural language rubric in R. In practice, this evaluation function is implemented by prompting a large language model with the query, candidate responses, and the rubric set R, tasking it to make a preference judgment. Directly solving for R_{task} in Equation 2 is intractable, as it requires searching an extremely large and unstructured space of natural language rules. To make this problem tractable, we introduce a practical, two-stage framework that automates the generation and aggregation of rubrics from a small number of samples, as detailed below.

3.2 QUERY-SPECIFIC RUBRIC GENERATION

Instead of requiring a large-scale dataset, our framework begins at a granular level by processing a small seed batch to infer high-quality rubrics for each individual preference pair (x_i, y_i^+, y_i^-) . The core of this process is an iterative **Propose-Evaluate-Revise** loop, which emphasizes verification to ensure rubric quality.

Formally, the process for a single preference pair begins with a proposal model, $\mathcal{M}_{propose}$, proposing an initial rubric set:

$$R_i^{(0)} \leftarrow \mathcal{M}_{\text{propose}}(x_i, y_i^+, y_i^-). \tag{3}$$

At each iteration t, an evaluation model, $\mathcal{M}_{\text{evaluate}}$, verifies the current rubric set $R_i^{(t)}$ by making a judgment:

$$y_{\text{pred}}^{(t)} \leftarrow \mathcal{M}_{\text{evaluate}}(x_i, y_i^+, y_i^-, R_i^{(t)}). \tag{4}$$

This verification step is necessary, acting as a quality gate. If the prediction does not match the ground-truth preference $(y_{\text{pred}}^{(t)} \neq y_i^+)$, the failed rubric set $R_i^{(t)}$ is used as negative feedback. A revision model, $\mathcal{M}_{\text{revise}}$, then produces an improved set:

$$R_i^{(t+1)} \leftarrow \mathcal{M}_{\text{revise}}(x_i, y_i^+, y_i^-, R_i^{(t)}). \tag{5}$$

This iterative refinement continues until verification succeeds or a maximum number of iterations E_{max} is reached.

Finally, for each sample (x_i, y_i^+, y_i^-) , we generate a query-specific rubric set R_i^* that captures the most relevant evaluation criteria for that specific instance. This process populates a large pool of candidate rubrics, $\mathcal{R}_{\text{pool}} = \bigcup_{i=1}^N R_i^*$.

3.3 QUERY-AGNOSTIC RUBRIC AGGREGATION

While the initial generation stage produces a rich pool of high-quality, query-specific rubrics, \mathcal{R}_{pool} , this collection on its own is insufficient. It is inherently redundant, with the same underlying rubrics

Algorithm 1 Batch-Iterative Rubric Extraction

Input: Seed preference dataset \mathcal{D}_0 .

Output: Structured query-agnostic rubric set R_{task} .

- 1: Initialize an empty core rubric set $R_{\text{core}} \leftarrow \emptyset$.
- 2: **for** each batch-iteration t = 1, 2, ..., T **do**
- 3: Sample a mini-batch of preferences D_{batch} from \mathcal{D}_0 .
- 4: Generate query-specific rubrics \mathcal{R}_{new} using the **Propose-Evaluate-Revise** loop (Sec. 3.2).
- 5: Form a candidate pool: $\mathcal{R}_{pool} \leftarrow R_{core} \cup \mathcal{R}_{new}$.
- 6: Update the core set R_{core} by selecting rubrics from $\mathcal{R}_{\text{pool}}(\text{Eq. 7,8})$.
- 7: end for
- 8: Structure the final set R_{core} to yield R_{task} .
- 9: **return** R_{task} .

(e.g., clarity) expressed in many slightly different phrasings, and fragmented, with many rubrics being too specific to their source query to be broadly useful. Therefore, a query-agnostic aggregation stage is crucial. The primary objective is to distill a minimal yet comprehensive core set of rubrics that enhances generalization and transferability to unseen queries. This is achieved by identifying and consolidating the most essential and recurrent evaluation rubrics from the raw, query-specific pool.

To achieve this, we select a subset that maximizes information gain, ensuring high semantic coverage while minimizing redundancy. Geometrically, this is equivalent to selecting a set of embedding vectors that span the largest possible volume, a process that naturally penalizes redundant (i.e., near-collinear) vectors. Our selection criterion, the maximization of the coding rate (Yu et al., 2020), directly operationalizes this principle. It is an information-theoretic measure defined on the rubric embeddings $\mathbf{E}_R \in \mathbb{R}^{d \times |R|}$:

$$C(\mathbf{E}_R, \varepsilon) = \frac{1}{2} \log \det \left(\mathbf{I} + \frac{1}{\varepsilon^2 |R|} \mathbf{E}_R^{\top} \mathbf{E}_R \right), \tag{6}$$

where $C \in \mathbb{R}$ and $\varepsilon > 0$ controls the trade-off between compression and fidelity. Maximizing this function is equivalent to maximizing the volume spanned by the rubric embedding vectors, thus promoting diversity. The optimization problem is to find the core set R_{core} :

$$R_{\text{core}}^* = \arg \max_{R \subseteq \mathcal{R}_{\text{pool}}, |R| \le m} \mathcal{C}(\mathbf{E}_R, \varepsilon), \tag{7}$$

where m is the desired size of the rubric set. Since this problem is NP-hard, we employ a greedy algorithm that iteratively selects the rubric providing the highest marginal information gain. Starting with an empty set $R_0 = \emptyset$, at each step k, we add the rubric r_{k+1} such that:

$$r_{k+1} = \arg \max_{r \in \mathcal{R}_{pool} \setminus R_k} \left[\mathcal{C}(\mathbf{E}_{R_k \cup \{r\}}, \varepsilon) - \mathcal{C}(\mathbf{E}_{R_k}, \varepsilon) \right]. \tag{8}$$

This process continues until convergence, which is determined by an early-stopping criterion: the marginal gain in coding rate must fall below a minimum threshold (τ_{\min}) for a set number of consecutive iterations (p_{patience}) to ensure the information content of the core set has saturated. Finally, the selected core set is structured into our interpretable "Theme-Tips" hierarchy by a structuring LLM. This two-stage framework can be viewed as an **online learning process**, where new batches of preference data are used to generate more query-specific rubrics, which in turn iteratively refine and expand the query-agnostic core set, leading to high sample efficiency. The specific prompts used for each stage of our pipeline are detailed in Appendix H.

3.4 A Framework for Rubric Analysis

To ensure the final rubric set is not only performant but also robust and well-structured, we introduce a quantitative analysis framework. This framework, a core part of our methodology, allows us to dissect the utility of each individual rubric within the final set R_{task} . By evaluating each rubric along three key dimensions, as defined in (Eq. 9,10,11), we can validate the effectiveness of our aggregation process and gain deeper insights into the evaluation mechanism.

For each rubric $r_i \in R_{task}$, we define the following metrics:

Coverage: The proportion of test samples where the rubric provides a discriminative signal. This
metric measures the rubric's generality and applicability.

$$Coverage(r_j) = \frac{1}{|D_{test}|} \sum_{i \in D_{test}} \mathbb{I}[eval_{\{r_j\}}(x_i, y_i^+, y_i^-) \neq tie]. \tag{9}$$

• **Precision:** The conditional probability that the rubric's judgment aligns with the ground truth, given that it provides a discriminative signal. This measures the rubric's reliability.

$$Precision(r_j) = P(eval_{\{r_j\}} \text{ is correct}|eval_{\{r_j\}} \neq tie). \tag{10}$$

• **Contribution:** The marginal impact of a rubric on the full set's performance, measured by the drop in overall accuracy upon its removal. This quantifies the rubric's unique value and non-redundancy.

$$Contribution(r_j) = Acc(R_{task}) - Acc(R_{task} \setminus \{r_j\}).$$
(11)

This analytical framework is crucial for verifying that our method produces a complementary set of rubrics, balancing general, high-coverage rubrics with specialized, high-precision ones.

4 EXPERIMENT

In this section, we conduct a series of experiments to validate our framework's core contributions. We aim to demonstrate its: (1) state-of-the-art performance on standard reward modeling benchmarks; (2) high data efficiency, evidenced by rapid convergence; and (3) ability to produce high-value, interpretable rubrics, validated by our novel analysis method.

4.1 EXPERIMENTAL SETTING

Datasets. We extract rubrics on two preference datasets: (1) **HelpSteer3-Preference** (Wang et al., 2025) provides open, human-annotated preferences spanning four domains: General, STEM, Code, and Multilingual. We focus on the General domain for rubric extraction. (2) **UltraFeedback-Binarized**(Cui et al., 2024) contains prompts with model completions scored by GPT-4 on rubric such as helpfulness and honesty.

Baselines. We compare our method against three classes of baselines: (1) **Base Models**: Zero-shot evaluation using various LLMs without any rubrics. (2) **In-Context Learning (ICL)**(Dong et al., 2022): The same base models prompted with k=5 examples to perform preference evaluation. (3) **Training-based Reward Models**: A comprehensive suite of state-of-the-art models, including ArmoRM(Wang et al., 2024), J1(Whitehouse et al., 2025), R3(Anugraha et al., 2025), RM-R1(Chen et al., 2025), and Skywork-Reward-V2(Liu et al., 2025a).

Evaluation Benchmarks. We evaluate on four standard benchmarks covering diverse domains: RewardBench (Lambert et al., 2024), Rewardbench2 (Malik et al., 2025), RM-Bench (Liu et al., 2025b), JudgeBench (Tan et al., 2025).

Models. Our training-free framework employs **Qwen3-32B** (Yang et al., 2025) throughout all stages of rubric construction—including proposal, evaluation, revision, and structuring. We further analyze the generalizability of the resulting rubrics across a range of LLMs and find that those generated by Qwen3-32B demonstrate the strongest cross-model applicability (see Appendix C). Detailed experimental settings and implementation details are provided in Appendix B.

4.2 Main Results

State-of-the-Art Performance Across Benchmarks. Our framework demonstrates state-of-theart performance, securing the top score on all four evaluation benchmarks(details in Table 1). Specifically, our Qwen3-235B model achieves top scores of **94.87**% on RewardBench, **86.46**% on RewardBench2, **89.58**% on RM-Bench, and **86.29**% on JudgeBench. This broad success highlights the robustness and general applicability of the extracted rubrics.

Table 1: Performance of Models on Four Key Benchmarks and Average Score (in Percent)

Model	RewardBench	RewardBench2	RM-Bench	JudgeBench	Avg.b
Base Models					
Qwen3-8B	92.93 74.37		86.90	73.14	81.84
Qwen3-14B	92.66	76.30	87.70	75.14	82.95
Qwen3-32B	92.96	75.55	85.67	75.71	82.47
Qwen3-235B-A22B-Instruct-2507	93.70	83.78	87.55	83.14	87.04
ICL (k=5)					
Qwen3-8B	90.18	72.57	72.57 86.83		79.32
Qwen3-14B	89.58	74.89	87.29	70.86	80.66
Qwen3-32B	90.82	75.24	85.91	74.00	81.49
Qwen3-235B-A22B-Instruct-2507	90.42	81.38	86.91	82.86	85.39
Training-based Reward Models					
ArmoRM-Llama3-8B-v0.1	90.40	66.50	69.30	59.70	71.48
J1-Llama-8B	85.70	_c	73.40	42.00	67.0
J1-Llama-70B	93.30	-	82.70	60.00	78.6
R3-QWEN3-8B-14K	87.50	-	82.10	-	84.80
R3-QWEN3-14B-LORA-4K	89.30	-	84.90	-	87.1
RM-R1-Qwen-Instruct-32B	92.90	92.90 -		-	86.0
RM-R1-DeepSeek-Distill-Qwen-32B	90.90	-	83.90	-	87.4
Skywork-Reward-V2-Qwen3-8B	93.70	78.20	82.60	73.40	81.98
Our Method (Rubrics from HelpSteer3)					
Qwen3-8B	93.50	80.91	88.28	75.71	84.60
Qwen3-14B	93.74	93.74 81.66		79.71	84.5'
Qwen3-32B	93.80	82.27	88.11	80.86	86.20
Qwen3-235B-A22B-Instruct-2507	94.87 ^a	86.46 ^a	89.51	85.43	89.0
Our Method (Rubrics from UltraFeedback)					
Qwen3-8B	93.10	80.54	88.60	75.43	84.4
Qwen3-14B	93.67	80.91	88.72	78.86	85.54
Qwen3-32B	93.03	80.69	87.50	79.14	85.09
Qwen3-235B-A22B-Instruct-2507	94.54	85.97	89.58 ^a	86.29 ^a	89.10

^a **Bold** indicates the top-performing model in a column.

Consistent Improvement Across Model Scales. As shown in Table 1, the rubric-enhanced models consistently outperform their base versions, with substantial average accuracy gains observed for Qwen3-14B (+2.59%), Qwen3-32B (+3.79%). Notably, our approach empowers smaller models to achieve exceptional performance. For instance, our rubric-guided Qwen3-8B not only surpasses the specialized, fully-trained Skywork-Reward-V2-Qwen3-8B on RewardBench2 (80.91% vs. 78.20%) but also demonstrates clear superiority on RM-Bench (88.28% vs. 82.60%), proving its enhancement is not limited to a single benchmark.

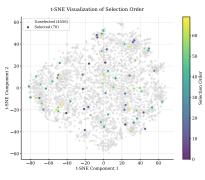
Robustness Across Rubric Source Datasets. The framework's generalization is strong, with rubrics derived from both human-annotated HelpSteer3 and AI-labeled UltraFeedback yielding competitive, state-of-the-art results. Despite nearly identical average scores on Qwen3-235B (**89.07**% vs. **89.10**%), each set excels on different benchmarks—HelpSteer3 on RewardBench/RewardBench2, and UltraFeedback on RM-Bench/JudgeBench. This proves the framework captures fundamental preference patterns from both human and AI annotations.

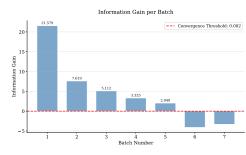
4.3 Data Efficiency and Convergence Analysis

A core claim of our work is achieving high performance with remarkable data efficiency. We demonstrate this by analyzing the convergence of our information-theoretic selection process, which iteratively draws batches of B=10 preference pairs from the 4,626-sample HelpSteer3 training dataset. Our framework employs an early-stopping mechanism where the information-theoretic selection terminates when the marginal gain in coding rate falls below $\tau_{\min}=0.002$ for $p_{\text{patience}}=2$ consections.

^b Average is calculated over benchmarks with available scores for each model.

^c Scores marked with '-' are unavailable as they were not reported in the original publications or on public leaderboards.





(a) t-SNE Visualization of Selection Order

(b) Information Gain per Batch

Figure 3: Analysis of the information-theoretic rubric selection process. (a) The t-SNE plot shows that early selections (darker points) are diverse and cover different semantic clusters. (b) Information gain (coding rate increment) saturates after a few batches, demonstrating rapid convergence and high sample efficiency.

utive iterations. Each preference pair undergoes a maximum of $E_{\rm max}=10$ epochs of the Propose-Evaluate-Revise cycle to ensure rubric quality. Figure 3 provides direct evidence for this efficiency, which we attribute to our selection process rapidly identifying a comprehensive and non-redundant rubric set from a small number of samples.

The t-SNE visualization in Figure 3a tracks the rubric selection order, showing that our algorithm actively promotes semantic diversity. Early selections (darker colors) are spread widely across different clusters, indicating that the framework prioritizes covering the entire semantic space over selecting similar, redundant rubrics. This ensures that each new rubric provides novel information, maximizing the value extracted from each sample.Additional analysis of the iterative refinement dynamics is provided in Appendix D, showing rapid convergence across different datasets.

This efficiency is quantified in the information gain plot (Figure 3b). The coding rate increment is highest in the first few batches and rapidly diminishes. Our early-stopping mechanism halted the process after 7 iterations, confirming that the core rubrics of preference could be captured from a very small fraction of the dataset. In total, only 70 samples were processed (1.5% of the source data) to distill the final, compact set of k=5 "Theme-Tips" rubrics.

4.4 ABLATION STUDIES

We conduct ablation studies to isolate the contribution of each core component of our framework, detailed in Table 2: (1) the iterative refinement of query-specific rubrics, (2) the information-theoretic selection of rubric subsets, (3) the final hierarchical structuring of the rubrics and (4) cross-model generalization capabilities.

Iterative Refinement. We test the necessity of our iterative refinement process by comparing our full feedback-driven approach against two baselines: a Single-pass Generation without refinement, and a Blind Revision that iterates without failed rubrics. The full iterative **Propose-Evaluate-Revise** process, with the help of the rubric evaluation and revision, it outperforms single-pass generation by +2.43% on RewardBench2 and +2.04% on RM-Bench. This confirms that a validation-driven feedback loop is essential for reliably improving rubric quality.

Rubric Selection Strategy. This ablation validates our information-theoretic selection strategy against a random selection baseline. The superiority of our approach is significant: our coding rate maximization strategy surpasses random selection by +3.16% on RewardBench2 and +1.31% on RM-Bench. This substantial performance gap confirms that efficiency and diversity-aware selection is essential for constructing a potent and non-redundant rubric set from a large pool of candidates.

Hierarchical Structure. We analyze the impact of rubric organization on evaluator performance by comparing our hierarchical Theme-Tips structure against flatter variants, including an unstruc-

tured list. The Theme-Tips format improves accuracy by +1.13% on RewardBench2 over a flat list, demonstrating that a balance of general rubrics (Themes) and specific guidance (Tips) is key to the effective application of the rubric.

Exceptional Cross-Model Generalization. To further validate the universality of our extracted rubrics, we conducted rigorous cross-model evaluations (see Appendix C, Figure 4 for full details). The results reveal that the rubrics generated by our framework are not only effective within their native model family but also demonstrate strong portability. Most notably, when applying the rubrics generated by Qwen3-32B to GPT-40, its performance on RewardBench2 surged from a baseline of 71.96% to 79.02%. This finding provides powerful evidence that our method captures fundamental and transferable rubrics of evaluation, rather than model-specific shortcuts or stylistic biases.

Component/Method	RewardBench2	RM-Bench	
Iterative Refinement Components			
Single-pass Generation	79.84	86.07	
Blind Revision (No Failed Rubric)	81.98 (+2.14)	85.79 (-0.28)	
Our Method (Full Iterative)	82.27 (+2.43)	88.11 (+2.04)	
Rubric Selection Methods			
Random Selection	79.11	86.80	
Our Method (Coding Rate)	82.27 (+3.16)	88.11 (+1.31)	
Rubric Structure Variants			
No Special Structure	81.14	87.41	
General (with optional rubrics)	80.01 (-1.13)	86.28 (-1.13)	
Theme (no tips)	80.77 (-0.37)	87.59 (+0.18)	
Theme-Tips	82.27 (+1.13)	88.11 (+0.70)	

Table 2: Comprehensive ablation studies across framework components.

4.5 Analysis of Core Rubrics

To validate that our framework produces high-value, interpretable data, we apply the analysis framework defined in our Methods 3.4 to the final extracted rubric set. This allows us to quantify the utility of each rubric and demonstrate that the final set is composed of complementary, non-redundant rubrics.

As shown in Table 3, foundational rubric like "Prioritize clarity" exhibit extremely high coverage (97.92%) and contribution (7.09% accuracy drop if removed), acting as the basis of the evaluation. In contrast, specialized rubric like "Ensure narrative fidelity" have lower coverage (71.91%) but the highest precision (68.24%), effectively handling niche scenarios that broader rubrics might miss. The significant contribution score of every rubric validates that our information-theoretic selection successfully produces a non-redundant set where each element plays a critical role. This analysis confirms that we are not just generating rubrics, but high-quality, structured evaluation knowledge. Complete rubric collections extracted from different datasets are presented in Appendix G.

Table 3: Analysis of the final rubrics, demonstrating their individual value.

Theme

Coverage (%) Precision (%) Contribution (A)

Rubric Theme	Coverage (%)	Precision (%)	Contribution (Δ Acc %)
Ensure factual accuracy.	91.91	62.78	4.42
Maintain strict adherence.	85.90	59.16	3.72
Prioritize clarity.	97.92	65.07	7.09
Deliver comprehensive.	97.16	65.92	4.78
Ensure narrative fidelity.	71.91	68.24	3.68

5 CONCLUSION

We introduced Auto-Rubric, a novel, training-free framework that successfully addresses the critical trade-off between performance, data efficiency, and interpretability in reward modeling. Our work demonstrates that the core criteria underlying human preferences can be automatically distilled into a compact, generalizable, and non-redundant set of "Theme-Tips" rubrics. The efficacy of this approach is notable: using just 70 preference pairs (1.5% of the source data), our extracted rubrics enabled a Qwen3-8B model to outperform specialized, fully-trained reward models, setting a new state-of-the-art for training-free methods on RewardBench2. These results provide a clear mandate: by shifting the focus from opaque **reward model learning** to transparent **rubric learning**, we can forge a more scalable, efficient, and trustworthy path for LLM alignment.

REFERENCES

- David Anugraha, Zilu Tang, Lester James V Miranda, Hanyang Zhao, Mohammad Rifqi Farhansyah, Garry Kuwanto, Derry Wijaya, and Genta Indra Winata. R3: Robust rubric-agnostic reward models. *arXiv preprint arXiv:2505.13388*, 2025.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, et al. Constitutional ai: Harmlessness from ai feedback, 2022.
- Ralph Allan Bradley and Milton E Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39:324–345, 1952.
- Xiusi Chen, Gaotang Li, Ziqi Wang, Bowen Jin, Cheng Qian, Yu Wang, Hongru Wang, Yu Zhang, Denghui Zhang, Tong Zhang, Hanghang Tong, and Heng Ji. Rm-r1: Reward modeling as reasoning, 2025. URL https://arxiv.org/abs/2505.02387.
- Ganqu Cui, Lifan Yuan, Ning Ding, Guanming Yao, Bingxiang He, Wei Zhu, Yuan Ni, Guotong Xie, Ruobing Xie, Yankai Lin, Zhiyuan Liu, and Maosong Sun. Ultrafeedback: Boosting language models with scaled ai feedback, 2024. URL https://arxiv.org/abs/2310.01377.
- DeepSeek-AI, Daya Guo, Dejian Yang, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning, 2025. URL https://arxiv.org/abs/2501.12948.
- Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Jingyuan Ma, Rui Li, Heming Xia, Jingjing Xu, Zhiyong Wu, Tianyu Liu, Baobao Chang, Xu Sun, Lei Li, and Zhifang Sui. A survey on incontext learning, 2022.
- Benjamin Feuer, Micah Goldblum, Teresa Datta, Sanjana Nambiar, Raz Besaleli, Samuel Dooley, Max Cembalest, and John P Dickerson. Style outweighs substance: Failure modes of LLM judges in alignment benchmarking. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=MzHNftnAM1.
- Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In *Proceedings of the 40th International Conference on Machine Learning*, ICML'23. JMLR.org, 2023.
- Jiaxin Guo, Zewen Chi, Li Dong, Qingxiu Dong, Xun Wu, Shaohan Huang, and Furu Wei. Reward reasoning model, 2025. URL https://arxiv.org/abs/2505.14674.
- Taneesh Gupta, Shivam Shandilya, Xuchao Zhang, Rahul Madhavan, Supriyo Ghosh, Chetan Bansal, Huaxiu Yao, and Saravan Rajmohan. Carmo: Dynamic criteria generation for context-aware reward modelling, 2025. URL https://arxiv.org/abs/2410.21545.
- Helia Hashemi, Jason Eisner, Corby Rosset, Benjamin Van Durme, and Chris Kedzie. LLM-rubric: A multidimensional, calibrated approach to automated evaluation of natural language texts. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar (eds.), *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 13806–13834, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.745. URL https://aclanthology.org/2024.acl-long.745/.
- Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, Noah A. Smith, and Hannaneh Hajishirzi. Rewardbench: Evaluating reward models for language modeling, 2024.
- Junlong Li, Shichao Sun, Weizhe Yuan, Run-Ze Fan, Hai Zhao, and Pengfei Liu. Generative judge for evaluating alignment, 2023.
- Chris Yuhao Liu, Liang Zeng, Yuzhen Xiao, Jujie He, Jiacai Liu, Chaojie Wang, Rui Yan, Wei Shen, Fuxiang Zhang, Jiacheng Xu, Yang Liu, and Yahui Zhou. Skywork-reward-v2: Scaling preference data curation via human-ai synergy, 2025a. URL https://arxiv.org/abs/2507.01352.

- Yantao Liu, Zijun Yao, Rui Min, Yixin Cao, Lei Hou, and Juanzi Li. RM-bench: Benchmarking reward models of language models with subtlety and style. In *The Thirteenth International Conference on Learning Representations*, 2025b. URL https://openreview.net/forum?id=QEHrmQPBdd.
- Saumya Malik, Valentina Pyatkin, Sander Land, Jacob Morrison, Noah A. Smith, Hannaneh Hajishirzi, and Nathan Lambert. Rewardbench 2: Advancing reward model evaluation, 2025. URL https://arxiv.org/abs/2506.01937.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35: 27730–27744, 2022. Also known as InstructGPT.
- Hanshi Sun, Momin Haider, Ruiqi Zhang, Huitao Yang, Jiahao Qiu, Ming Yin, Mengdi Wang, Peter Bartlett, and Andrea Zanette. Fast best-of-n decoding via speculative rejection, 2024.
- Sijun Tan, Siyuan Zhuang, Kyle Montgomery, William Yuan Tang, Alejandro Cuadron, Chenguang Wang, Raluca Popa, and Ion Stoica. Judgebench: A benchmark for evaluating LLM-based judges. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=G0dksFayVq.
- Vijay Viswanathan, Yanchao Sun, Shuang Ma, Xiang Kong, Meng Cao, Graham Neubig, and Tongshuang Wu. Checklists are better than reward models for aligning language models, 2025. URL https://arxiv.org/abs/2507.18624.
- Haoxiang Wang, Wei Xiong, Tengyang Xie, Han Zhao, and Tong Zhang. Interpretable preferences via multi-objective reward modeling and mixture-of-experts, 2024. URL https://arxiv.org/abs/2406.12845.
- Tevin Wang and Chenyan Xiong. Autorule: Reasoning chain-of-thought extracted rule-based rewards improve preference learning, 2025. URL https://arxiv.org/abs/2506.15651.
- Zhilin Wang, Jiaqi Zeng, Olivier Delalleau, Hoo-Chang Shin, Felipe Soares, Alexander Bukharin, Ellie Evans, Yi Dong, and Oleksii Kuchaiev. Helpsteer3-preference: Open human-annotated preference data across diverse tasks and languages, 2025. URL https://arxiv.org/abs/2505.11475.
- Chenxi Whitehouse, Tianlu Wang, Ping Yu, Xian Li, Jason Weston, Ilia Kulikov, and Swarnadeep Saha. J1: Incentivizing thinking in llm-as-a-judge via reinforcement learning, 2025. URL https://arxiv.org/abs/2505.10320.
- An Yang, Anfeng Li, Baosong Yang, et al. Qwen3 technical report, 2025.
- Yaodong Yu, Kwan Ho Ryan Chan, Chong You, Chaobing Song, and Yi Ma. Learning diverse and discriminative representations via the principle of maximal coding rate reduction. *Advances in neural information processing systems*, 33:9422–9434, 2020.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023. URL https://openreview.net/forum?id=uccHPGDlao.
- Lianghui Zhu, Xinggang Wang, and Xinlong Wang. Judgelm: Fine-tuned large language models are scalable judges, 2023.
- Xiao Zhu, Chenmien Tan, Pinzhen Chen, Rico Sennrich, Yanlin Zhang, and Hanxu Hu. Charm: Calibrating reward models with chatbot arena scores, 2025. URL https://arxiv.org/abs/2504.10045.

A THE USE OF LARGE LANGUAGE MODELS

In the preparation of this manuscript, we leveraged several large language models—including Google's Gemini, Alibaba's Qwen, and Anthropic's Claude—to assist with language editing and textual refinement. The role of these models was strictly limited to enhancing the clarity, grammatical correctness, fluency, and stylistic consistency of the manuscript. Specific tasks included refining sentence structure, suggesting alternative phrasings for improved readability, and harmonizing terminology and tone across sections. All outputs generated or suggested by these models were carefully evaluated, critically revised, and ultimately approved by the authors. The authors retain full responsibility for the scientific content, accuracy, and integrity of the final manuscript.

B EXPERIMENT SETTING DETAILS

Implementation Details. Our rubric extraction pipeline processes data in batches of B=10. The Propose-Evaluate-Revise cycle for each sample runs for a maximum of $E_{\rm max}=10$ epochs. The information-theoretic selection terminates when the marginal gain in coding rate falls below $\tau_{\rm min}=0.002$ for $p_{\rm patience}=2$ consecutive iterations. The final core set is structured into k=5 "Theme-Tips" rubrics. For evaluation, we use accuracy as the primary metric and employ voting strategies tailored to each benchmark's stability (e.g., voting@10 for RewardBench2, voting@5 for RewardBench and JudgeBench, and voting@1 for RM-Bench) to balance result reliability with computational efficiency. A comprehensive test-time scaling analysis examining the trade-offs between voting numbers and performance is presented in Appendix E.

C ANALYSIS ON THE GENERALIZABILITY OF MODEL-GENERATED RUBRICS

To select the optimal LLM for our framework, we analyzed the generalizability of evaluation rubrics generated by three leading models: **Qwen3-32B**, **GPT-40**, **and Claude-4-Sonnet**. We benchmarked each model's performance as an evaluator, both in a baseline condition (no rubric) and when guided by the rubrics from each of the three generators. The results in Figure 4 reveal clear patterns in rubric quality and cross-model utility.

The findings confirm two main points. First, in all scenarios, applying a model-generated rubric provides a significant performance uplift over the baseline. Second, and more critically, **Qwen3-32B generated rubrics exhibit the strongest generalizability**. This is most evident in the cross-model tests; for example, Qwen3-32B's rubric boosts GPT-4o's performance on RewardBench2 to **0.7902** and significantly higher than the score achieved with its own rubric (0.7453). While Claude-4-Sonnet consistently posts the highest absolute scores, proving it is a powerful standalone evaluator, the superior and consistent performance uplift that **Qwen3-32B's rubrics provide to** *other* **models** makes it the unambiguous choice for generating a robust, universally applicable set of rubric for our main experiments.

D QUERY-SPECIFIC ACCURACY IMPROVEMENT ANALYSIS

To further understand the learning dynamics of our rubric extraction framework, we analyze the query-specific accuracy improvement across training epochs for both datasets used in our experiments. Figure 5 illustrates the progressive enhancement in accuracy as our iterative refinement process generates and refines rubrics.

The results reveal several key insights about our framework's learning dynamics:

Rapid Initial Convergence. Both datasets exhibit steep accuracy improvements in the first 2-3 epochs, with HelpSteer3-Preference jumping from 86.1% to 92.7% (epoch 0 to 2) and UltraFeedback-Binarized improving from 93.9% to 97.4%. This rapid initial improvement demonstrates the effectiveness of our iterative refinement process in quickly identifying fundamental evaluation rubrics that govern human preferences.

Dataset-Specific Characteristics. UltraFeedback-Binarized consistently achieves higher accuracy levels and faster convergence, reaching 99.20% by epoch 9 compared to HelpSteer3-

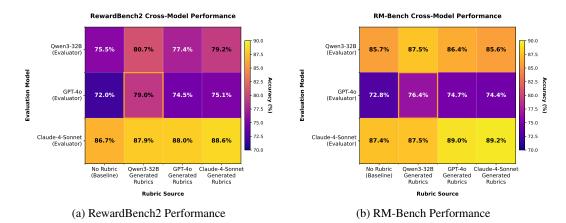


Figure 4: Cross-model rubric generalization analysis. The heatmaps show performance when evaluation models use rubrics generated by different LLMs. The orange borders highlight the best cross-model transferability: GPT-40 evaluator with Qwen3-32B rubrics achieves 79.02% on RewardBench2 and 76.37% on RM-Bench, demonstrating superior generalizability of Qwen3-32B generated rubrics.

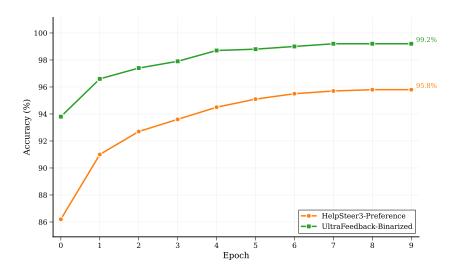


Figure 5: Query-specific generation accuracy improvement across epochs. The figure shows the progressive enhancement in accuracy for both HelpSteer3-Preference (orange line) and UltraFeedback-Binarized (green line) datasets. Both curves demonstrate rapid initial improvement followed by convergence to high accuracy levels, with UltraFeedback-Binarized achieving 99.20% and HelpSteer3-Preference reaching 95.80% by epoch 9. The steep initial gradient indicates the effectiveness of our iterative Propose-Evaluate-Revise mechanism in quickly identifying and refining core evaluation rubrics.

Preference's 95.80%. This difference likely reflects the distinct annotation methodologies: Ultra-Feedback's GPT-4-based scoring may exhibit more consistent patterns compared to HelpSteer3's human annotations, which naturally contain more subjective variance.

Convergence Stability. Both curves demonstrate saturation behavior after epoch 6, with minimal improvements in subsequent iterations. This validates our adaptive stopping mechanism and suggests that the core evaluation rubrics underlying human preferences can be effectively captured within a limited number of refinement cycles.

Cross-Dataset Validation. The consistent improvement patterns across both datasets support our core hypothesis about rubric convergence—despite different domains, annotation methods, and preference distributions, the underlying evaluation rubrics exhibit similar optimization dynamics, confirming the generalizability of our approach.

E TEST-TIME SCALING ANALYSIS

To evaluate the robustness and stability of our rubric-based evaluation framework, we investigate how performance scales with increasing voting numbers during test-time inference on Reward-Bench2. This analysis provides crucial insights into the trade-offs between computational cost and evaluation reliability.

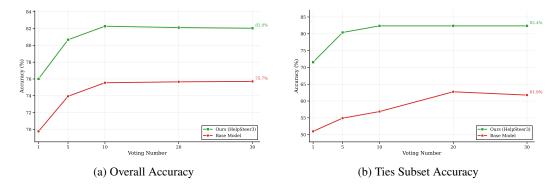


Figure 6: Test-time scaling analysis on RewardBench2 benchmark. (a) Overall accuracy improvement shows our rubric-enhanced HelpSteer3-Preference approach (green line) consistently outperforming the base model (red line) by 6-7 percentage points across all voting strategies (1, 5, 10, 20, 30). (b) Ties subset accuracy focuses on challenging cases where the base model produces tie decisions, demonstrating dramatic improvements of approximately 20 percentage points, highlighting the critical role of explicit rubrics in resolving ambiguous preference decisions.

The test-time scaling analysis reveals several important characteristics of our framework:

Consistent Performance Advantage. Figure 6a demonstrates that our rubric-enhanced approach maintains a consistent 6-7 percentage point advantage over the base model across all voting strategies. This systematic improvement suggests that our extracted rubrics provide fundamental evaluation capabilities that are orthogonal to the benefits of ensemble voting, creating additive performance gains.

Rapid Convergence with Low Voting Numbers. Both approaches show the most significant improvements when scaling from voting@1 to voting@5, with diminishing returns thereafter. This pattern indicates that the primary benefits of ensemble voting can be captured with relatively modest computational overhead. For practical deployment, voting@5 to voting@10 appears to offer the optimal balance between performance and efficiency.

Superior Performance on Challenging Cases. Figure 6b provides particularly compelling evidence for our framework's effectiveness. On the ties subset—representing the most challenging evaluation scenarios where base models struggle to make decisive judgments—our rubric-enhanced approach shows dramatic improvements of approximately 20 percentage points. This substantial gap highlights the critical role of explicit rubrics in providing discriminative power precisely where it is most needed.

Plateau Behavior and Computational Efficiency. Both figures demonstrate plateau behavior beyond voting@10, suggesting that additional computational investment yields marginal returns. This finding has important practical implications: our framework achieves near-optimal performance with moderate ensemble sizes, making it computationally efficient for real-world deployment while maintaining high evaluation quality.

Robustness Across Difficulty Levels. The consistent performance patterns across both overall accuracy and ties subset accuracy indicate that our rubrics provide robust evaluation capabilities that scale effectively across different difficulty levels. This robustness is crucial for practical applications where evaluation systems must handle diverse query types and ambiguous cases reliably.

F DETAILED EXPERIMENTAL ANALYSIS

To provide comprehensive insights into our framework's effectiveness, we conduct detailed analyses across multiple benchmarks and evaluation dimensions. This section examines where our rubric-guided approach provides the most significant value, focusing on challenging evaluation scenarios and domain-specific performance patterns.

F.1 CROSS-BENCHMARK PERFORMANCE ANALYSIS

Our detailed analysis encompasses two complementary benchmarks that together provide a comprehensive view of rubric effectiveness: RM-Bench, which allows us to examine performance on samples of varying difficulty levels, and RewardBench2, which offers diverse evaluation dimensions including challenging edge cases.

RM-Bench: Difficulty-Stratified Analysis. We conduct a stratified analysis on RM-Bench to understand how our rubrics perform across different difficulty levels (Table 4). The results reveal a clear and consistent pattern: our rubrics excel at resolving the most challenging cases, where base models struggle to make accurate preference judgments.

The difficulty-stratified analysis shows that hard samples benefit substantially more from rubric guidance (+4.68%) compared to the overall improvement (+2.45%). This 2x amplification effect on difficult cases demonstrates that our rubrics provide crucial discriminative power precisely where it is most needed—in scenarios where implicit evaluation rubrics are insufficient.

Domain-specific patterns further illuminate our framework's targeted strengths. The **Chat** domain exhibits the most dramatic improvement (+13.95% on hard samples), highlighting our rubrics' effectiveness in the notoriously subjective area of conversational evaluation where nuanced judgment rubric are critical. Substantial gains are also observed in **Math** (+4.54%) and **Safety-Refuse** (+3.64%), demonstrating broad applicability across diverse reasoning and safety scenarios.

Table 4: Performance analysis on RM-Bench using Qwen3-32B across domains and difficulty levels. All values are accuracies (%). Δ denotes the improvement.

	Overall (%)			Hard Samples (%)		
Domain	Base	Rubric	Δ	Base	Rubric	Δ
Overall	85.67	88.12	+2.45	77.74	82.42	+4.68
Chat	70.37	77.00	+6.63	36.95	50.90	+13.95
Math	89.29	91.89	+2.60	85.44	89.98	+4.54
Code	74.90	76.95	+2.05	72.08	73.54	+1.46
Safety-Refuse	94.68	96.64	+1.96	91.55	95.19	+3.64
Safety-Response	85.42	85.35	-0.07	68.58	72.61	+4.03

RewardBench2: Evaluation Dimension Analysis. To complement our difficulty-focused RM-Bench analysis, we examine performance across diverse evaluation dimensions on RewardBench2 (Table 5). RewardBench2 provides a more challenging and comprehensive evaluation setting, allowing us to understand where rubric-guided evaluation provides the most significant advantages across different types of evaluation rubrics.

The results reveal consistent and substantial improvements across all evaluation dimensions, with our rubrics achieving a remarkable overall improvement of +6.72% (from 75.55% to 82.27%). This substantial gain on a challenging benchmark demonstrates the robust effectiveness of our framework across diverse evaluation scenarios.

The most notable finding is the significant improvement on the **Ties** subset (+25.49%), jumping from 56.86% to 82.35%. This substantial gain represents the most challenging evaluation scenarios—where base models struggle to make decisive judgments—and highlights the critical discriminative power that explicit rubrics provide in ambiguous cases. The **Safety** domain also shows significant enhancement (+10.34%), demonstrating our rubrics' effectiveness in navigating nuanced safety considerations that require careful balance between multiple competing factors.

Importantly, even domains where base models already achieve strong performance show meaningful improvements: **Factuality** gains +8.84% and **Precise IF** improves by +5.62%. This pattern indicates that our rubrics provide value across the full spectrum of evaluation difficulty, from challenging edge cases to well-established domains, confirming the broad applicability and robustness of our approach.

Table 5: Performance analysis on RewardBench2 using Qwen3-32B across evaluation dimensions. All values are accuracies (%). Δ denotes the improvement.

Dimension	Base (%)	Rubric (%)	Δ (%)
Overall	75.55	82.27	+6.72
Precise IF	46.88	52.50	+5.62
Ties	56.86	82.35	+25.49
Factuality	65.68	74.53	+8.84
Focus	92.93	93.13	+0.20
Safety	77.13	87.47	+10.34
Math	85.79	86.34	+0.55

G EXTRACTED RUBRIC COLLECTIONS

This section presents the complete sets of query-agnostic rubrics extracted by our framework from different datasets and experimental configurations. These rubrics demonstrate the structured "Theme-Tips" hierarchy that emerges from our information-theoretic selection and thematic induction processes.

G.1 HelpSteer3-Preference Dataset Rubrics

The following rubrics were extracted from the HelpSteer3-Preference dataset:

Theme 1: Factual Accuracy and Canonical Consistency

Theme: Ensure factual accuracy, canonical consistency, and avoid fabrication or hallucination in responses.

- **Tip 1:** For queries about *Undertale*, ensure all character motivations and gameplay mechanics align with established lore, avoiding speculative or contradictory claims.
- **Tip 2:** When discussing historical milestones like early synchronized sound cartoons, correctly attribute "Steamboat Willie" instead of "My Old Kentucky Home" to maintain reliability.
- Tip 3: In responses involving Hogwarts students, include only canonically portrayed students with academically accurate achievements, excluding professors or non-student figures.
- **Tip 4:** Avoid inventing Sumerian texts or fabricated survey links; instead, acknowledge missing context and request clarification when necessary, especially for niche cultural references.

Theme 2: Strict Adherence to Prompt Requirements

Theme: Maintain strict adherence to prompt structure, formatting, and explicit user requirements.

- **Tip 1:** When asked for a single word, provide exactly one word without redundancy or additional suggestions, as in responses requiring minimal output.
- **Tip 2:** For prompts specifying 100 items, deliver a complete list even if the topic is broad, proactively selecting a relevant subject to fulfill the quantitative requirement.
- **Tip 3:** In tagline creation, directly incorporate core technology benefits like "distance at impact" and avoid vague or redundant phrasing that dilutes product relevance.
- **Tip 4:** When the prompt requires the word "scenery" followed by a colon and a one-word term, follow this exact syntactic structure without deviation.

Theme 3: Clarity and Structured Organization

Theme: Prioritize clarity, conciseness, and structured organization to enhance readability and directness.

- **Tip 1:** For a "Thank you" prompt, respond with a concise acknowledgment and an open invitation for further questions, avoiding assumptions about the user being a student or lawyer.
- **Tip 2:** When summarizing steps for building a dropshipping agent business, use bullet points or numbered lists to present key points logically and avoid hallucinated information.
- **Tip 3:** In audit findings related to deposit insurance boards, structure responses with precise, actionable items and conclude with a concise summary emphasizing implications.
- **Tip 4:** Avoid excessive formatting like bold text or unnecessary punctuation when explaining grammatical correctness, maintaining a straightforward and professional tone.

Theme 4: Comprehensive and Detailed Analysis

Theme: Deliver comprehensive, detailed, and thematically coherent narratives or analyses that fully address all prompt elements.

- **Tip 1:** For a CFA Institute Investment Foundations® Certificate explanation, include curriculum, eligibility, exam format, preparation resources, benefits, and continuing education with specific examples.
- **Tip 2:** In a fantasy story response, incorporate rich narrative detail, distinct character development, and immersive world-building such as vivid settings and dynamic interactions.
- **Tip 3:** When addressing a tax-proportional legislature, outline mechanics, implications, data collection, representation quotas, equity concerns, and constitutional considerations comprehensively.
- **Tip 4:** For a horror anime scene, use INT./EXT. designations, emphasize atmospheric tension, and describe creature details like a rhombus tail and chameleon-like head to align with anime style.

Theme 5: Narrative and Contextual Fidelity

Theme: Ensure narrative and contextual fidelity by preserving character dynamics, tone, and worldbuilding consistency.

- **Tip 1:** In responses involving Jade's character, maintain her authoritative yet professional tone, avoiding hostile shifts that contradict established behavior.
- **Tip 2:** For stories featuring Emily from KikoRiki, preserve her role as a mischievous prankster and integrate the whimsical tone when describing her failed morph into Rosa and the orange rear end mishap.

- **Tip 3:** When continuing a narrative about diaper use over potty training, maintain a playful, child-friendly tone and avoid contradictions with the original theme.
- **Tip 4:** In therapeutic role-play scenarios, prioritize immersive engagement with the patient's imaginative world through dialogue and validation, rather than clinical checklists.

G.2 ULTRAFEEDBACK-BINARIZED DATASET RUBRICS

The following rubrics were extracted from the UltraFeedback-Binarized dataset:

Theme 1: Factual Accuracy and Domain-Specific Knowledge

Theme: The answer must be factually accurate and grounded in correct domain-specific knowledge, avoiding misconceptions, logical errors, or speculative assumptions.

- **Tip 1:** Correctly apply scientific, technical, or mathematical principles (e.g., gravity, regex syntax, Pig Latin rules) with precision.
- **Tip 2:** Avoid perpetuating false premises (e.g., birds producing seeds) and instead clarify biological or conceptual inaccuracies.
- **Tip 3:** Use verified data, proper citations, and accurate terminology (e.g., Azure workflows, MLA formatting, product design details).
- **Tip 4:** When faced with ambiguity, seek clarification rather than making unfounded assumptions.
- **Tip 5:** Preserve original information in translations without adding, omitting, or distorting meaning.

Theme 2: Explicit Requirement Fulfillment

Theme: The answer must directly fulfill the user's explicit requirements in structure, content, and format, adhering strictly to all stated constraints.

- **Tip 1:** Follow prescribed structural elements (e.g., opening phrases, question framing, section order).
- **Tip 2:** Respect formatting rules (e.g., LaTeX, APA, SQL schema limits, phone number patterns).
- **Tip 3:** Address every component of multi-part queries (e.g., examples, explanations, code, citations).
- **Tip 4:** Use only valid functions, libraries, or commands within the correct technical context (e.g., Streamlit, PL/pgSQL).
- **Tip 5:** Extract or generate responses using only permitted sources (e.g., exact text spans, background passages).

Theme 3: Clarity and Logical Organization

Theme: The answer must provide clarity, coherence, and completeness through well-structured, concise, and logically organized reasoning.

- **Tip 1:** Offer step-by-step explanations that make reasoning transparent and verifiable.
- **Tip 2:** Maintain grammatical correctness and preserve original language or formatting conventions.
- **Tip 3:** Avoid unnecessary elaboration, redundancy, or irrelevant details that distract from the core task.
- **Tip 4:** Ensure responses are self-contained and understandable without external context.
- **Tip 5:** Use precise connectors and descriptive language to maintain fidelity in translation or interpretation.

Theme 4: Depth and Contextual Relevance

Theme: The answer must demonstrate depth and richness by integrating specific examples, actionable strategies, and contextual relevance.

- **Tip 1:** Include concrete, scenario-specific illustrations (e.g., AR gameplay mechanics, cultural program metrics).
- **Tip 2:** Provide practical implementation guidance with technical detail (e.g., iOS frameworks, OpenGL code).
- **Tip 3:** Link abstract concepts to real-world applications (e.g., symbolism in literature, ESG factors in market entry).
- **Tip 4:** Show progression or transformation (e.g., habit formation plans, historical scientific impact).
- **Tip 5:** Balance breadth and depth by covering multiple dimensions while offering nuanced analysis.

Theme 5: Ethical Responsibility and User Alignment

Theme: The answer must prioritize ethical responsibility, user alignment, and functional utility in its approach and tone.

- **Tip 1:** Reframe potentially offensive or harmful terms proactively to maintain respectful communication.
- **Tip 2:** Focus on actionable solutions rather than dismissive or overly theoretical responses.
- **Tip 3:** Tailor advice to the user's role, goals, or identity (e.g., UK lawyer, developer, educator).
- **Tip 4:** Encourage engagement through clear invitations or follow-up prompts when interaction is intended.
- **Tip 5:** Enhance transparency with confidence indicators or explicit justifications for conclusions.

H PROMPT TEMPLATES

```
## Overview
You are an expert rubric writer for open-ended question. Your job
generate a self-contained set of evaluation criteria ("rubrics")
   for choosing a better answer from candidate answers to a given
   query. Rubrics can cover aspects such as factual correctness,
   depth of reasoning, clarity, completeness, style, helpfulness,
   and common pitfalls. Each rubric item must be fully self-
   contained so that non-expert readers need not consult any
   external information.
I will give you:
1. the query (maybe contains history messages)
2. candidate answers
3. which answer is better than others
4. critics by the human experts, and you need to carefully read the
    critics provided by human experts and summarize the rubrics.
NOTE: The number of rubrics should be LESS THAN OR EQUAL TO {number
## Query
{query}
## Candidate Answers
<answer_1>{answer_1}</answer_1>
<answer_2>{answer_2}</answer_2>
## Better Answer
Answer {preference} is better than others.
## Critics
<critic>{critic}</critic>
## Output Format Requirements
<rubrics>your rubrics without index</rubrics>
```

Figure 7: Prompt for generating query-specific rubrics.

```
## Task Description
I will provide you with a set of rubrics, along with the current
   query and two responses. These rubrics are the primary basis for
    selecting the best answer. You must follow the steps specified
   in the Evaluation Process when conducting your evaluation
   process.
## Rubrics
{rubrics}
## Process
1. Confirm the task scenario of the current query and select the
   corresponding evaluation rubrics.
2. Identify the best response that meets the most selected rubrics.
## Query
{query}
## Response A
{response_a}
## Response B
{response_b}
## Output Requirement
Please choose the better response. Response "A", "B", or "tie"
   within the tags.
erence>A/B/tie</preference>
```

Figure 8: Prompt for rubric-based pairwise evaluation.

```
## Overview
You are an expert rubric writer for open-ended question. A self-
   contained set of evaluation criteria ("rubrics") is needed for
   choosing a better answer from candidate answers to a given query
   . Since the rubrics generated in the previous round failed to
   correctly select a better answer, you need to revise the rubrics
    . rubrics can cover aspects such as factual correctness, depth
   of reasoning, clarity, completeness, style, helpfulness, and
   common pitfalls. Each rubric item must be fully self-contained
   so that non-expert readers need not consult any external
   information.
I will give you:
1. the query (maybe contains history messages)
2. candidate answers
3. which answer is better than others
4. critics by the human experts, and you need to carefully read the
    critics provided by human experts and summarize the rubrics.
5. previous round rubrics that should to be improved
NOTE: The number of rubrics should be LESS THAN OR EQUAL TO {number
## Query
{query}
## Candidate Answers
<answer_1>
{answer_1}
</answer_1>
<answer_2>
{answer_2}
</answer_2>
## Better Answer
Answer {preference} is better than others.
## Previous Round rubrics
<rubric_1>
{previous_rubric_1}
</rubric_1>
## Output Format Requirements
Note: Ensure all outputs are placed within the tags like <tag>...</
   tag> as required!!!
<rubrics>
your improved rubrics without index
</rubrics>
```

Figure 9: Prompt for revising query-specific rubrics based on evaluation feedback.

```
##Task Description
Your task is to generate a set of evaluation rubrics to identify
   the best answer, based on the suggestions for determining from
   the examples. I will give you some examples, and every example
   contains the query and suggestion which has been verified to
   help select the best answer.
## Requirements
- Rubrics must be fully self-contained so that non-expert readers
   need not consult any external information.
- Each rubric should assess an independent dimension and be non-
   contradictory with others.
- Rubrics ensure that the overall judgment remains aligned and
   consistent for all examples.
- The number of rubrics should be LESS THAN OR EQUAL TO 5. The
   number of tips for each rubric should be LESS THAN OR EQUAL TO 5.
- Must strictly adhere to the Rubrics Format.
## Rubric Format
Each rubric consists of two parts:
- Theme: A concise and clear statement that captures the core focus
    of the rubric, and must be **necessary** for all queries with
   no assumption.
- Tips: Multiple bullet points that expand on or supplement the
   rubric and only focuses on some specific queries.
Here is an example of a rubric:
Theme: [Concise theme statement]
-Tip 1:
-Tip 2:
-Tip 3:
-(Optional: More tips as needed)
## Process
1. Based on the query and suggestions of each example, summarize
   the rubric of each example.
2. summarize the rubrics of each example, taking care to strictly
   adhere to the Requirements.
NOTE: The number of rubrics should be LESS THAN OR EQUAL TO 5. The
   number of tips for each rubric should be LESS THAN OR EQUAL TO
## Output Format Requirements
<rubrics>
Theme: [Concise theme statement]
-Tip 1: [Specific tip for certain queries]
-Tip 2: [Another specific tip]
-Tip 3: [Additional tip if needed]
Theme: [Another theme statement]
-Tip 1: [Related tip]
-Tip 2: [Another tip]
</rubrics>
```

Figure 10: Prompt for structuring the core rubric set into a "Theme-Tips" hierarchy.