Evaluating Medical LLMs by Levels of Autonomy: A Survey Moving from Benchmarks to Applications

Xiao Ye¹ Jacob Dineen¹ Zhaonan Li¹ Zhikun Xu¹ Weiyu Chen¹ Shijie Lu¹ Yuxi Huang¹ Ming Shen¹ Phu Tran² Ji–Eun Irene Yum² Muhammad Ali Khan² Muhammad Umar Afzal² Irbaz Bin Riaz² Ben Zhou¹

¹Arizona State University ²Mayo Clinic xiaoye2@asu.edu

Abstract

Medical Large language models achieve strong scores on standard benchmarks; however, the transfer of those results to safe and reliable performance in clinical workflows remains a challenge. This survey reframes evaluation through a levels-of-autonomy lens (L0-L3), spanning informational tools, information transformation and aggregation, decision support, and supervised agents. We align existing benchmarks and metrics with the actions permitted at each level and their associated risks, making the evaluation targets explicit. This motivates a levelconditioned blueprint for selecting metrics, assembling evidence, and reporting claims, alongside directions that link evaluation to oversight. By centering autonomy, the survey moves the field beyond score-based claims toward credible, risk-aware evidence for real clinical use.

1 Introduction

Large language models (LLMs) have advanced rapidly on medical benchmarks (Singhal et al., 2023; Hendrycks et al., 2020; Nazi and Peng, 2024). Both general-purpose (Qwen Team, 2025; DeepSeek-AI, 2025; OpenAI, 2025) and domain-specialized (Singhal et al., 2025; Bolton et al., 2024) LLMs now achieve high scores on licensing-style examinations and medical Q&A benchmarks, and they perform well on clinical summarization tasks (Oliveira et al., 2025; Tang et al., 2023). These headline results suggest that LLMs could meaningfully assist clinicians and patients across a range of information-centric workflows.

However, benchmark correctness alone is not sufficient for clinical use (Hager et al., 2024; Ma et al., 2025). Clinical deployment requires consistency, fairness, auditability, calibrated uncertainty, and demonstrably safe clinical reasoning (Omiye et al., 2023; Fehr et al., 2024). Most benchmarks are Q&A-centric, which rarely probe these aspects, allowing unsafe reasoning and missing context to

go undetected (Soroush et al., 2024; Asgari et al., 2025). A rigorous multidimensional evaluation that spans factual grounding, reasoning quality, uncertainty calibration, safety, and human preferences is required (Tam et al., 2024; Shool et al., 2025).

This survey reviews the state of LLM evaluation in the medical field and identifies its limitations. We first summarize how LLMs are being applied in the medical field and what current benchmarks measure. While these scores are informative, they only provide quick capability snapshots and overlook integration into real workflows, calibration, and traceable provenance. Therefore, we move from scores to applications, treating evaluation as a means of showing that a system is sufficient for a defined purpose and scope at a specific autonomy level. For each level, we define the scope, typical applications, evaluation focus, boundaries, and challenges: L0 Inform (no personalized advice); L1 Information Transformation & Aggregation (structure, summarize); L2 Decision Support (recommendations and personalized advice); and L3 Agents Under Human Supervision (plan + invoke tools/APIs to enqueue actions under explicit review). This organization makes evaluation targets explicit: factual grounding at L0/L1; calibrated reasoning and coverage at L2; tool-use safety and auditability at L3. Regarding the challenges in each level, they are cumulative, not isolated: each higher autonomy level inherits unresolved issues from lower levels. As autonomy and permitted actions expand, new risks emerge that are specific to that level's capabilities. Finally, we outline future work and recommendations for developing more robust evaluation frameworks to ensure that LLMbased systems can be trusted in clinical practice.

2 Related Work

Contemporary medical-LLM surveys mostly list datasets (e.g., USMLE/OKAP), task scenarios,

	L0-Inform	L1-Transformation & Aggregation	L2-Decision Support	L3-Agents under Supervision
Scope	Explain / Inform; NO advice	Structured extraction and aggregation	Recommendation; NO action	Plan, act under supervision
Application	Heath Search QA, Patient Summary	EHR Transformation, Retrieve & Aggregate	Diagnostic Reasoning, Patient Education	EHR co-pilot, Clinical Simulation
Example	What is MRI?	From this discharge note, extract the medication list.	Given the summary, is MRI appropriate?	Assemble prior results, draft an order, queue for sign-off.
Challenge	Hallucination, Bias	Attribution, Privacy	Reasoning Consistency, Fairness	Tool-use, Human-Al Interaction

Figure 1: Overview of autonomy levels (L0–L3) for medical LLMs, showing for each level the scope, typical applications, an example question/task, and key challenges. The rightward arrow indicates increasing autonomy and risk.

and evaluation modes, but they rarely organize evaluation targets by autonomy level or permitted clinical actions. Representative examples include a review contrasting closed and open-ended tasks and discussing agentic settings (Chen et al., 2025b), and a systematic review highlighting reliance on general-purpose models and accuracy metrics with limited calibration and safety assessment (Shool et al., 2025). Parallel strands propose conversation-quality and safety evaluations (Abbasian et al., 2024), human-rater rubrics such as QUEST (Tam et al., 2024), multi-dimensional criteria in SCORE (Tan et al., 2024), and reporting guidance in TRIPOD-LLM (Gallifant et al., 2025). We address this gap by mapping evaluation objectives and metrics to autonomy levels (L0-L3). This mapping clarifies what evidence is sufficient for a system's intended role, which risks must be tested and where human oversight is required. It also provides a practical blueprint for selecting datasets, metrics, and the identification of possible risks for different clinical actions.

Autonomy scales outside medical-LLM evaluation exist but target system design or human factors, not evaluation blueprints. In clinical decision research, levels of autonomy delineate who acts and who bears responsibility (Festor et al., 2021). In the broader agent literature, five-level design frameworks define autonomy via user roles (operator → observer) (Feng et al., 2025a), and industry taxonomies describe L0–L5 agentic behavior (Kirkovska et al., 2025). None of these prescribe autonomy-conditioned metrics or align evaluation with healthcare oversight and risk. Our survey fills that gap by integrating autonomy scales with con-

crete measurement choices for medical LLMs, so assessment can evolve with the agent's permitted actions and governance requirements.

3 Evaluation Methodologies

We start by listing current benchmarks and metrics for medical LLMs. This sets clear measurement boundaries before we move to the applications. For LLMs in the medical field, benchmarks provide low-cost, repeatable evidence about specific capabilities within a bounded scope; they surface failure modes early, track progress over time, and help align claims with a system's intended role (L0–L3). First, we group benchmarks by task and summarize the usefulness of each and the L0–L3 level(s) they inform (§3.1). We then summarize automated and human metrics and their limitations (§3.2). The list here is not exhaustive; a complete, expanded table appears in Appendix A.1.

3.1 Benchmarks

Exam Q&A. Task definition: answer multiple-choice or short-answer questions derived from professional medical examinations across specialties. Typical benchmarks: MedQA (USMLE) (Jin et al., 2021); MedMCQA (Indian medical entrance exams) (Pal et al., 2022). Usefulness: low-cost unit tests of factual breadth and specialty coverage that gate L0–L1 informational tools and surface coarse knowledge gaps.

Summarization. Task definition: generate faithful summaries, patient-friendly simplifications, or structural rewrites of clinical text. **Typical benchmarks:** MS² (multi-document evidence summaries) (Wang et al., 2022); PubMed long-

document summarization (Cohan et al., 2018). **Usefulness:** evaluates whether the output is faithful to the source (no omissions or fabrications) and structurally complete for explanation tasks (L0–L1).

Retrieval-augmented generation (RAG). Task definition: answer questions grounded in retrieved documents with explicit attribution to evidence passages. Typical benchmarks: BEIR (heterogeneous zero-shot retrieval) (Thakur et al., 2021); TREC-COVID (pandemic literature search) (Roberts et al., 2021). Usefulness: demonstrates that responses are supported by cited sources and do not contradict them, enabling freshness checks on evolving topics (L0–L1).

Information extraction. Task definition: extract and normalize clinical entities, relations, and codes from notes and reports. **Typical benchmarks:** CBLUE (NER and diagnosis normalization) (Zhang et al., 2021) **Usefulness:** establishes reliable structuring and normalization to support downstream aggregation in L1 systems.

Decision Support. Task definition: make thresholded recommendations or triage decisions in vignettes or multi-turn clinical scenarios scored by clinician rubrics Typical benchmarks: Health-Bench (multi-turn, physician-authored rubrics) (Arora et al., 2025); MedHELM (holistic medical evaluation suite) (Bedi et al., 2025). Usefulness: targets selective reliability for L2 systems by quantifying calibrated behavior at deployable thresholds and surfacing harm-proximal errors in simulation. Clinical Dialogue. Task definition: Conduct multi-turn conversations with patients or clinicians to achieve task goals while communicating safely and clearly. Typical benchmarks: MedDialog (EN/ZH) (Zeng et al., 2020); MedDG (ZH) (Liu et al., 2022). Usefulness: assesses communication effectiveness and safety across L0–L2.

3.2 Automated vs. Human Metrics

Automated Metrics. Automated metrics quantify answer correctness, calibration, faithfulness, and retrieval quality. For Q&A, Exact Match and F1 summarize accuracy, while ECE, Brier, and NLL capture probability calibration (ECE compares confidence to accuracy; Brier is the mean-squared error of predicted probabilities; NLL penalizes overconfident errors) (Guo et al., 2017; Brier, 1950; Manning et al., 2008). In selective prediction, risk-coverage curves relate error to the answered fraction and justify abstention on uncertain cases (Geifman and El-Yaniv, 2017; Geifman and El-Yaniv,

2019; Traub et al., 2024). For summarization, lexical overlap (ROUGE/chrF) and embedding-based similarity (BERTScore/BLEURT) are commonly complemented with source-grounded checks for omissions and contradictions (Lin, 2004; Popović, 2015; Zhang et al., 2020; Sellam et al., 2020). In RAG settings, retrieval ranking is assessed with Recall@k, MRR, and nDCG (Manning et al., 2008). In practice, care is needed: ECE can be sensitive to binning and mis-rank models unless robust estimators or verified calibration are used (Nixon et al., 2019; Kumar et al., 2019); risk-coverage summaries are hard to compare across tasks without principled area measures (Traub et al., 2024); overlap/embedding metrics correlate only weakly with factuality, motivating explicit factuality/entailment auditing (Maynez et al., 2020); and strong retrieval alone does not ensure correct answers, with automatic attribution evaluation remaining challenging (Li et al., 2024d; Joren et al., 2025).

Human Evaluation. Clinician raters typically apply behaviorally anchored rubrics (BARS-style) to score clinical correctness, coverage, contextualization, reasoning transparency, uncertainty handling, readability, actionability/safety, and empathy; inter-rater agreement is summarized with Cohen's κ or ICC (Holland et al., 2022; McHugh, 2012). For patient-facing text, readability targets are checked with Flesch/Flesch-Kincaid or SMOG indices (Singh et al., 2024; Badarudeen and Sabharwal, 2010). Reporting should also acknowledge known issues: κ can appear low despite high raw agreement when class prevalence or bias shifts (Feinstein and Cicchetti, 1990; Cicchetti and Feinstein, 1990); and readability formulas may disagree by several grade levels on the same passage and primarily capture surface difficulty rather than genuine comprehension, so multiple indices and taskspecific checks are advisable (Wang et al., 2013). LLM-as-Judge (LAJ). To scale open-ended assessments, many studies adopt LLM-as-judge: strong models render pairwise preferences or rubric-based scores, often with prompts that request quoted evidence. To control bias, studies randomize order or blind positions and, when feasible, ensemble multiple judges. (Zheng et al., 2023; Li et al., 2024c; Gu et al., 2024; Zhu et al., 2025; Tan et al., 2025). LAJ can track human preferences well in aggregate but is vulnerable to position and verbosity biases and self-preference effects; accordingly, order randomization, style/length normalization, and regular human spot-checks are recommended (Chen et al., 2024; Ye et al., 2025a; Wataoka et al., 2024).

4 Applications and Autonomy Levels

From Scores to Applications. Static benchmarks provide helpful, quick overviews of the model capabilities, but they do not capture real workflow context, calibration and abstention, provenance and audit trails, or role-specific needs. We therefore treat evaluation as showing that a system is **sufficient for a defined purpose and scope** at a given autonomy level. Concretely, each autonomy level is organized into five parts: (1) **Definition & scope**; (2) **Typical applications**-representative real tasks; (3) **Evaluation focus**-what to test and report; (4) **Scope boundaries**-what the level does not cover; and (5) **Challenges**. Benchmarks remain useful as ingredients in this assessment, not the destination.

4.1 L0 - Inform

Definition and Scope. At autonomy level L0, the system functions purely as an informational tool: it explains medical concepts and provides a general background in plain language. It neither tailors advice to an individual patient nor initiates clinical decisions. Outputs are educational in nature and include an explicit non-advice disclaimer.

Typical Applications. Representative L0 tasks include answering common health questions (e.g., "What is MRI?"), producing lay summaries of technical passages, and simplifying lab reports. Public datasets used as proxies include consumer Q&A corpora and patient-facing summarization sets: HealthSearchQA (commonly searched lay questions) within the MultiMedQA (Singhal et al., 2023) suite, the TREC 2017 LiveQA Medical task (Abacha et al., 2017), PubMedQA for abstract-level research comprehension, and consumer-question summarization datasets such as MeQSum and MEDIQA'21 (Ben Abacha and Demner-Fushman, 2019; Ben Abacha et al., 2021).

Evaluation Focus. L0 evaluation should primarily focus on accuracy, completeness (covering key points), and readability (appropriate and wellorganized) (Srinivasan et al., 2025). Some benchmarks also include structured rubrics or provide reference contexts that enable additional checks: HealthBench (Arora et al., 2025) grades freetext answers with physician-written rubrics, Multi-MedQA reports human ratings for long-form consumer answers, and PubMedQA uses the abstract

as the reference context. We reference these only to motivate the axes here, not to require source citation as a core L0 practice.

Scope Boundaries. L0 outputs are intentionally non-patient-specific and involve minimal reasoning: they recall and lightly synthesize facts but do not conduct case workups, triage, or make recommendations (those belong to L2 Decision Support). As a result, L0 content may omit personspecific contraindications or time-sensitive context and should avoid language that could be interpreted as advice. Readability should target plain-language norms (e.g., consumer-health instruments) to reduce misunderstanding, and common safety disclaimers may be shown without tailoring. Some L0 benchmarks include rubrics or reference contexts for scoring (e.g., physician-written criteria or abstract-based contexts), but we do not treat source citation as a core requirement for L0.

Challenges. Hallucination: Even at L0, models often produce fluent but unfaithful text: Broad surveys separate hallucinations into intrinsic and extrinsic, link them to training/decoding choices and weak grounding, and recommend source- and taskaware evaluations instead of generic overlap scores (Ji et al., 2023). In summarization specifically, human studies show that systems invent unsupported details and that standard n-gram metrics miss these errors, motivating faithfulness-oriented checks instead (Maynez et al., 2020). For patient-facing summaries, clinical audits advise using medical rubrics that check whether each claim is supported by the underlying notes and whether important facts are missing-rather than relying on readability alone. (Asgari et al., 2025). However, because these tools audit rather than eliminate hallucinations, the problem persists. **Bias:** We use bias to mean systematic tendencies that push outputs away from truth or intended scope (e.g., agreement pressure, overconfidence, or sensitivity to superficial cues), such that the same input intent can yield subtly distorted explanations. These tendencies plausibly arise from the pretraining objective (next-token prediction on skewed web corpora) and are further shaped by alignment procedures that optimize against human preferences (Ouyang et al., 2022; Sharma et al., 2023; Xu et al., 2025; Shen et al., 2025). In L0 applications, such biases surface as assistants mirroring user beliefs rather than correcting them (Sharma et al., 2023) and shifting outputs under small, meaning-preserving prompt changes such as formatting tweaks, option order in multiple-choice

settings, or early anchoring hints (Sclar et al., 2023; Pezeshkpour and Hruschka, 2023; Lou et al., 2024; Zhou et al., 2024; Li et al., 2024b; Ye et al., 2025b; Li et al., 2024a; RRV et al., 2025). They persist in practice because preference-optimized objectives and prompt conventions are integral to how LLMs are used (Kadavath et al., 2022; Steyvers et al., 2025). Consequently, mitigations such as reporting performance ranges across prompt formats, simple option-order calibrations, anchor-aware templates, and a conservative tone can reduce bias at L0, but they do not reliably remove it. (Sclar et al., 2023; Pezeshkpour and Hruschka, 2023; Lou et al., 2024; Steyvers et al., 2025).

4.2 L1 - Information Transformation & Aggregation

Definition and Scope This stage turns raw, heterogeneous clinical data into standardized, computable representations and then combines them with external evidence to produce grounded outputs. In practice, health systems map local EHR fields to an interoperability standard (e.g., HL7 FHIR) or a research schema (e.g., OMOP CDM), attach machine-readable provenance for auditability, and operate over de-identified corpora such as MIMIC-IV that illustrate the target tables (encounters, labs, medications) and privacy constraints (Abacha et al., 2021; Alsentzer et al., 2023; Zhang et al., 2024a; Liu et al., 2021).

Typical Applications. EHR data transformation: Typical pipelines mix schema harmonization (FHIR/OMOP ETL), clinical NLP to extract entities/attributes from notes, concept normalization to standard vocabularies, and de-identification. i2b2/n2c2 shared tasks supply widely used deidentified note sets for de-identification, concept extraction, relation labeling, and medication-change context, enabling objective measurement of spanlevel and mapping accuracy (Li et al., 2016; Abacha et al., 2021; Nowak et al., 2023; Mahajan et al., 2023; Henry et al., 2020). Retrieve & **Aggregate:** On top of the transformed corpus, systems index structured facts (problem lists, meds, labs) and unstructured notes, then pair them with external sources (guidelines, reviews) via retrievalaugmented generation (RAG). Retrieval metrics (Recall@k, nDCG, MRR) assess whether the right evidence is fetched; generation metrics (faithfulness/attribution, grounded-answer rate) check that outputs rely on retrieved passages rather than model priors (Cohan et al., 2020; Zhang et al., 2024a; Tang et al., 2023). Representative resources include BEIR for generalizable retrieval evaluation, TREC-COVID for high-stakes, rapidly evolving topics, and domain-specific retrievers such as Med-CPT for biomedical search (Thakur et al., 2021; Roberts et al., 2021; Jin et al., 2023).

Evaluation Focus. Transformation: report extraction/normalization scores (e.g., span-level F1, concept mapping accuracy), coverage/completeness of key fields, and lineage completeness. Retrieval & Aggregation: report Recall@k/nDCG/MRR; grounded-answer and attribution rates to retrieved passages; contradiction-to-source; and selective prediction/abstention rates under uncertainty (Alsentzer et al., 2023; Abacha et al., 2021; Zhang et al., 2024a).

Scope Boundaries. L1 improves structure, traceability, and access to evidence, but it does not perform the patient-specific reasoning required for diagnosis, test selection, or treatment trade-offs. Real clinical workups must integrate temporality, comorbidities, contraindications, and uncertaintycapabilities not captured by schema or Recall@k alone. Studies also show that (i) retrieval can surface conflicting or outdated sources (e.g., during pandemics) and (ii) LLMs may misattribute or overtrust citations; recent medical RAG evaluations highlight these gaps even when retrieval quality is strong (Roberts et al., 2021; Peng et al., 2023; Hueber and Kleyer, 2023; Golan et al., 2023; Dhanvijay et al., 2023; Sezgin et al., 2023; Team et al., 2024; Chen et al., 2023). Thus, L0+L1 should be paired with higher-level (L2+) decision-focused assessments before deployment.

Challenges. Attribution: Attribution concerns whether aggregated outputs actually rely on and are traceable to the retrieved sources. Automated audits in the medical domain report that models often cite papers that are only loosely relevant or that do not fully support the generated claims (Wu et al., 2025). Even with explicit attribution metrics and fine-grained factuality scoring, models can pass retrieval tests while still weaving in priors or mixing multiple sources in ways that obscure provenance (Yang et al., 2025). Techniques like Self-RAG improve on-demand retrieval and selfcritique, yet do not eliminate misattribution when sources disagree, are low quality, or when prompts nudge the model toward fluent synthesis over faithful quotation (Asai et al., 2023; Jung et al., 2024). Thus, attribution still remains as a challenge. Completeness: Completeness addresses whether transformed corpora and their aggregations cover all clinical facts (problems, meds, labs, temporality, context) without omissions. Comparative studies find that concept-recognition tools are inconsistent and often miss negations, misread abbreviations, and struggle with ambiguity or misspellings, which leads to missing or distorted facts downstream (Lossio-Ventura et al., 2023). Beyond extraction, mapping text spans to standard concepts is fragile because benchmark datasets contain many ambiguous terms and do not align well with UMLS coverage, so the 'correct' code is often unclear from the beginning (Newman-Griffis et al., 2021). In practice, long EHR narratives, events spread across notes, and uneven local coding leave gaps that retrieval can't fill when the structured substrate is incomplete. These issues persist because gold standards underrepresent edge cases, annotation guidelines vary, and many L1 evaluations emphasize span-level F1 or Recall@k over end-toend coverage of clinically critical fields (Lossio-Ventura et al., 2023). **Privacy & lineage:** It evaluates whether L1 transformations are governable and safe to share. De-identification of clinical notes reduces direct identifiers but does not preclude membership inference against downstream models; moreover, LLM-generated synthetic notes can carry similar privacy risks when they closely match real data utility (Sarkar et al., 2024). At the model layer, training data extraction and related attacks demonstrate that LLMs can memorize and regurgitate snippets of their training corpora (Carlini et al., 2021). These realities make rigorous lineage essential: machine-readable provenance (what data, which ETL/normalizers, which retrievers/rankers) should be recorded using established schemas so that organizations can audit and reproduce outputs (Lebo et al., 2013; Mitchell et al., 2019; Gebru et al., 2021). Yet provenance remains challenging in practice because pipelines are multi-hop and frequently updated; components are swapped or fine-tuned; and evidence bases evolve. Without disciplined, standard documentation, downstream users cannot reliably trace how a particular claim was produced. (Yang et al., 2025).

4.3 L2 - Decision Support

Definition and Scope. At autonomy level L2, the system provides patient-specific recommendations that can assist clinical decision making. By design, L2 depends on upstream EHR transformation and retrieval/aggregation (L1) but adds reasoning over

the individual's data and clinical context.

Typical Applications. Diagnostic reasoning: it reads a patient's history, exam, labs, and imaging to propose differentials with brief rationales and possible next steps; evaluations commonly use vignette-based case sets and prompting schemes that elicit stepwise clinical reasoning (Goh and colleagues, 2024; Savage et al., 2024a). Medication decision support: it integrates a patient's active medications, allergies, problems, and recent labs to surface potential adverse drug events, drug-drug interactions, and dosing/contraindication checks; widely used resources include the 2018 n2c2 ADE shared task (concept extraction, relation classification, and end-to-end pipelines) and the SemEval-2013 DDIExtraction task for literature-based DDI detection (Henry et al., 2020; Segura-Bedmar et al., 2013). Patient education. Patient education tools support patient-clinician communication by turning medical information into clear, usable messages and scaffolding two-way conversations. Typical functions include generating plain-language explanations of conditions, tests, and procedures; prompting patients to ask key questions; and using teach-back so patients restate key points to confirm understanding (Shoemaker et al., 2013; Centers for Disease Control and Prevention, 2020; Agency for Healthcare Research and Quality, 2023; Institute for Healthcare Improvement, n.d.; IPDAS Collaboration, 2024; Stacey et al., 2021). Taken together, these tasks are L2 because they require reasoning over an individual patient's context to generate recommendations or tailored explanations.

Evaluation Focus. For L2 decision support, evaluation should foreground three axes beyond L0's accuracy/completeness/readability: reasoning consistency, calibration and abstention, and safety. For reasoning, score the process, not just final answers: stepwise diagnostic logic and evidence use should be judged against clinician-authored rubrics or case rationales. For calibration, report reliability at the case level, plus selective prediction curves (risk—coverage) with a tunable "don't know / escalate" option. For safety, it is essential to track contraindication and guideline-violation rates.

Scope Boundaries. At L2, systems give decision support by combining a patient's EHR context with external evidence to produce advisory recommendations. However, they do not plan tasks, call tools or APIs, or change the record, so a clinician must review and decide. This preserves clinical authority but forces clinicians to turn text into orders and

messages, which adds workload and risks transcription errors. These limits motivate a shift to agentic L3 configurations that keep a human in the loop while reducing cognitive burden.

Challenges. Reasoning consistency & faithfulness: It concerns whether patient-specific answers are stable across prompt phrasings and whether the rationales actually support the recommendation. Comparative guideline evaluations show that changing formats or instructions can swing model outputs and agreement, underscoring prompt sensitivity in medical settings (Wang et al., 2024). Even when stepwise prompts are used to elicit reasoning, studies find that the generated "explanations" can be plausible yet unfaithful to the features that truly drove the prediction (Turpin et al., 2023; Madsen et al., 2024; Kuang et al., 2025; RRV et al., 2025). In diagnostic contexts, structured prompting can improve transparency but does not guarantee faithful causal grounding of the final answer (Savage et al., 2024b; Dineen et al., 2025). The result in practice is that two seemingly careful L2 prompts may yield different plans with rationales that read well but do not reliably reflect the model's decision process, which keeps consistency and faithfulness an open problem for deployment (Turpin et al., 2023). Confidence calibration: Here the question is whether stated or implicit confidence tracks correctness so that uncertain cases can be flagged or deferred. Recent medical evaluations show that simple proxies (e.g., log-probabilities) correlate only weakly with error and that token-probability or sampling-based methods improve ranking of uncertainty yet still leave pockets of overconfidence (Bentegeac et al., 2025). Semantic-entropy approaches detect confabulations and better prioritize abstention, but they often require multiple generations or added computation, and their reliability varies by task and domain (Farquhar et al., 2024; Kossen et al., 2024; Penny-Dimri et al., 2025). In real L2 use, these trade-offs mean systems may sound certain on incorrect recommendations or abstain too rarely on edge cases (Bentegeac et al., 2025; Farquhar et al., 2024; Feng et al., 2025b). Fairness: The core issue is whether recommendations generalize equitably across patient subgroups and clinical contexts. Specialized benchmarks and audit tools show that biases appear in longform answers and clinical recommendations that plain accuracy scores don't catch (Prakash et al., 2024). Purpose-built bias benchmarks for clinical LLMs report subgroup-linked shifts in outputs,

while triage studies using counterfactual tests reveal intersectional differences across sex and race (Zhang et al., 2024b; Lee et al., 2025b). A recent scoping review highlights uneven coverage across medical fields and limited clinician-in-the-loop evaluation-leaving blind spots in where and how biases manifest (Cheng et al., 2025). Consequently, subgroup reliability remains a persistent challenge (Prakash et al., 2024; Cheng et al., 2025).

4.4 L3 - Agents under human supervision

Definition and Scope. We define L3 agents as systems that plan and invoke tools/APIs to initiate actions in clinical workflows while keeping a clinician explicitly "in the loop" for review, modification, and sign-off. Core capabilities are task planning, retrieval, and safe tool use (e.g., querying/reading/writing to EHR endpoints), with human oversight enforced at key checkpoints (e.g., order "draft" states, queued messages/referrals).

Typical Applications. Clinical Copilot: clinical copilots plan tasks, fetch chart context, and then draft orders, messages, referrals, or care plans via tool/API calls, pausing for human sign-off-e.g., constellation designs that split a patient-facing agent from specialist agents (Polaris), pharmacy verification agents that stage indication/dose/interaction checks (Rx Strategist), and EHR copilots that navigate records and place draft orders (Almanac Copilot) (Mukherjee et al., 2024; Van et al., 2024; Zakka et al., 2024). Sandboxed **Simulation:** It embeds agents in controlled clinics to probe plan-act-check loops with audit trailsbenchmarks like AgentClinic (multimodal encounters with tool use), AI Hospital (multi-agent roles), ClinicalLab (multi-department diagnostics), and 3MDBench (telemedicine dialogue with assessor agents) (Schmidgall et al., 2024; Fan et al., 2024; Yan et al., 2024; Sviridov et al., 2025; Yue et al., 2025). Operation & EHR Automation: It coordinates non-diagnostic workflows (e.g., prior authorization) and return proposed actions for approval (RxLens; multi-agent prior-auth pipelines) (Jagatap et al., 2025). In contrast to L2 (textual recommendations), these L3 systems invoke tools to initiate actions but keep a clinician in the loop for review and sign-off.

Evaluation Focus. At L3, evaluation shifts from answer quality to supervised action quality: studies should first demonstrate end-to-end task success under human oversight. Second, they must verify tool-use correctness: every API call, order, and

parameter matches clinical intent. Finally, they should require auditability via machine-readable provenance of data, prompts, models, tools, parameters, and approvals, leveraging standards such as FHIR Provenance and AuditEvent so that actions can be traced for post-hoc review and governance. Scope Boundaries. L3 agents are limited to draftand-queue actions under explicit human oversight: they may plan tasks and call approved tools/APIs to prepare orders, messages, referrals, or documentation, but execution requires a clinician's independent review. Looking ahead, closed-loop system deployments can allow agents to autonomously execute pre-approved, low-risk steps when explicit policies and fine-grained access controls are satisfied, rather than requiring per-action approval.

Challenges. Tool-use failures: These arise when an agent plans correctly but issues the wrong API call or constructs malformed parameters, so the action fails even if the reasoning was sound. Recent EHR-agent benchmarks show that agents often get confused by FHIR's linked resources and multi-step queries, and many actions still fail because the API calls they issue are invalid or not FHIR-conformant (Lee et al., 2025a). This persists because health data are heterogeneous and nested, vendor FHIR implementations vary, and small prompt or formatting shifts can flip tool behavior with little semantic validation in the loop. Human-AI interaction dynamics: These refer to how clinicians review and sign off on queued actions. Studies of decision support show overreliance on automated suggestions and alert fatigue, where users ignore or over-accept system outputs, producing oversight gaps even when accuracy is reasonable (Khera et al., 2023; Abdelwanis et al., 2024). The problem endures because busy workflows, long sessions, and uneven interface cues make it hard to calibrate attention and to sustain critical review at every step. Auditability and accountability: These require a traceable record of what data were accessed, which tools were called with what parameters, and who approved the final action. Standards such as FHIR AuditEvent and Provenance define the necessary primitives, and governance frameworks emphasize logging and traceability for post-hoc review (HL7 International, 2025a,b; National Institute of Standards and Technology, 2023; World Health Organization, 2025). Yet multi-tool, multi-service agent stacks often lack end-to-end lineage across steps, so reconstructing a failure or near-miss remains difficult in practice.

5 Future Work

Closed-looped System We expect clinical deployments to shift from single-model helpers to closed-loop, hospital-scale systems composed of cooperating, role-specialized agents that escalate to clinicians at predefined gates (Schmidgall et al., 2024; Borkowski and Ben-Ari, 2025). We call for simple, auditable protocols for handoff, disagreement resolution, and safety gating of tool use, preceded by simulation and shadow deployments. We recommend reporting operational outcomes—deferral rates, near-misses, and governance practices—so the community can compare architectures and converge on safe patterns (Liu et al., 2020).

Guarantees. We call for reframing "good" performance around risk-controlled selectivity: in clinical settings, systems should act only when a target risk can be met and otherwise defer, evaluated by risk-coverage rather than average accuracy. In this agenda, reliable high accuracy on a subset with explicit deferral (e.g., 99% accuracy on 20% task) is preferable to broad moderate accuracy (e.g., 80% on 80%), because the former enables out-of-scope detection and smooth human routing while the latter obscures which cases are wrong. We recommend that future evaluations report calibrated confidence on the acted-on subset, and coverage at target risk with slice breakdowns (Guo et al., 2017).

6 Conclusion

This survey reframes medical LLMs evaluation around levels of autonomy (L0–L3) so that what a system is allowed to do matches the evidence required to trust it. We emphasize risk coverage rather than average accuracy: safe systems act only when they can meet a target risk and defer otherwise. At L0–L1, the focus is factual fidelity, bias, and structural correctness with clear grounding to sources. At L2–L3, the bar rises to calibrated uncertainty, selective answering, subgroup robustness, tool-use correctness, and verified human oversight.

Taken together, this level-conditioned lens turns benchmark scores into credible, clinically relevant claims. Evaluations are most persuasive when they make the level explicit, pair target risk with achieved coverage, report performance across slices and shifts, and verify both confidence calibration and oversight checkpoints. We hope this provides a clear language for building medical LLMs that are not only capable, but reliably useful—and safe—in practice.

Limitations

This survey is necessarily selective and timebounded; model releases, datasets, and guidance evolve rapidly, so some details may become outdated. Evidence in the literature remains uneven: many studies emphasize lab benchmarks over prospective or randomized evaluations, and reporting quality is inconsistent.

References

- Asad Aali, Dave Van Veen, Yamin Arefeen, Jason Hom, Christian Bluethgen, Eduardo Pontes Reis, Sergios Gatidis, Namuun Clifford, Joseph Daws, Arash Tehrani, Jangwon Kim, and Akshay Chaudhari. 2025. Mimic-iv-ext-bhc: Labeled clinical notes dataset for hospital course summarization. PhysioNet Repository.
- A. B. Abacha, E. Agichtein, Y. Pinter, and D. Demner-Fushman. 2017. Overview of the medical question answering task at trec 2017 liveqa. *TREC*, pages 1–12
- A. B. Abacha, Y. M'rabet, Y. Zhang, C. Shivade, C. Langlotz, and D. Demner-Fushman. 2021. Overview of the medical 2021 shared task on summarization in the medical domain,. *Proceedings of the 20th Workshop on Biomedical Language Processing*, pages 74–85.
- Mahyar Abbasian, Elahe Khatibi, Iman Azimi, David Oniani, Zahra Shakeri Hossein Abad, Alexander Thieme, Ram Sriram, Zhongqi Yang, Yanshan Wang, Bryant Lin, Olivier Gevaert, Li-Jia Li, Ramesh Jain, and Amir M. Rahmani. 2024. Foundation metrics for evaluating effectiveness of healthcare conversations powered by generative ai. *npj Digital Medicine*, 7(1):82.
- Moustafa Abdelwanis, Hamdan Khalaf Alarafati, Maram M. S. Tammam, and Mecit Can Emre Simsekler. 2024. Exploring the risks of automation bias in healthcare artificial intelligence applications: A bowtie analysis. *Journal of Safety Science and Resilience*, 5:460–469.
- Agency for Healthcare Research and Quality. 2023. Tool: Teach-back.
- E. Alsentzer, M. J. Rasmussen, R. Fontoura, A. L. Cull, B. Beaulieu-Jones, K. J. Gray, D. W. Bates, and V. P. Kovacheva. 2023. Zero-shot interpretable phenotyping of postpartum hemorrhage using large language models,. *NPJ Digital Medicine*, 6(1).
- Rahul K Arora, Jason Wei, Rebecca Soskin Hicks, Preston Bowman, Joaquin Quiñonero-Candela, Foivos Tsimpourlas, Michael Sharman, Meghan Shah, Andrea Vallone, Alex Beutel, and 1 others. 2025. Healthbench: Evaluating large language models towards improved human health. *arXiv preprint arXiv:2505.08775*.

- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. 2023. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. *arXiv preprint arXiv:2310.11511*.
- Elham Asgari, Nina Montaña-Brown, Magda Dubois, Saleh Khalil, Jasmine Balloch, Joshua Au Yeung, and Dominic Pimenta. 2025. A framework to assess clinical safety and hallucination rates of LLMs for medical text summarisation. *npj Digital Medicine*, 8(1):274.
- Sameer Badarudeen and Sanjeev Sabharwal. 2010. Assessing readability of patient education materials: Current role in orthopaedics. *Clinical Orthopaedics and Related Research*, 468(10):2572–2580.
- Suhana Bedi, Hejie Cui, Miguel Fuentes, Alyssa Unell, Michael Wornow, Juan M Banda, Nikesh Kotecha, Timothy Keyes, Yifan Mai, Mert Oez, and 1 others. 2025. Medhelm: Holistic evaluation of large language models for medical tasks. *arXiv preprint arXiv:2505.23802*.
- Asma Ben Abacha and Dina Demner-Fushman. 2019. On the summarization of consumer health questions. In *Proceedings of ACL*, pages 2228–2234.
- Asma Ben Abacha, Dina Demner-Fushman, Shweta Yadav, and Deepak Gupta. 2021. MEDIQA 2021: Consumer health question summarization, multi-answer summarization, and radiology report summarization. https://sites.google.com/view/mediqa2021. NAACL-BioNLP shared task site.
- Asma Ben Abacha, Wen-wai Yim, Yadan Fan, and Thomas Lin. 2023. An empirical study of clinical note generation from doctor—patient encounters. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*.
- Raphaël Bentegeac, Bastien Le Guellec, Grégory Kuchcinski, Philippe Amouyel, and Aghiles Hamroun. 2025. Token probabilities to mitigate large language models overconfidence in answering medical questions: Quantitative study. *Journal of Medical Internet Research*, 27:e64348.
- Elliot Bolton, Abhinav Venigalla, Michihiro Yasunaga, David Hall, Betty Xiong, Tony Lee, Roxana Daneshjou, Jonathan Frankle, Percy Liang, Michael Carbin, and Christopher D. Manning. 2024. Biomedlm: A 2.7b parameter language model trained on biomedical text. *Preprint*, arXiv:2403.18421.
- Andrew A. Borkowski and Alon Ben-Ari. 2025. Multiagent ai systems in health care: Envisioning next-generation intelligence. *Federal Practitioner*, 42(5):188–193.
- Glenn W. Brier. 1950. Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1):1–3.

- Nathan Brown, Marco Fiscato, Marwin Segler, and Alain C. Vaucher. 2019. Guacamol: Benchmarking models for de novo molecular design. *Journal* of Chemical Information and Modeling, 59(3):1096– 1108.
- Maria Ana Cardei, Josephine Lamp, Mark Derdzinski, and Karan Bhatia. 2025. Dexbench: Benchmarking llms for personalized decision making in diabetes management. *Preprint*, arXiv:2510.00038.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. 2021. Extracting training data from large language models. In *USENIX Security Symposium*, pages 2633–2650.
- Centers for Disease Control and Prevention. 2020. Clear communication index user guide.
- Guiming Hardy Chen, Shunian Chen, Ziche Liu, Feng Jiang, and Benyou Wang. 2024. Humans or Ilms as the judge? a study on judgement bias. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8301–8327. Association for Computational Linguistics.
- K. Chen, T. Litfin, J. Singh, J. Zhan, and Y. Zhou. 2023. The master database of all possible rna sequences and its integration with rnacmap for rna homology search,. *bioRxiv*, pages 2023–02.
- Tong Chen, Zimu Wang, Yiyi Miao, Haoran Luo, Yuanfei Sun, Wei Wang, Zhengyong Jiang, Procheta Sen, and Jionglong Su. 2025a. Medfact: A large-scale chinese dataset for evidence-based medical fact-checking of llm responses. *Preprint*, arXiv:2509.17436.
- Xiaolan Chen, Jiayang Xiang, Shanfu Lu, Yexin Liu, Mingguang He, and Danli Shi. 2025b. Evaluating large language models and agents in healthcare: key challenges in clinical applications. *Intelligent Medicine*.
- Lionel Tim-Ee Cheng, Jasmine Chiat Ling Ong, Zhen Ling Teo, Ting Fang Tan, Narrendar RaviChandran, Fei Wang, Leo Anthony Celi, Marcus Eng Hock Ong, and Nan Liu. 2025. A scoping review and evidence gap analysis of clinical AI fairness. *npj Digital Medicine*. PMCID: PMC12167363.
- Domenic V. Cicchetti and Alvan R. Feinstein. 1990. High agreement but low kappa: Ii. resolving the paradoxes. *Journal of Clinical Epidemiology*, 43(6):551–558.
- A. Cohan, S. Feldman, I. Beltagy, D. Downey, and D. S. Weld. 2020. Specter: Document-level representation learning using citation-informed transformers,. *arXiv* preprint arXiv:2004.07180.
- Arman Cohan, Franck Dernoncourt, Doo Soon Kim, Trung Bui, Seokhwan Kim, Walter Chang, and Nazli Goharian. 2018. A discourse-aware attention model

- for abstractive summarization of long documents. In Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers), pages 615–621, New Orleans, Louisiana. Association for Computational Linguistics.
- DeepSeek-AI. 2025. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *Preprint*, arXiv:2501.12948.
- Jay DeYoung, Iz Beltagy, Madeleine van Zuylen, Bailey Kuehl, and Lucy Lu Wang. 2021. Ms2: Multi-document summarization of medical studies. Preprint, arXiv:2104.06486.
- A. K. D. Dhanvijay, M. J. Pinjar, N. Dhokane, S. R. Sorte, A. Kumari, H. Mondal, and A. K. Dhanvijay. 2023. Performance of large language models (chatgpt, bing search, and google bard) in solving case vignettes in physiology, *Cureus*, 15(8).
- Jacob Dineen, Aswin RRV, Qin Liu, Zhikun Xu, Xiao Ye, Ming Shen, Zhaonan Li, Shijie Lu, Chitta Baral, Muhao Chen, and Ben Zhou. 2025. Qa-lign: Aligning llms through constitutionally decomposed qa. *Preprint*, arXiv:2506.08123.
- Zhihao Fan, Jialong Tang, Wei Chen, Siyuan Wang, Zhongyu Wei, Jun Xi, Fei Huang, and Jingren Zhou. 2024. Ai hospital: Benchmarking large language models in a multi-agent medical interaction simulator. *Preprint*, arXiv:2402.09742.
- Faiha Fareez, Tishya Parikh, Christopher Wavell, Saba Shahab, Meghan Chevalier, Scott Good, Isabella De Blasi, Rafik Rhouma, Christopher McMahon, Jean-Paul Lam, Thomas Lo, and Christopher W. Smith. 2022. A dataset of simulated patient—physician medical interviews with a focus on respiratory cases. *Scientific Data*, 9:313.
- Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. 2024. Detecting hallucinations in large language models using semantic entropy. *Nature*, 630(8017):625–630.
- Jana Fehr, Brian Citro, Rohit Malpani, Christoph Lippert, and Vince I. Madai. 2024. A trustworthy ai reality-check: the lack of transparency of artificial intelligence products in healthcare. *Frontiers in Digital Health*, 6:1267290.
- Alvan R. Feinstein and Domenic V. Cicchetti. 1990. High agreement but low kappa: I. the problems of two paradoxes. *Journal of Clinical Epidemiology*, 43(6):543–549.
- Kevin J. Feng, David W. McDonald, and Amy X. Zhang. 2025a. Levels of autonomy for ai agents. arXiv:2506.12469.
- Yu Feng, Ben Zhou, Weidong Lin, and Dan Roth. 2025b. Bird: A trustworthy bayesian inference framework for large language models. *Preprint*, arXiv:2404.12494.

- Paul Festor, Ibrahim Habli, Yan Jia, Anthony Gordon,
 A. Aldo Faisal, and Matthieu Komorowski. 2021.
 Levels of autonomy and safety assurance for ai-based clinical decision systems. In SAFECOMP 2021 Workshops, pages 291–296. Springer.
- Jack Gallifant, Majid Afshar, Saleem Ameen, Yindalon Aphinyanaphongs, Sherri Chen, Giovanni E. Cacciamani, Dina Demner-Fushman, Dmitriy Dligach, Roxana Daneshjou, Christopher Fernandes, and 1 others. 2025. The tripod-llm reporting guideline for studies using large language models. *Nature Medicine*, 31(1):60–69.
- Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. 2021. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92.
- Yonatan Geifman and Ran El-Yaniv. 2017. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems*, volume 30.
- Yonatan Geifman and Ran El-Yaniv. 2019. Selectivenet: A deep neural network with an integrated reject option. In *Proceedings of the 36th International Conference on Machine Learning (ICML 2019)*, volume 97 of *Proceedings of Machine Learning Research*, pages 2151–2159. PMLR.
- Elizabeth Goh and colleagues. 2024. Large language model influence on diagnostic reasoning. *JAMA Network Open*.
- R. Golan, S. J. Ripps, R. Reddy, J. Loloi, A. P. Bernstein, Z. M. Connelly, N. S. Golan, R. Ramasamy, S. Ripps, and R. V. Reddy. 2023. Chatgpt's ability to assess quality and readability of online medical information: evidence from a cross-sectional study,. *Cureus*, 15(7).
- Jiawei Gu, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, Honghao Liu, Saizhuo Wang, Kun Zhang, Yuanzhuo Wang, Wen Gao, Lionel Ni, and Jian Guo. 2024. A survey on llm-as-a-judge. *arXiv* preprint arXiv:2411.15594.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML 2017)*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR.
- Paul Hager, Friederike Jungmann, Robbie Holland, Kunal Bhagat, Inga Hubrecht, Manuel Knauer, Jakob Vielhauer, Marcus Makowski, Rickmer Braren, Georgios Kaissis, and Daniel Rueckert. 2024. Evaluation and mitigation of the limitations of large language models in clinical decision-making. *Nature Medicine*, 30(9):2613–2622.

- D. Hendrycks, C. Burns, S. Basart, A. Zou, M. Mazeika, D. Song, and J. Steinhardt. 2020. Measuring massive multitask language understanding,. arXiv preprint arXiv:2009.03300.
- Sam Henry, Kevin Buchan, Michele Filannino, Amber Stubbs, and Özlem Üzuner. 2020. 2018 n2c2 shared task on adverse drug events and medication extraction in electronic health records. *Journal of the American Medical Informatics Association*, 27(1):3–12.
- HL7 International. 2025a. Hl7 fhir auditevent resource. https://build.fhir.org/auditevent.html. Accessed Oct 5, 2025.
- HL7 International. 2025b. Hl7 fhir provenance resource. https://build.fhir.org/provenance.html. Accessed Oct 5, 2025.
- Jaycelyn R. Holland, Donald H. Arnold, Holly R. Hanson, Barbara J. Solomon, Nicholas E. Jones, Tucker W. Anderson, Wu Gong, Christopher J. Lindsell, Travis W. Crook, and Daisy A. Ciener. 2022. Reliability of the behaviorally anchored rating scale (bars) for assessing non-technical skills of medical students in simulated scenarios. *Medical Education Online*, 27(1).
- A. J. Hueber and A. Kleyer. 2023. Quality of citation data using the natural language processing tool chatgpt in rheumatology: creation of false references,. *RMD open*, 9(2).
- Institute for Healthcare Improvement. n.d. Ask me 3: Good questions for your good health. Accessed 2025-10-05.
- IPDAS Collaboration. 2024. International patient decision aid standards (ipdas) collaboration. Accessed 2025-10-05.
- Akshay Jagatap, Srujana Merugu, and Prakash Mandayam Comar. 2025. Rxlens: Multi-agent Ilmpowered scan and order for pharmacy. In Proceedings of the 2025 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 3: Industry Track), pages 822–832, Albuquerque, New Mexico. Association for Computational Linguistics.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Delong Chen, Wenliang Dai, Ho Shu Chan, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12).
- D. Jin, E. Pan, N. Oufattole, W.-H. Weng, H. Fang, and P. Szolovits. 2021. What disease does this patient have? a large-scale open domain question answering dataset from medical exams,. *Applied Sciences*, 11(14).
- Q. Jin, W. Kim, Q. Chen, D. C. Comeau, L. Yeganova, W. J. Wilbur, and Z. Lu. 2023. Medcpt: Contrastive pre-trained transformers with large-scale pubmed

- search logs for zero-shot biomedical information retrieval,. *Bioinformatics*, 39(11).
- A. E. Johnson, T. J. Pollard, S. J. Berkowitz, N. R. Greenbaum, M. P. Lungren, C.-y. Deng, R. G. Mark, and S. Horng. 2019. Mimic-cxr, a de-identified publicly available database of chest radiographs with free-text reports,. *Scientific data*, 6(1).
- Hailey Joren, Jianyi Zhang, Chun-Sung Ferng, Da-Cheng Juan, Ankur Taly, and Cyrus Rashtchian. 2025. Sufficient context: A new lens on retrieval augmented generation systems. *Preprint*, arXiv:2411.06037.
- Dongwon Jung, Qin Liu, Tenghao Huang, Ben Zhou, and Muhao Chen. 2024. Familiarity-aware evidence compression for retrieval-augmented generation. *Preprint*, arXiv:2409.12468.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Samuel Bowman, Stanislav Fort, and 17 others. 2022. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*.
- Nikhil Khandekar, Qiao Jin, Guangzhi Xiong, Soren Dunn, Serina S. Applebaum, Zain Anwar, Maame Sarfo-Gyamfi, Conrad W. Safranek, Abid A. Anwar, Andrew Zhang, Aidan Gilson, Maxwell B. Singer, Amisha Dave, Andrew Taylor, Aidong Zhang, Qingyu Chen, and Zhiyong Lu. 2024. Medcalcbench: Evaluating large language models for medical calculations. *Preprint*, arXiv:2406.12036.
- Rohan Khera, Melissa A. Simon, and Joseph S. Ross. 2023. Automation bias and assistive ai: Risk of harm from ai-driven clinical decision support. *JAMA*, 330(23):2255–2257.
- Anita Kirkovska, Nico Finelli, and David Vargas. 2025. The six levels of agentic behavior. Vellum AI Blog.
- Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. 2024. Semantic entropy probes: Robust and cheap hallucination detection in llms. arXiv preprint arXiv:2406.15927.
- Jiayi Kuang, Haojing Huang, Yinghui Li, Xinnian Liang, Zhikun Xu, Yangning Li, Xiaoyu Tan, Chao Qu, Meishan Zhang, Ying Shen, and Philip S. Yu. 2025. Atomic thinking of llms: Decoupling and exploring mathematical reasoning abilities. *Preprint*, arXiv:2509.25725.
- Ananya Kumar, Percy Liang, and Tengyu Ma. 2019. Verified uncertainty calibration. In *Advances in Neural Information Processing Systems*, volume 32.
- Timothy Lebo, Satya S. Sahoo, Deborah L. McGuinness, and 1 others. 2013. PROV-O: The PROV ontology. W3C Recommendation.

- Gyubok Lee, Elea Bach, Eric Yang, Tom Pollard, Alistair Johnson, Edward Choi, Yugang Jia, and Jong Ha Lee. 2025a. Fhir-agentbench: Benchmarking llm agents for realistic interoperable ehr question answering. *Preprint*, arXiv:2509.19319.
- Joseph Lee, Tianqi Shang, Jae Young Baik, Duy Duong-Tran, Shu Yang, Lingyao Li, and Li Shen. 2025b. From promising capability to pervasive bias: Assessing large language models for emergency department triage. *arXiv* preprint arXiv:2504.16273.
- Bangzheng Li, Ben Zhou, Xingyu Fu, Fei Wang, Dan Roth, and Muhao Chen. 2024a. Famicom: Further demystifying prompts for language models with task-agnostic performance estimation. *Preprint*, arXiv:2406.11243.
- Bangzheng Li, Ben Zhou, Fei Wang, Xingyu Fu, Dan Roth, and Muhao Chen. 2024b. Deceptive semantic shortcuts on reasoning chains: How far can models go without hallucination? *Preprint*, arXiv:2311.09702.
- Haitao Li, Qian Dong, Junjie Chen, Huixue Su, Yujia Zhou, Qingyao Ai, Ziyi Ye, and Yiqun Liu. 2024c. Llms-as-judges: A comprehensive survey on llm-based evaluation methods. arXiv preprint arXiv:2412.05579.
- Jiao Li, Yueping Sun, Robin J. Johnson, Daniela Sciaky, Chih-Hsuan Wei, Robert Leaman, Allan Peter Davis, Carolyn J. Mattingly, Thomas C. Wiegers, and Zhiyong Lu. 2016. Biocreative v cdr task corpus: a resource for chemical–disease relation extraction. *Database*, 2016:baw068.
- Yifei Li, Xiang Yue, Zeyi Liao, and Huan Sun. 2024d. Attributionbench: How hard is automatic attribution evaluation? *Preprint*, arXiv:2402.15089.
- Chin-Yew Lin. 2004. ROUGE: A package for automatic evaluation of summaries. In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain. Association for Computational Linguistics.
- G. Liu, Y. Liao, F. Wang, B. Zhang, L. Zhang, X. Liang, X. Wan, S. Li, Z. Li, and S. Zhang. 2021. Medical-vlbert: Medical visual language bert for covid-19 ct report generation with alternate learning. *IEEE Transactions on Neural Networks and Learning Systems*, 32(9):3786–3797.
- W. Liu, J. Tang, Y. Cheng, W. Li, Y. Zheng, and X. Liang. 2022. Meddg: an entity-centric medical consultation dataset for entity-aware medical dialogue generation,. CCF International Conference on Natural Language Processing and Chinese Computing, pages 447–459.
- Xiaoxuan Liu, S. C. Rivera, D. Moher, and et al. 2020. Consort-ai extension: Reporting guidelines for clinical trial reports for interventions involving artificial intelligence. *Nature Medicine*, 26:1364–1374.

- Juan Antonio Lossio-Ventura, Ran Sun, Sebastien Boussard, and Tina Hernandez-Boussard. 2023. Clinical concept recognition: Evaluation of existing systems on ehrs. *Frontiers in Artificial Intelligence*, 5:1051724.
- Zhong Lou, Chen Zhao, Wenhao Li, Chenyu Zhang, Qi Zhu, Xuehai Pan, Mohit Bansal, and Kai-Wei Chang. 2024. Anchoring effects in large language models. *arXiv preprint arXiv:2412.06593*.
- Ling Luo, Po-Ting Lai, Chih-Hsuan Wei, Cecilia N. Arighi, and Zhiyong Lu. 2022. Biored: A rich biomedical relation extraction dataset. *Briefings in Bioinformatics*.
- Zizhan Ma, Wenxuan Wang, Guo Yu, Yiu-Fai Cheung, Meidan Ding, Jie Liu, Wenting Chen, and Linlin Shen. 2025. Beyond the leaderboard: Rethinking medical benchmarks for large language models. *Preprint*, arXiv:2508.04325.
- Andreas Madsen, Sarath Chandar, and Siva Reddy. 2024. Are self-explanations from large language models faithful? *arXiv preprint arXiv:2401.07927*.
- Diwakar Mahajan, Jennifer J. Liang, Ching-Huei Tsou, and Özlem Uzuner. 2023. Overview of the 2022 n2c2 shared task on contextualized medication event extraction in clinical notes. *Journal of Biomedical Informatics*, 144:104432.
- Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. 2008. *Introduction to Information Retrieval*. Cambridge University Press, Cambridge, UK.
- Joshua Maynez, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. 2020. On faithfulness and factuality in abstractive summarization. In *Proceedings* of the 58th Annual Meeting of the Association for Computational Linguistics, pages 1906–1919, Online. Association for Computational Linguistics.
- Mary L. McHugh. 2012. Interrater reliability: the kappa statistic. *Biochemia Medica*, 22(3):276–282.
- Nikita Mehandru, Niloufar Golchini, David Bamman, Travis Zack, Melanie F. Molina, and Ahmed Alaa. 2025. Er-reason: A benchmark dataset for llm-based clinical reasoning in the emergency room. *Preprint*, arXiv:2505.22919.
- Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. 2019. Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT*)*, pages 220–229.
- Subhabrata Mukherjee, Paul Gamble, Markel Sanz Ausin, Neel Kant, Kriti Aggarwal, Neha Manjunath, Debajyoti Datta, Zhengliang Liu, Jiayuan Ding, Sophia Busacca, Cezanne Bianco, Swapnil Sharma, Rae Lasko, Michelle Voisard, Sanchay Harneja, Darya Filippova, Gerry Meixiong, Kevin

- Cha, Amir Youssefi, and 7 others. 2024. Polaris: A safety-focused llm constellation architecture for healthcare. *Preprint*, arXiv:2403.13313.
- National Institute of Standards and Technology. 2023. Artificial intelligence risk management framework (ai rmf 1.0). https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf. Accessed Oct 5, 2025.
- Zabir Al Nazi and Wei Peng. 2024. Large language models in healthcare and medical domain: A review. *Informatics*, 11(3):57.
- Denis Newman-Griffis, Guy Divita, Bart Desmet, Ayah Zirikly, Carolyn P Rosé, and Eric Fosler-Lussier. 2021. Ambiguity in medical concept normalization: An analysis of types and coverage in electronic health record datasets. *Journal of the American Medical Informatics Association*, 28(3):516–532.
- Jeremy Nixon, Mike Dusenberry, Ghassen Jerfel, Timothy Nguyen, Jeremiah Liu, Linchuan Zhang, and Dustin Tran. 2019. Measuring calibration in deep learning. arXiv preprint arXiv:1904.01685.
- S. Nowak, D. Biesner, Y. Layer, M. Theis, H. Schneider, W. Block, B. Wulff, U. Attenberger, R. Sifa, and A. Sprinkart. 2023. Transformer-based structuring of free-text radiology report databases,. *European Radiology*, 33(6):4228–4236.
- Juliana Damasio Oliveira, Henrique D. P. Santos, Ana Helena D. P. S. Ulbrich, Julia Colleoni Couto, Marcelo Arocha, Joaquim Santos, Manuela Martins Costa, Daniela Faccio, Fabio O. Tabalipa, and Rodrigo F. Nogueira. 2025. Development and evaluation of a clinical note summarization system using large language models. *Communications Medicine*, 5(1):376.
- Jesutofunmi A. Omiye, Jenna C. Lester, Simon Spichak, Veronica Rotemberg, and Roxana Daneshjou. 2023. Large language models propagate race-based medicine. *npj Digital Medicine*, 6:195.
- OpenAI. 2025. Introducing GPT-5. https://openai.com/index/introducing-gpt-5/. Accessed: 2025-10-02.
- L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, and A. Ray. 2022. Training language models to follow instructions with human feedback, *Advances in neural information processing systems*, 35.
- A. Pal, L. K. Umapathi, and M. Sankarasubbu. 2022. Medmcqa: A large-scale multi-subject multi-choice dataset for medical domain question answering, *Conference on health, inference, and learning*, pages 248–260.
- C. Peng, X. Yang, A. Chen, K. E. Smith, N. PourNejatian, A. B. Costa, C. Martin, M. G. Flores, Y. Zhang, and T. Magoc. 2023. A study of generative large language model for medical research and healthcare,. *NPJ Digital Medicine*, 6(1).

- Jahan C. Penny-Dimri, Magdalena Bachmann, William R. Cooke, Sam Mathewlynn, Samuel Dockree, John Tolladay, Jannik Kossen, Lin Li, Yarin Gal, and Gabriel Davis Jones. 2025. Reducing large language model safety risks in women's health using semantic entropy. *Preprint*, arXiv:2503.00269.
- Pouya Pezeshkpour and Estevam Hruschka. 2023. Large language models sensitivity to the order of options in multiple-choice questions. *arXiv* preprint *arXiv*:2308.11483.
- Daniil Polykovskiy, Alexander Zhebrak, Benjamin Sanchez-Lengeling, Sergey Golovanov, Oktai Tatanov, Stanislav Belyaev, Rauf Kurbanov, Aleksey Artamonov, Vladimir Aladinskiy, Mark Veselov, Artur Kadurin, Simon Johansson, Hongming Chen, Sergey Nikolenko, Alan Aspuru-Guzik, and Alex Zhavoronkov. 2018. Molecular sets (moses): A benchmarking platform for molecular generation models. *Preprint*, arXiv:1811.12823.
- Maja Popović. 2015. chrf: character n-gram f-score for automatic MT evaluation. In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395, Lisbon, Portugal. Association for Computational Linguistics.
- Sushant Prakash, Katherine Heller, Alan Karthike-salingam, Christopher Semturs, Joelle Barral, Greg Corrado, Yossi Matias, Jamila Smith-Loud, Ivor Horn, and Karan Singhal. 2024. A toolbox for surfacing health equity harms and biases in large language models. *Nature Medicine*, 30:3590–3600.
- Qwen Team. 2025. Qwen3 technical report. arXiv preprint arXiv:2505.09388.
- Kirk Roberts, Tasmeer Alam, Steven Bedrick, Dina Demner-Fushman, Kyle Lo, Ian Soboroff, Ellen Voorhees, Lucy Lu Wang, and William R. Hersh. 2021. Searching for scientific evidence in a pandemic: An overview of trec-covid. *Journal of Biomedical Informatics*, 121:103865.
- Aswin RRV, Jacob Dineen, Divij Handa, Md Nayem Uddin, Mihir Parmar, Chitta Baral, and Ben Zhou. 2025. Thinktuning: Instilling cognitive reflections without distillation. *Preprint*, arXiv:2508.07616.
- Atiquer Rahman Sarkar, Yao-Shun Chuang, Noman Mohammed, et al., and Xiaoqian Jiang. 2024. Deidentification is not enough: a comparison between de-identified and synthetic clinical notes. *Scientific Reports*, 14(29669).
- Thomas Savage, Ashwin Nayak, Robert Gallo, Ekanath Rangan, and Jonathan H. Chen. 2024a. Diagnostic reasoning prompts reveal the potential for large language model interpretability in medicine. *npj Digital Medicine*, 7(20).
- Thomas Savage, Ashwin Nayak, Robert Gallo, Ekanath Rangan, and Jonathan H. Chen. 2024b. Diagnostic reasoning prompts reveal the potential for large language model interpretability in medicine. *npj Digital Medicine*, 7(1):20.

- Samuel Schmidgall, Rojin Ziaei, Carl Harris, Eduardo Reis, Jeffrey Jopling, and Michael Moor. 2024. Agentclinic: a multimodal agent benchmark to evaluate ai in simulated clinical environments. *Preprint*, arXiv:2405.07960.
- Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2023. Quantifying language models' sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. *arXiv* preprint *arXiv*:2310.11324. ICLR 2024 camera-ready version.
- Isabel Segura-Bedmar, Paloma Martínez, and David Herrero-Zazo. 2013. Semeval-2013 task 9: Extraction of drug-drug interactions from biomedical texts (ddiextraction 2013). In *Proceedings of SemEval 2013*.
- Thibault Sellam, Dipanjan Das, and Ankur Parikh. 2020. BLEURT: Learning robust metrics for text generation. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7881–7892. Association for Computational Linguistics.
- E. Sezgin, F. Chekeni, J. Lee, and S. Keim. 2023. Clinical accuracy of large language models and google search responses to postpartum depression questions: cross-sectional study,. *Journal of Medical Internet Research*, 25.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2023. Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*.
- Ming Shen, Zhikun Xu, Jacob Dineen, Xiao Ye, and Ben Zhou. 2025. Bow: Reinforcement learning for bottlenecked next word prediction. *Preprint*, arXiv:2506.13502.
- Sarah J. Shoemaker, Michael S. Wolf, and Cindy Brach. 2013. *The Patient Education Materials Assessment Tool (PEMAT) and User's Guide (Version 1.0)*. Agency for Healthcare Research and Quality, Rockville, MD. AHRQ Publication No. 14-0002-EF. Updated August 2014.
- Sina Shool, Sara Adimi, Reza Saboori Amleshi, Ehsan Bitaraf, Reza Golpira, and Mahmood Tara. 2025. A systematic review of large language model (llm) evaluations in clinical medicine. *BMC Medical Informatics and Decision Making*, 25(117).
- Som Singh, Aleena Jamal, and Fawad Qureshi. 2024. Readability metrics in patient education: Where do we innovate? *Clinics and Practice*, 14(6):2341–2349.

- K. Singhal, S. Azizi, T. Tu, S. S. Mahdavi, J. Wei, H. W. Chung, N. Scales, A. Tanwani, H. Cole-Lewis, and S. Pfohl. 2023. Large language models encode clinical knowledge,. *Nature*, 620(7972):172–180.
- Karan Singhal, Tao Tu, Juraj Gottweis, Rory Sayres, Ellery Wulczyn, Mohamed Amin, Le Hou, Kevin Clark, Stephen R. Pfohl, Heather Cole-Lewis, Darlene Neal, Qazi Mamunur Rashid, Mike Schaekermann, Amy Wang, Dev Dash, Jonathan H. Chen, Nigam H. Shah, Sami Lachgar, Philip Andrew Mansfield, and 16 others. 2025. Toward expert-level medical question answering with large language models. *Nature Medicine*, 31(3):943–950.
- Ali Soroush, Benjamin S. Glicksberg, Eyal Zimlichman, Yiftach Barash, Robert Freeman, Alexander W. Charney, Girish N. Nadkarni, and Eyal Klang. 2024. Large language models are poor medical coders—benchmarking of medical code querying. *NEJM AI*, 1(5).
- Adarsh Srinivasan, Jacob Dineen, Muhammad Umar Afzal, Muhammad Uzair Sarfraz, Irbaz B. Riaz, and Ben Zhou. 2025. Recap: Transparent inference-time emotion alignment for medical dialogue systems. *Preprint*, arXiv:2509.10746.
- Dawn Stacey, Robert J. Volk, and IPDAS Evidence Update Leads. 2021. The international patient decision aid standards (ipdas) collaboration: Evidence update 2.0. *Medical Decision Making*, 41(7):729–733.
- Mark Steyvers, Heliodoro Tejeda, Aakriti Kumar, Catarina Belem, Sheer Karny, Xinyue Hu, Lukas W. Mayer, and Padhraic Smyth. 2025. What large language models know and what people think they know. *Nature Machine Intelligence*, 7:221–231.
- Ivan Sviridov, Amina Miftakhova, Artemiy Tereshchenko, Galina Zubkova, Pavel Blinov, and Andrey Savchenko. 2025. 3mdbench: Medical multimodal multi-agent dialogue benchmark. arXiv preprint arXiv:2504.13861.
- Thomas Yu Chow Tam, Sonish Sivarajkumar, Sumit Kapoor, Alisa V. Stolyar, Katelyn Polanska, Karleigh R. McCarthy, Hunter Osterhoudt, Xizhi Wu, Shyam Visweswaran, Sunyang Fu, Piyush Mathur, Giovanni E. Cacciamani, Cong Sun, Yifan Peng, and Yanshan Wang. 2024. A framework for human evaluation of large language models in healthcare derived from literature review. *npj Digital Medicine*, 7(1):258.
- Sijun Tan, Siyuan Zhuang, Kyle Montgomery, William Y. Tang, Alejandro Cuadron, Chenguang Wang, Raluca Ada Popa, and Ion Stoica. 2025. Judgebench: A benchmark for evaluating llm-based judges. In *Proceedings of the International Conference on Learning Representations (ICLR)*. ICLR 2025 (published version); arXiv:2410.12784.
- Ting Fang Tan, Kabilan Elangovan, Jasmine Ong, Nigam Shah, Joseph Sung, Tien Yin Wong, Lan Xue, Nan Liu, Haibo Wang, Chang Fu Kuo, Simon

- Chesterman, Zee Kin Yeong, and Daniel S. W. Ting. 2024. A proposed s.c.o.r.e. evaluation framework for large language models: Safety, consensus, objectivity, reproducibility and explainability. *arXiv preprint arXiv:2407.07666*.
- Liyan Tang, Zhaoyi Sun, Betina Idnay, Jordan G. Nestor, Ali Soroush, Pierre A. Elias, Ziyang Xu, Ying Ding, Greg Durrett, Justin F. Rousseau, Chunhua Weng, and Yifan Peng. 2023. Evaluating large language models on medical evidence summarization. *npj Digital Medicine*, 6(1):158.
- G. Team, T. Mesnard, C. Hardin, R. Dadashi, S. Bhupatiraju, S. Pathak, L. Sifre, M. Rivi'ere, M. S. Kale, and J. Love. 2024. Gemma: Open models based on gemini research and technology, arXiv preprint arXiv:2403.08295.
- N. Thakur, N. Reimers, A. R""uckl'e, A. Srivastava, and I. Gurevych. 2021. Beir: A heterogenous benchmark for zero-shot evaluation of information retrieval models.. arXiv preprint arXiv:2104.08663.
- Jeremias Traub, Till J. Bungert, Carsten T. Lüth, Michael Baumgartner, Klaus H. Maier-Hein, Lena Maier-Hein, and Paul F. Jaeger. 2024. Overcoming common flaws in the evaluation of selective classification systems. In *Advances in Neural Information Processing Systems*. NeurIPS 2024 (spotlight).
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Özlem Uzuner, Brett R. South, Shuying Shen, and Scott L. DuVall. 2011. 2010 i2b2/va challenge on concepts, assertions, and relations in clinical text. *Journal of the American Medical Informatics Association*, 18(5):552–556.
- Phuc Phan Van, Dat Nguyen Minh, An Dinh Ngoc, and Huy Phan Thanh. 2024. Rx strategist: Prescription verification using Ilm agents system. *Preprint*, arXiv:2409.03440.
- Li Wang, Xi Chen, XiangWen Deng, Hao Wen, MingKe You, WeiZhi Liu, Qi Li, and Jian Li. 2024. Prompt engineering in consistency and reliability with the evidence-based guideline for LLMs. *npj Digital Medicine*, 7(1):41.
- Lih-Wern Wang, Michael J. Miller, Michael R. Schmitt, and Frances K. Wen. 2013. Assessing readability formula differences with written health information materials: application, results, and recommendations. *Research in Social and Administrative Pharmacy*, 9(5):503–516.
- Lucy Lu Wang, Jay DeYoung, and Byron Wallace. 2022. Overview of MSLR2022: A shared task on multidocument summarization for literature reviews. In

- Proceedings of the Third Workshop on Scholarly Document Processing, pages 175–180, Gyeongju, Republic of Korea. Association for Computational Linguistics.
- Xiaosong Wang, Yifan Peng, Le Lu, and 1 others. 2017. Chestx-ray8: Hospital-scale chest x-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 3462–3471.
- Koki Wataoka, Tsubasa Takahashi, and Ryokan Ri. 2024. Self-preference bias in llm-as-a-judge. *arXiv* preprint arXiv:2410.21819.
- World Health Organization. 2025. Ethics and governance of artificial intelligence for health: Guidance on large multi-modal models. https://www.who.int/publications/i/item/9789240084759. Accessed Oct 5, 2025.
- K. Wu, R. Wang, H. Quan, C. Lin, E. Lo, H. Tu, F. Yu, and H. T. Lin. 2025. An automated framework for assessing how well large language models cite relevant medical references. *Nature Communications*, 16(1):427.
- Guangzhi Xiong, Qiao Jin, Zhiyong Lu, and Aidong Zhang. 2024. Benchmarking retrieval-augmented generation for medicine. *Preprint*, arXiv:2402.13178.
- Zhikun Xu, Ming Shen, Jacob Dineen, Zhaonan Li, Xiao Ye, Shijie Lu, Aswin RRV, Chitta Baral, and Ben Zhou. 2025. Tow: Thoughts of words improve reasoning in large language models. *Preprint*, arXiv:2410.16235.
- Weixiang Yan, Haitian Liu, Tengxiao Wu, Qian Chen, Wen Wang, Haoyuan Chai, Jiayi Wang, Weishan Zhao, Yixin Zhang, Renjun Zhang, Li Zhu, and Xuandong Zhao. 2024. Clinicallab: Aligning agents for multi-departmental clinical diagnostics in the real world. *Preprint*, arXiv:2406.13890.
- Rui Yang, Yilin Ning, Emilia Keppo, Mingxuan Liu, Chuan Hong, Danielle S. Bitterman, Jasmine C. L. Ong, Daniel S. W. Ting, and Nan Liu. 2025. Retrieval-augmented generation for generative artificial intelligence in health care. *npj Health Systems*, 2(2):2.
- Jiayi Ye, Yanbo Wang, Yue Huang, Dongping Chen, Qihui Zhang, Nuno Moniz, Tian Gao, Werner Geyer, Chao Huang, Pin-Yu Chen, Nitesh V. Chawla, and Xiangliang Zhang. 2025a. Justice or prejudice? quantifying biases in llm-as-a-judge. In Proceedings of the International Conference on Learning Representations (ICLR). ICLR 2025 Poster; arXiv:2410.02736.
- Xiao Ye, Shaswat Shrivastava, Zhaonan Li, Jacob Dineen, Shijie Lu, Avneet Ahuja, Ming Shen, Zhikun Xu, and Ben Zhou. 2025b. Cc-learn: Cohort-based consistency learning. *Preprint*, arXiv:2506.15662.

- Wen-wai Yim, Yujuan Fu, Asma Ben Abacha, Neal Snider, Thomas Lin, and Meliha Yetisgen. 2023. Acibench: A novel ambient clinical intelligence dataset for benchmarking automatic visit note generation. *Scientific Data*, 10.
- Matthew Yue, Zhikun Xu, Vivek Gupta, Thao Ha, Liesal Sharabi, and Ben Zhou. 2025. Relate-sim: Leveraging turning point theory and llm agents to predict and understand long-term relationship dynamics through interactive narrative simulations. *Preprint*, arXiv:2510.00414.
- Cyril Zakka, Joseph Cho, Gracia Fahed, Rohan Shad, Michael Moor, Robyn Fong, Dhaman-preet Kaur, Vishnu Ravi, Oliver Aalami, Roxana Daneshjou, Akshay Chaudhari, and William Hiesinger. 2024. Almanac copilot: Towards autonomous electronic health record navigation. *Preprint*, arXiv:2405.07896.
- Guangtao Zeng, Wenmian Yang, Zeqian Ju, Yue Yang, Sicheng Wang, Ruisi Zhang, Meng Zhou, Jiaqi Zeng, Xiangyu Dong, Ruoyu Zhang, Hongchao Fang, Penghui Zhu, Shu Chen, and Pengtao Xie. 2020. Meddialog: Large-scale medical dialogue datasets. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9241–9250, Online. Association for Computational Linguistics.
- N. Zhang, M. Chen, Z. Bi, X. Liang, L. Li, X. Shang, K. Yin, C. Tan, J. Xu, and F. Huang. 2021. Cblue: A chinese biomedical language understanding evaluation benchmark,. *arXiv preprint arXiv:2106.08087*.
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. BERTScore: Evaluating text generation with BERT. In *International Conference on Learning Representations* (ICLR 2020).
- Y. Zhang, M. Lang, J. Jiang, Z. Gao, F. Xu, T. Litfin, K. Chen, J. Singh, X. Huang, and G. Song. 2024a. Multiple sequence alignment-based rna language model and its application to structural inference,. *Nucleic Acids Research*, 52(1).
- Yubo Zhang, Shudi Hou, Mingyu Derek Ma, Wei Wang, Muhao Chen, and Jieyu Zhao. 2024b. CLIMB: A benchmark of clinical bias in large language models. *Preprint*, arXiv:2407.05250.
- Zhengyun Zhao, Qiao Jin, Fangyuan Chen, Tuorui Peng, and Sheng Yu. 2023. A large-scale dataset of patient summaries for retrieval-based clinical decision support systems. *Scientific Data*, 10(1):909.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena. In NeurIPS 2023 Datasets and Benchmarks Track.

Ben Zhou, Hongming Zhang, Sihao Chen, Dian Yu, Hongwei Wang, Baolin Peng, Dan Roth, and Dong Yu. 2024. Conceptual and unbiased reasoning in language models. *Preprint*, arXiv:2404.00205.

Lianghui Zhu, Xinggang Wang, and Xinlong Wang. 2025. Judgelm: Fine-tuned large language models are scalable judges. In *Proceedings of the International Conference on Learning Representations (ICLR)*. ICLR 2025 Spotlight; arXiv:2310.17631.

A Appendix

A.1 Full Table

A.2 LLM Usage

We used LLMs for language polishing and organization only; all technical claims and citations were authored and verified by the authors

Table 1: Expanded benchmarks by task class. Each row lists what the task can establish for clinical readiness, common evaluation metrics, the target level (L0-L3), and representative benchmarks.

Task	Usefulness	Metric	Level	Benchmark
Knowledge recall / exam Q&A	Breadth of factual recall and clinical reasoning on multiple-choice or open-ended exam questions across specialties and difficulty levels. Evaluates recall under constrained formats and coverage gates, and whether models hallucinate when unsure (Arora et al., 2025).	Exact match/F1; accuracy; subject/difficulty slices; contamination checks; calibration on unanswerable or uncertain questions.	LO	MedMCQA (Pal et al., 2022); MultiMedQA (MedQA, MedMCQA, PubMedQA, MMLU) (Singhal et al., 2023); Mirage RAG suite (MMLU-Med, MedQA-US, PubMedQA, BioASQ Y/N) (Xiong et al., 2024)
Summarization / transformation	Fidelity and completeness of summaries or transformations of clinical notes, conversations, or research literature. Tests whether models can produce coherent, structurally complete summaries without omissions or hallucinations (DeYoung et al., 2021; Aali et al., 2025).	Hallucination/omission rate; ROUGE/BERTScore/chrF; section completeness; clinical correctness; expert ratings.	L1	MS ² (De Young et al., 2021); MIMIC-IV-BHC (Aali et al., 2025); MTS-Dialog (Ben Abacha et al., 2023); ACI-Bench (Yim et al., 2023)
Retrieval-augmented QA	Attribution and faithfulness of answers to retrieved sources and freshness/recency of information. Evaluates how well models retrieve and ground answers in relevant documents (Xiong et al., 2024).	source-contradiction rate; Recall@k, nDCG, MRR; answer correctness;	L0,L1	Mirage RAG benchmark (MMLU-Med, MedMCQA, PubMedQA, BioASQ) (Xiong et al., 2024); HealthSearchQA (part of MultiMedQA) (Singhal et al., 2023)
Evidence-based fact-checking	Reliability of claims and ability to verify or refute medical statements using evidence. Useful for ensuring LLM outputs do not propagate misinformation.	Claim classification accuracy; evidence recall/precision; F1 for true/false/unfounded labels; citation quality.	L0,L1	MedFact (Chen et al., 2025a)
Information extraction / coding	Structured accuracy on entity recognition, relation extraction, coding, and normalization tasks. Establishes ability to extract structured data from unstructured texts (Luo et al., 2022; Li et al., 2016; Uzuner et al., 2011; Zhang et al., 2021).	Mention/cluster F1; relation F1; coding/normalization accuracy; entity-linking accuracy.	L1	BioRED (Luo et al., 2022); BC5CDR (Li et al., 2016); n2c2 2010 (i2b2) (Uzuner et al., 2011); CBLUE (Zhang et al., 2021)

Table 1: Expanded benchmarks by task class (continued).

Task	Usefulness	Metric	Level	Benchmark
Decision support / triage (simulation)	Selective reliability for clinical decision making: calibration at deployable thresholds, risk—coverage trade-offs, harm proxies, and quantitative reasoning. Includes simulation of triage, diagnosis, personalized diabetes management, and medical calculations (Arora et al., 2025; Bedi et al., 2025; Cardei et al., 2025; Mehandru et al., 2025; Khandekar et al., 2024).	ECE; Brier; NLL; risk-coverage curves; contraindication/near- miss rates; accuracy, groundedness, safety, clarity, actionability; MAE for calculations.	L2	HealthBench (Arora et al., 2025); MedHELM (Bedi et al., 2025); DexBench (Cardei et al., 2025); ER-Reason (Mehandru et al., 2025); MedCalc-Bench (Khandekar et al., 2024)
Clinical dialogue	Communication quality and human factors in multi-turn doctor–patient conversations or simulated OSCE interviews. Measures goal completion, uncertainty marking, empathy, and adherence to safety rails; also covers note generation from visit dialogues (Zeng et al., 2020; Yim et al., 2023; Fareez et al., 2022).	Goal completion; uncertainty/hedging tags; rubric-based ratings; guideline-contradiction flags; empathy/communication scores; note-generation quality; ROUGE/BERTScore.	L2,L3	MedDialog (Chinese/English) (Zeng et al., 2020); MTS-Dialog (Ben Abacha et al., 2023); ACI-Bench (Yim et al., 2023); OSCE simulated interview dataset (Fareez et al., 2022)
Multimodal (imaging + text)	Linkage between radiology images and free-text reports or classification labels; evaluates image understanding, report generation, and cross-modal retrieval (Johnson et al., 2019; Wang et al., 2017).	Report correctness; finding detection/linking; classification accuracy; precision/recall/F1; bounding- box/segmentation metrics.	L2,L3	MIMIC-CXR, MIMIC-CXR-JPG (Johnson et al., 2019); NIH ChestX-ray (Wang et al., 2017)
Patient retrieval	Ability to retrieve relevant literature or similar patient summaries to support clinicians. Tests retrieval quality and ranking of semantically similar patients or articles (Zhao et al., 2023).	Recall@k; nDCG; MRR; patient-similarity accuracy; retrieval precision.	L3	PMC-Patients (Zhao et al., 2023)
Molecular / drug discovery	Validity and diversity of generated molecules and optimization of proxy properties for research settings.	Validity/diversity/novelty proxies; property-optimization success; synthetic accessibility; logP and QED scores.	L3	MOSES (Polykovskiy et al., 2018); GuacaMol (Brown et al., 2019)