Robobench: A Comprehensive Evaluation Benchmark for Multimodal Large Language Models as Embodied Brain

Yulin Luo[1,2*]   **Chun-Kai Fan**[1*]   **Menghang Dong**[1*]   **Jiayu Shi**[1,2*]   **Mengdi Zhao**[3*]
Bo-Wen Zhang[4*]   **Cheng Chi**[2]   **Jiaming Liu**[1]   **Gaole Dai**[1]   **Rongyu Zhang**[1]
Ruichuan An[1]   **Kun Wu**[5]   **Zhengping Che**[5]   **Shaoxuan Xie**[2]   **Guocai Yao**[2]   **Zhongxia Zhao**[1,2]
Pengwei Wang[2]   **Guang Liu**[2]   **Zhongyuan Wang**[2]   **Tiejun Huang**[1,2]   **Shanghang Zhang**[1,2✉]

[1] State Key Laboratory of Multimedia Information Processing, School of Computer Science, Peking University
[2] Beijing Academy of Artificial Intelligence   [3] Institute for Brain and Intelligence, Fudan University
[4] University of Science and Technology Beijing   [5] Beijing Innovation Center of Humanoid Robotics

yulin@stu.pku.edu.cn   shanghang@pku.edu.cn
Project page: https://robo-bench.github.io/

Abstract

Building robots that can perceive, reason, and act in dynamic, unstructured environments remains a core challenge. Recent embodied systems often adopt a dual-system paradigm, where System 2 handles high-level reasoning while System 1 executes low-level control. In this work, we refer to System 2 as the *embodied brain*, emphasizing its role as the cognitive core for reasoning and decision-making in manipulation tasks. Given this role, systematic evaluation of the embodied brain is essential for advancing robotic intelligence. Yet existing benchmarks emphasize execution success, or when targeting high-level reasoning, suffer from incomplete dimensions and limited task realism, offering only a partial picture of cognitive capability. To bridge this gap, we introduce RoboBench, a benchmark that systematically evaluates multimodal large language models (MLLMs) as embodied brains. Motivated by the critical roles across the full manipulation pipeline, RoboBench defines five dimensions—instruction comprehension, perception reasoning, generalized planning, affordance prediction, and failure analysis—spanning 14 capabilities, 25 tasks, and 6092 QA pairs. To ensure realism, we curate datasets across diverse embodiments, attribute-rich objects, multi-view scenes, and memory-driven navigation, drawing from large-scale real robotic data and in-house collection. For planning, RoboBench introduces an evaluation framework that uses an MLLM as a world simulator. It moves beyond symbolic matching to evaluate embodied feasibility by simulating whether predicted plans can achieve critical object-state changes under physical and visual constraints, enabling faithful assessment of long-horizon reasoning. Experiments on 14 state-of-the-art MLLMs reveal fundamental limitations: difficulties with implicit instruction comprehension, spatiotemporal reasoning, cross-scenario planning, fine-grained affordance understanding, and execution failure diagnosis. RoboBench provides a comprehensive scaffold to quantify high-level cognition, clarify the role of the embodied brain, and guide the development of next-generation MLLMs toward more robust robotic intelligence.

1   Introduction

Manipulation in dynamic, unstructured environments remains a core challenge for building generalizable embodied intelligence Sun et al. (2024); Clark (1998); Liu et al. (2024c). Such tasks demand not only precise motor control, but also high-level cognition: understanding instructions, perceiving over surroundings, formulating long-horizon plans, inferring affordances, and reflecting on failures Chen et al. (2025c); Zhang et al. (2024a). In this scenario, Multimodal large language models (MLLMs) have shown strong potential for these roles, thanks to their strengths in instruction-following, com-

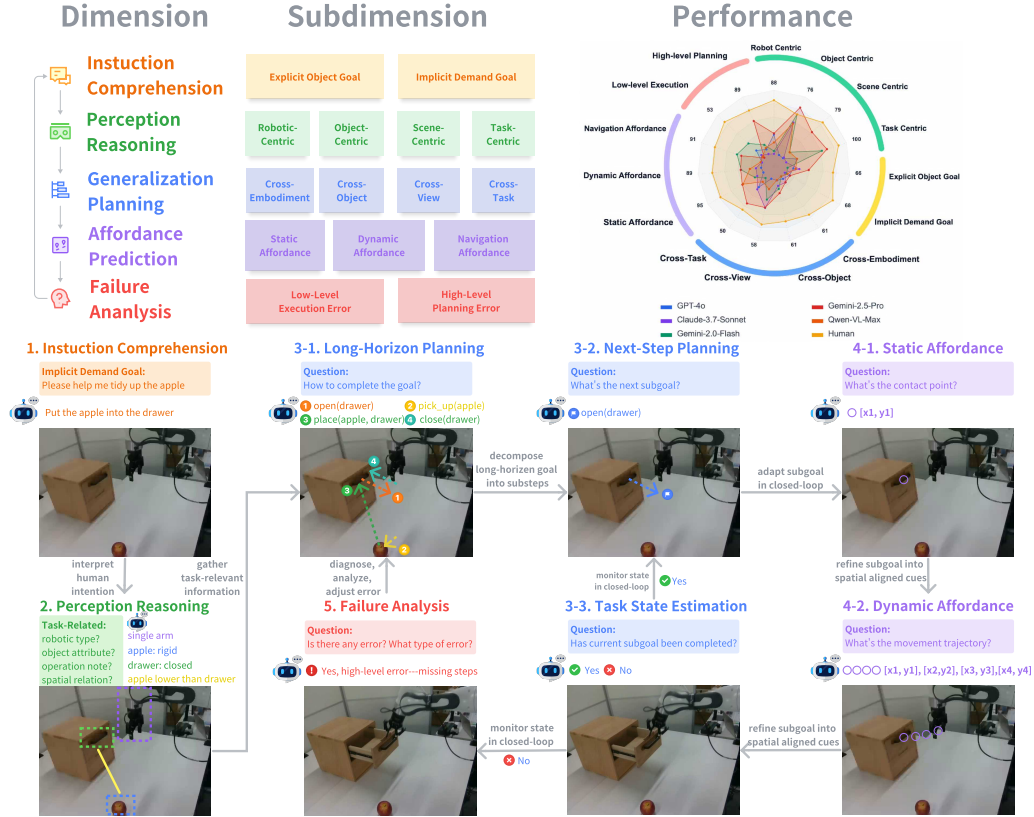---

* Equal contribution   ✉ Corresponding author

1

Figure 1: **Overview of RoboBench** We evaluates MLLMs as embodied brains across 5 dimensions, 14 subdimensions, and 25 tasks, with tasks color-coded by type **(top left)**. These dimensions follow the embodied execution pipeline **(bottom)**—from understanding intent, perceiving the environment, planning and adapting actions, refining subgoals via affordances, diagnosing failures—capturing the core cognitive roles of System 2. Performance comparison **(top right)** highlights significant gaps among state-of-the-art MLLMs, with Gemini-2.5-Pro achieving the best results.

monsense reasoning, and general planning Li et al. (2024a); Wang et al. (2024). To leverage these capabilities, recent embodied systems integrate MLLMs through a dual-system design Driess et al. (2023); Black et al. (2024); Bjorck et al. (2025), where System 2 performs high-level reasoning and System 1 handles low-level control Liu et al. (2025); Chen et al. (2025a). In vision–language–action (VLA) models, MLLMs are fine-tuned as backbones Black et al. (2024), while in multi-agent frameworks, they serve as high-level planners guiding specialized executors Driess et al. (2023); Huang et al. (2023; 2024); Ji et al. (2025). In this work, we refer to System 2 as the *embodied brain*, emphasizing its role as the cognitive core for reasoning and decision-making.

Given this design, systematic evaluation of the embodied brain is essential, yet current benchmarks remain inadequate. Most existing efforts focus narrowly on overall task success, offering little insight into underlying reasoning James et al. (2020); Liu et al. (2023a); Chen et al. (2025d). Even benchmarks that explicitly target embodied cognition exhibit three major shortcomings: (1) fragmented coverage of cognitive abilities, typically isolating perception Majumdar et al. (2024), planning Sermanet et al. (2024); Chen et al. (2024a); Li et al. (2024d), or error reflection Liu et al. (2023b); Duan et al. (2024) rather than assessing them as an integrated whole; (2) limited task realism and complexity, relying heavily on simulations Shridhar et al. (2020a;b); Yang et al. (2025); Cheng et al. (2025); Zhang et al. (2024a), or ignoring practical challenges such as diverse embodiment, object properties, and occlusion; and (3) simplistic planning evaluation, often reduced to multiple-choice Chen et al. (2024a); Azzolini et al. (2025); Team et al. (2025) or text similarity metrics like BLEU Sermanet et al. (2024) and generic LLM scoring Chi et al. (2024), failing to capture embodied priors such as skill–object dependencies, execution-order flexibility, and embodied feasibility.
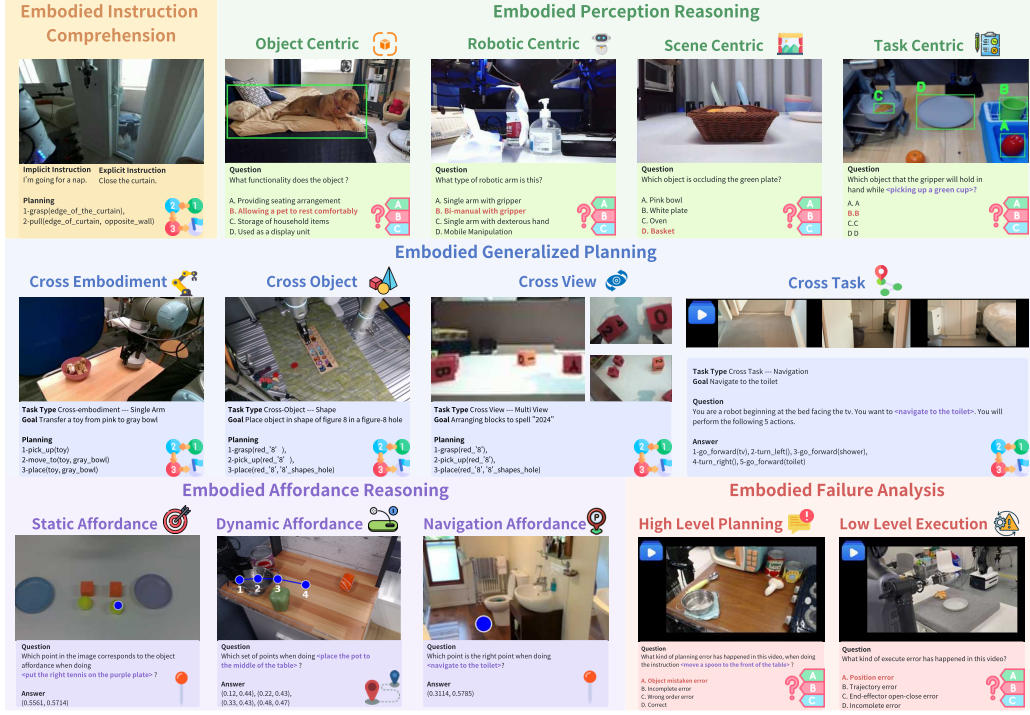
Figure 2: **Examples of RoboBench** Our benchmark covers 5 dimensions, 14 capabilities, and 25 tasks. We design diverse question formats, including multiple-choice, planning, and point prediction.

To overcome these gaps, we introduce **RoboBench**, a benchmark systematically designed to evaluate MLLMs as the cognitive core for robotic manipulation. Specifically, we highlight three key contributions of RoboBench: (1) **Comprehensive evaluation dimensions.** RoboBench defines five dimensions—Instruction Comprehension, Perception Reasoning, Generalized Planning, Affordance Prediction, and Failure Analysis—that together capture the interdependent capabilities required for embodied cognition. These dimensions come from the critical roles of the embodied brain in manipulation by tracing the full execution pipeline, shown in Figure 1. Before acting, it should interpret human intent in context Zhang et al. (2024a). It then perceives the environment to gather task-relevant information. During execution, the brain decomposes long-horizon goals into sequential steps and adapts them in closed-loop. Each subgoal is further refined into spatially aligned cues to guide low-level controllers Liu et al. (2024a); Ji et al. (2025); Team (2025). When failures occur, the system should diagnose, analyze and adjust errors to maintain robustness Duan et al. (2024). (2) **Realistic and diverse tasks.** We construct task settings spanning single-arm, dual-arm, and mobile manipulation; objects with material attributes, physical properties, and world knowledge; multi-view scenes with occlusion and memory-driven navigation Yang et al. (2024). Data combine large-scale real-world robotic datasets Fang et al. (2023); Collaboration et al. (2023); Liu et al. (2024c); Wu et al. (2024); Bu et al. (2025) with curated in-house collection, narrowing the sim-to-real gap. (3) **World-simulation rollout for planning evaluation.** Moving beyond symbolic matching, we propose an novel MLLM-as-world-simulator framework to evaluate long-horizon planning. From the initial scene image, a reference plan list, and a human-annotated directed acyclic graph capturing inter-step dependencies, the evaluator constructs a structured task representation, identifies critical object-state milestones, and rolls out predicted plans step-by-step under visual and physical constraints. This assesses the embodied feasibility—whether the plan can be executed successfully in the real world.

We evaluate 14 state-of-the-art MLLMs across five embodied dimensions and uncover several critical limitations. First, models exhibit severe degradation on implicit instructions, with average scores dropping by 30%, indicating weak grounding of indirect human goals. Second, models perform poorly on robot-view recognition and temporal grounding, revealing fragile spatial–temporal perception and limited embodiment awareness. Third, complex planning remains the primary bottleneck: models struggle with dual-arm coordination and rare-object reasoning, and while multi-view inputs yield

Table 1: Comparison of various benchmarks for embodied AI, evaluating their features across comprehensiveness, reality, complexity, and planning evaluation methodology.

Benchmark	Comprehensiveness					Reality		Complexity		Planning Evaluation method	Size
	Language Instruction	Perception Reasoning	Generalized Planning	Affordance Prediction	Failure Analysis	Real Scene	Real Robot	Cross Embodiment	Multi View		
RoboVQA Sermanet et al. (2024)	✓	×	✓	×	×	✓	✓	✓	×	BLEU	18,248
EgoPlanBench Chen et al. (2024a)	×	✓	✓	×	×	✓	×	×	×	Multi-Choice	3,355
MMRo Li et al. (2024c)	×	✓	✓	×	×	✓	✓	×	×	LLM Score	26,175
EAI Li et al. (2024d)	×	✓	✓	×	×	×	×	×	×	Simulator Task SR	438
OpenEQA Majumdar et al. (2024)	×	✓	✓	×	×	✓	×	×	×	LLM Score	1,636
EgoThink Cheng et al. (2024b)	×	✓	✓	×	×	✓	×	×	×	LLM Score	700
VideoEgoThink Cheng et al. (2024a)	×	✓	✓	×	×	✓	×	×	×	LLM Score	4,993
EmbodiedEval Cheng et al. (2025)	×	✓	✓	×	✓	×	×	×	×	Multi-Choice	327
EmbodiedBench Yang et al. (2025)	×	✓	✓	×	✓	×	×	✓	×	Simulator Task SR	1,128
VLABench Zhang et al. (2024a)	✓	✓	✓	×	×	✓	×	✓	✓	Simulator Task SR	2,427
Robobench (Ours)	✓	✓	✓	✓	✓	✓	✓	✓	✓	MLLM World Simulator	6,092

modest gains (e.g., GPT-4o 33.66→38.51, Claude-3.7-Sonnet 44.51→48.19), they do not close the gap to human performance. Fourth, models show shallow affordance prediction—Gemini-2.5-Pro reaches 65.21, outperforming others but still 17 points behind humans (82.63). Finally, execution-level failure diagnosis remains the hardest dimension, with scores of 10–20 compared to 40–60 on planning error analysis, and even human evaluators achieving only 47.30, underscoring the intrinsic difficulty. Overall, these findings indicate that, despite their promise as embodied brains, current MLLMs exhibit only superficial embodied reasoning and world-modeling capabilities.

In summary, our contributions are as below:

- We introduce RoboBench, a comprehensive benchmark for evaluating MLLMs as embodied brains across 5 key dimensions, spanning 14 capabilities, 25 tasks, and 6092 QA pairs.

- We design realistic datasets and tasks across embodiments, objects, viewpoints, and task settings, providing precise reflections of real-world embodied interaction.

- We propose MLLM-as-world-simulator evaluation framework that faithfully assesses long-horizon planning by simulating whether plans achieve critical object-state milestones.

- We conduct large-scale systematic evaluation of SOTA MLLMs, establishing a leaderboard and offering actionable insights for advancing the embodied brain in robotic AI.

2 RELATED WORKS

Embodied Cognition Benchmark With the rapid development of embodied AI, the number of benchmarks has also grown Sermanet et al. (2024); Chen et al. (2024a); Li et al. (2024c); Cheng et al. (2024a); Yang et al. (2025). However, current benchmarks still have significant limitations in evaluating the capabilities of embodied intelligence. We analyze these limitations from four aspects: comprehensiveness, realism, complexity, and evaluation of planning tasks. Existing benchmarks related to embodied cognition tend to assess individual abilities separately — for example, perception Majumdar et al. (2024); Cheng et al. (2024b), planning Sermanet et al. (2024); Chen et al. (2024a); Li et al. (2024d), and affordance prediction Yuan et al. (2024); Zhao et al. (2025) or failure analysis Liu et al. (2023b); Duan et al. (2024); Zhou et al. (2025b). None of them comprehensively evaluate Multimodal Large Language Models (MLLMs) as integrated embodied brains. In terms of realism, some benchmarks rely on simulators, leading to a simulation-to-reality gap Yang et al. (2025); Cheng et al. (2025); Zhang et al. (2024a). Others use human daily-activity data Chen et al. (2024a); Cheng et al. (2024a), which still differ from real robotic manipulation. Regarding complexity, many benchmarks fail to account for cross-embodiment and cross-viewpoint scenarios Li et al. (2024c); Chen et al. (2024a); Majumdar et al. (2024). For planning evaluation, existing methods also have drawbacks. Text-similarity metrics such as BLEU Sermanet et al. (2024) only measure surface-level linguistic similarity, ignoring the structural and physical feasibility of plans, which can misjudge logically sound but phrased-differently results. Multiple-choice-style evaluations Chen et al. (2024a) deviate from real-world tasks, where solutions are open-ended and not constrained by predefined options. Other scoring approaches that rely on LLM-based judgments Chi et al. (2024) introduce subjectivity and inconsistency, as LLMs may fail to recognize logically valid plans that differ linguistically from reference answers. A comprehensive comparison between our benchmark and prior works is presented in Table 1.
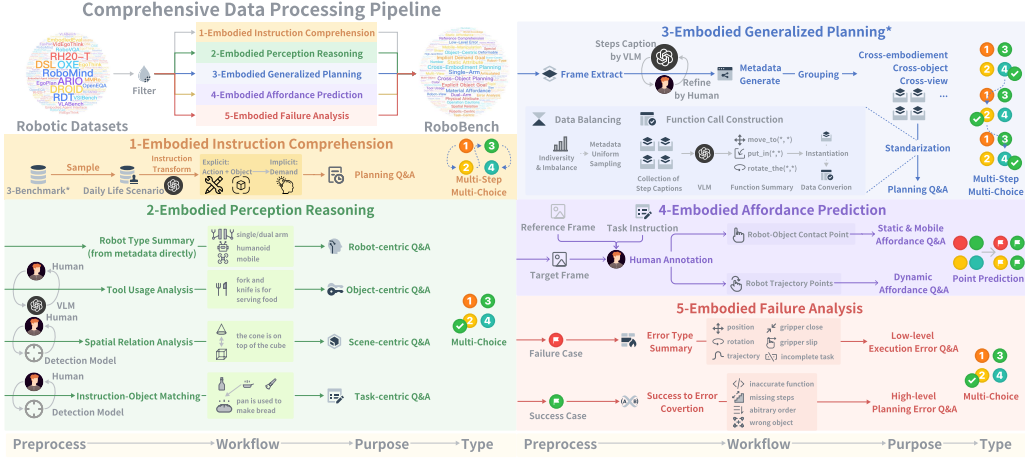
Figure 3: **Dataset Construction Pipeline.** RoboBench customizes data workflows for each dimension: orange for instruction, green for perception, blue for planning, purple for affordance, and red for reflection. Each workflow generally follows three stages—preprocessing, annotation with VLMs, detection models or human experts, and Q&A format generation.

3 ROBOBENCH

3.1 CORE CAPABILITIES

Through a detailed analysis, we identify five critical dimensions of the MLLM-powered embodied brain, each aligned with the task execution pipeline: understanding human intent, perceiving the environment, formulating and adapting plans, refining actions via affordance prediction, and diagnosing failures. RoboBench evaluates these capabilities to uncover bottlenecks in embodied cognition.

Embodied Instruction Comprehension: *Can the embodied brain understand human intention?* Most embodied tasks rely on explicit instructions where actions and targets are clearly specified Sermanet et al. (2024); Yang et al. (2025); Li et al. (2025). In contrast, real-world instructions are often implicit (e.g., "I'm thirsty" instead of "Retrieve a drink") Zhang et al. (2024a). This dimension evaluate whether models can infer both explicit and implicit instructions into actionable plans.

Embodied Perception Reasoning: *Can the embodied brain perceive the environment to gather task-relevant information?* Reliable planning and execution rely on accurate perception Majumdar et al. (2024); Li et al. (2024c); Song et al. (2025); Chen et al. (2025b), RoboBench evaluates this through four aspects: Robotic-centric considers embodiment type and viewpoint. Object-centric examines static and functional attributes. Scene-centric evaluate spatial relations, temporal grounding, and causality analysis. Task-centric assesses identifying instruction-relevant objects.

Embodied Generalized Planning: *Can the embodied brain generalize planning across embodiments, objects, views, and tasks?* Planning begins with decomposing long-horizon goals into subgoals, and continues during execution through predicting the next subtask, monitoring completion, and adapting subsequent steps Sermanet et al. (2024); Yang et al. (2025). We evaluates generalized planning across four aspects: embodiments (single-arm, dual-arm, mobile manipulator, and humanoid robot), objects (material attribute, physical attribute, and commonsense knowledge), views (multi-view integration under occlusion), and tasks (navigation planning using spatial cues from videos Yang et al. (2024)).

Embodied Affordance Prediction: *Can the embodied brain refine subtask plans through spatial affordances?* Beyond high-level planning, each subgoal should be transformed into spatial cues to guide low-level execution Liu et al. (2024a); Li et al. (2024b); Yuan et al. (2024); Zhou et al. (2025a); Team et al. (2025); Lee et al. (2025). Affordance prediction links subgoals with the embodiment, objects, and environment, enabling System 2 to instruct System 1 more effectively Li et al. (2024b). RoboBench evaluates three types: Static affordance identifys contact points (e.g., grasp an apple); Dynamic affordance predicts motion trajectories (e.g., open a drawer); Navigation affordance determines robot base position (e.g., approach a microwave on a distant table).
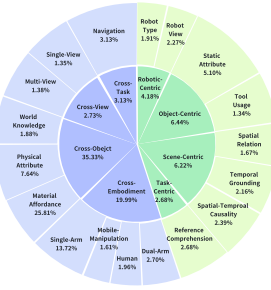
Figure 4: **Dimension Distribution of RoboBench.**

Statistic	Number
Total items	4038
Total questions	6092
Perception, Affordance and Failure questions	
Multiple-choice questions	1875
Instruction and Planning questions	
Q1 questions	1973
Q2 questions	842
Q3 questions	1192
Affordance questions	
Points questions	252
Points set questions	150
Instruction and Planning details-Q1	
Avg. steps	6.74
Unique task instructions	1403
Unique answers	1462
Perception details	
Avg. question lengths	131.3
Avg. choice lengths	13.8
Failure details	
Avg. question lengths	92.9
Avg. choice lengths	17.0

Table 2: **Dataset statistics.**

Embodied Failure Analysis: *Can the embodied brain detect and correct failures?* Open-world manipulation inevitably introduces errors, requiring the brain not only to identify them but also to diagnose causes and suggest corrections Liu et al. (2023b); Duan et al. (2024). RoboBench evaluates low-level execution errors (e.g., position misalignment, trajectory deviations, gripper failures, and incomplete actions), and high-level planning errors (e.g., wrong object, missing steps, incorrect ordering), providing insights for more robust and generalizable execution.

3.2 BENCHMARK CONSTRUCTION

Dataset Collection and Processing Pipeline   To evaluate the five cognitive dimensions, we build datasets from open-source and in-house robotic data, enriched with MLLM- and human-provided annotations, aligning each sub-benchmark with the embodied task pipeline for realism. The examples and construction pipeline are shown in Figures 2 and 3.

Instruction Comprehension: Evaluated through planning tasks, this dimension adopts a paired explicit–implicit design. Explicit instructions are drawn from daily-life scenarios with clearly specified actions and targets, whereas implicit instructions are obtained by rewriting them into demand-based requests using LLMs. This contrast tests a model's ability to interpret human intentions.

Perception Reasoning: Accurate perception is essential for reliable planning and execution. We construct datasets across four aspects: robotic-centric, using real robot data with type and view meta-data for template-based QA; object-centric, combining curated static attributes Gao et al. (2024) with GPT-generated functional properties and distractors; scene-centric, leveraging Gemini-segmented video steps for temporal grounding, together with manual annotations of relative positions and key-point changes for spatial and causal reasoning; and task-centric, with human-labeled bounding boxes linking long-horizon instructions to target objects. All data are standardized into multiple-choice QA.

Generalized Planning: We build a planning pool from high-quality robotic videos Collaboration et al. (2023); Lin et al. (2024a); Liu et al. (2024b;c); Fang et al. (2023); Wu et al. (2024); Mu et al. (2023); Yang et al. (2024), extracting frame sequences as standardized inputs. Gemini Comanici et al. (2025) generates structured annotations—task summaries, step-wise instructions with timestamps, and metadata (objects, actions, scenes, embodiments)—which are refined by human annotators. Each step is then mapped into function templates (e.g., pick_up(object), move_to(object, target)) grouped into manipulation or navigation skill lists Cheng et al. (2024a), guiding models toward structured plan generation. Evaluation covers three types Sermanet et al. (2024): (1) Q1: long-horizon planning, predicting the full action sequence from the first frame and instruction; (2) Q2: next-step planning, forecasting the (n+1)-th step from prior steps; and (3) Q3: task state estimation, deciding whether a given subtask has been completed.

Affordance Prediction: Affordances refine high-level subgoals into spatial cues for low-level execution Liu et al. (2024a). From the planning pool, representative frames are sampled and annotated with three types of affordance: static (contact points), dynamic (motion trajectories), and mobile (base positions). Models predict the corresponding point or trajectory given the task instruction.
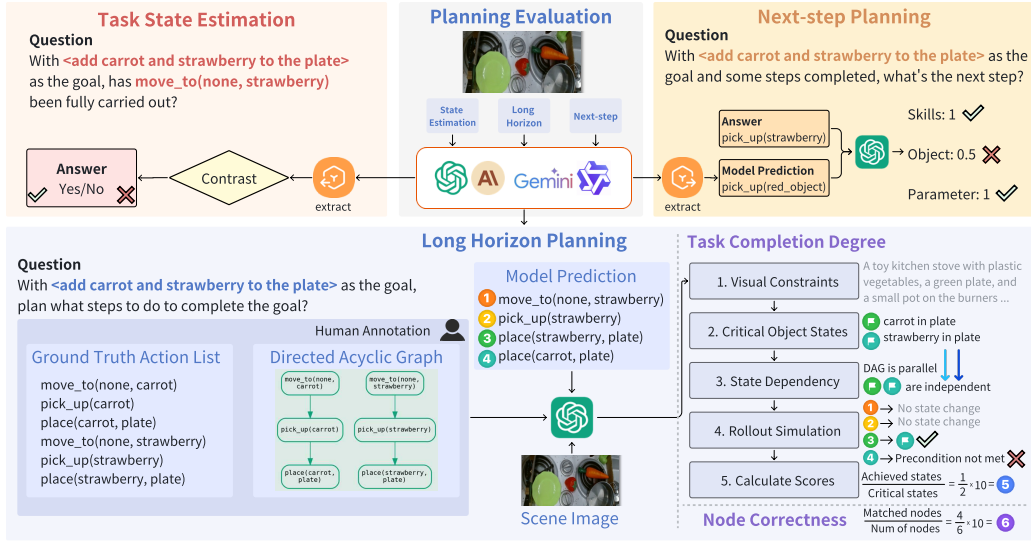
Figure 5: **Planning Evaluation Pipeline.** The planning benchmark includes three task types: long-horizon planning, next-step planning, and task state estimation. They are evaluated respectively with our proposed MLLM-as-world-simulator framework, LLM scoring, and binary accuracy.

Failure Analysis: It evaluates whether models can detect and reason about errors during execution Liu et al. (2023b); Duan et al. (2024). Execution-level failures (e.g., position misalignment, trajectory deviations, gripper errors) are collected from RoboMIND Wu et al. (2024) and labeled by experts, while planning-level failures are synthesized by perturbing correct instructions (wrong object, missing step, wrong order) due to lack of real-world planning failure data.

Quality Control We adopt a two-stage quality control process to ensure benchmark quality: data filtering during construction and post-construction validation. In the construction stage,

We apply both general and task-specific filters. General criteria address image quality and task validity, while sub-benchmarks have tailored rules—such as excluding single-arm data from dual-arm tasks. After data construction, we recruit 20 professionally trained annotators to conduct comprehensive quality inspections, validating linguistic clarity, answerability, and correctness Liu et al. (2024d). We further adopt a majority-vote strategy Liu et al. (2024d): items that all models solve correctly are removed, whereas items all models fail to solve undergo manual review and correction.

Dataset Statistics RoboBench consists of 6092 samples and 4333 unique items, providing a balanced mix of diversity and complexity to evaluate embodied brain capabilities. It covers 5 primary dimensions, 14 secondary meta-tasks, and 25 subtasks, ensuring comprehensive and challenging assessments. A detailed distribution is demonstrated in Figure 4. The detailed question statistics in different dimensions can be found in Table 2.

4 EVALUATION METRICS

To capture the diverse cognitive demands of embodied manipulation, we tailor metrics to each dimension. Perception reasoning and failure analysis are evaluated by multiple-choice accuracy Azzolini et al. (2025); Team et al. (2025). Affordance prediction adopts Euclidean distance for point prediction and RMSE for trajectory prediction Team (2025). To align with other metrics where higher is better, we transform the distance error into a score using $score = 100 \times (1 - d)^{\alpha}$ ($\alpha = 2.5$). Planning is assessed across three tasks (Q1–Q3). Q1 serves as the core, leveraging a MLLM world simulator to evaluate both structural correctness and embodied feasibility.

Q1 – Long-horizon planning. We model a task as a partially ordered set of *atomic actions* parameterized by $\langle$skill, object, args$\rangle$. Inter-step constraints are encoded by a Directed Acyclic Graph (DAG) $G = (V, E)$, where each $v \in V$ is an action node and each edge $(u \to v) \in E$ enforces precedence. The DAG exposes milestones (e.g., `drawer=open`) while allowing valid permutations. Manually annotated G serves as the reference for evaluation; the evaluation pipeline is shown in Figure 5.

We calculate two metrics—*NodeCorrectness* and *TaskCompletion*—both computed via an MLLM. The final score is LongHorizon = (NodeCorrectness + TaskCompletion)/20 $\in [0, 1]$.

Node correctness. Let $V^\star$ denote ground-truth nodes and $\hat{V}$ predicted nodes. An MLLM performs one-to-one matching under exact alignment on {skill, object, parameter}. The calculation is:

$$\mathrm{NodeCorrectness} = \left\lfloor \frac{|V^\star \cap \hat{V}|}{|V^\star|} \times 10 \right\rfloor .$$

Task completion. An MLLM conducts a lightweight world-simulation rollout using: (i) first-frame image I_0 for spatial/physical priors, (ii) reference action list $A^\star$, and (iii) reference DAG G. We define *critical states* as object-level state changes (e.g., $drawer = $ `open`, $apple$ on $plate$). Let $S^\star$ be the set of such milestones and $\hat{S}$ those achieved during rollout. The MLLM proceeds as follows:

1. **Visual constraint analysis:** from I_0, identify objects, spatial relations, and physical constraints to form the initial world W_0 (e.g., $drawer = $ `closed`, $apple$ on $table$).

2. **Critical object-state detection:** parse $A^\star$ and $G = (V, E)$ to extract state-transition predicates from action effects, then aggregate into $S^\star$.

3. **State order & concurrency validation:** use G to enforce precedence and allowable parallelism among states in $S^\star$, and establish causal links.

4. **Rollout simulation:** execute predicted actions step-by-step, checking preconditions against W_t, updating to W_{t+1}, and marking $s \in S^\star$ as *achieved & protected* when it becomes true and remains valid to its last consumer, accumulating $\hat{S} \subseteq S^\star$.

The calculation is as follows:

$$\mathrm{TaskCompletion} = \left\lfloor \frac{|\hat{S}|}{|S^\star|} \times 10 \right\rfloor .$$

Q2 – Next-step planning. Given a observation I_t, the model predicts the next action $\hat{a}_{t+1}$. An MLLM judges against $a_{t+1}^\star$: skill must *exactly* match ($s_{\mathrm{skill}} \in \{0, 1\}$); object and parameters are scored by reasonableness under visual constraints ($s_{\mathrm{obj}}, s_{\mathrm{param}} \in \{0, 0.5, 1\}$). Per-sample score:

$$\mathrm{NextStep} = (s_{\mathrm{skill}} + s_{\mathrm{obj}} + s_{\mathrm{param}})/3 \in [0, 1]$$

Q3 – Task state estimation. Given $I_{t-1} or I_t$, the model decides whether a given subtask $\hat{a}_t$ is achieved: $\hat{y} \in \{0, 1\}$. With ground truth y, we report binary accuracy:

$$\mathrm{StateAcc} = \mathbb{1}\{\hat{y}_i = y_i\}.$$

5 EXPERIMENT

We evaluate several MLLMs on RoboBench, including closed-source MLLMs, open-source multi-image MLLMs, and open-source embodied MLLMs. Detailed descriptions of the models are provided in the supplementary materials. Additionally, we introduce text-only LLM and human as reference.

5.1 OVERALL RESULTS

Gemini-2.5-Pro leads but still lags behind humans: Among all evaluated models, Gemini-2.5-Pro exhibits the most consistent advantages across dimensions, clearly outperforming both other closed- and open-source MLLMs. For instance, in Table 3, it achieves 62.96 in perception reasoning—well

Table 3: **Results on Perception Reasoning(%)**Attr. stand for Attribute; Temp. stand for Temporal; Refer. Comprehen. stand for Reference Comprehension.

Model	Robotic-centric		Object-centric		Scene-centric			Task-centric	Avg
	Robot-type	Robot-view	Static Attr.	Functional Attr.	Spatial Relation	Temp. Grounding	Causality	Refer. Comprehen.	
Basic Reference									
Human Evaluation	80.67	79.08	43.77	83.89	70.91	51.61	91.22	93.22	74.30
GPT-4o-text-only	20.51	13.77	5.18	35.37	25.74	18.32	25.52	22.09	20.81
Closed-Source MLLMs									
GPT-4o-Mini	38.75	18.84	26.43	53.66	30.36	22.65	34.25	39.67	33.08
GPT-4o	**64.96**	39.38	24.92	46.75	42.24	20.61	33.10	41.31	39.16
Claude-3.5-Sonnet	41.31	36.23	29.13	62.60	34.98	21.88	36.09	25.36	35.95
Claude-3.7-Sonnet	40.46	32.37	45.20	71.14	36.63	21.09	40.92	28.02	39.48
Gemini-2.0-Flash	56.69	20.77	49.08	78.46	42.57	21.37	51.72	72.40	49.13
Gemini-2.5-Flash	62.39	39.38	**55.02**	77.24	57.43	33.58	70.34	74.64	58.75
Gemini-2.5-Pro	64.30	41.71	54.83	**82.27**	60.44	49.68	71.73	**78.68**	62.96
Qwen-VL-Plus	28.21	21.74	34.63	58.54	27.72	21.37	31.03	34.36	32.20
Qwen-VL-Max	47.86	**43.48**	39.70	75.20	50.17	27.45	37.93	41.53	45.42
Open-Source Multi-Image MLLMs									
LLaVA-OneVision-0.5B	30.34	23.68	37.08	49.66	27.27	18.42	23.65	19.21	28.66
LLaVA-OneVision-7B	44.83	30.26	33.43	75.84	45.45	23.68	25.68	44.63	40.48
Qwen2.5-VL-7B-Ins	23.93	26.81	37.86	46.34	31.68	22.90	14.48	36.81	30.10
Qwen2.5-VL-72B-Ins	47.72	42.75	41.74	72.95	48.51	27.87	40.32	42.13	45.50
Embodied MLLMs									
RoboBrain-2.0-7B	44.97	24.84	40.43	79.19	48.18	23.48	41.22	53.67	44.50

Table 4: Results on Instruction Comprehension and Generalized Planning Q1. manip. stands for manipulation, Afford. stands for affordance, Knowl. stands for Knowledge.

Model	Instruction Comprehension			Generalized Planning									Cross-Task Planning	Avg
	Explicit	Implicit	Avg	Cross-Embodiment Planning				Cross-Object Planning			Cross-View Planning			
				Single-arm	Dual-arm	Mobile-manip.	Human	Material Afford.	Physical Attr.	World Knowl.	Multi	Single	Navigation Plan	
Basic Reference														
Human Evaluation	59.94	61.13	60.54	72.50	41.93	41.55	62.28	56.70	58.98	49.36	52.82	51.59	45.23	54.50
GPT-4o-text-only	38.80	11.10	24.95	26.70	33.32	43.65	37.86	36.58	22.33	37.68	44.35	38.11	36.90	33.95
Closed-Source MLLMs														
GPT-4o-Mini	41.21	14.95	28.08	27.47	25.21	37.98	31.72	33.75	38.46	42.56	39.11	33.29	34.04	33.31
GPT-4o	45.60	19.04	32.32	28.28	32.65	52.69	35.71	39.93	46.09	41.34	38.51	33.66	39.41	37.74
Claude-3.5-Sonnet	42.11	14.85	28.48	30.18	33.65	50.29	41.05	38.28	40.67	39.63	45.95	40.43	39.77	38.07
Claude-3.7-Sonnet	47.77	14.53	31.15	29.86	38.69	50.39	37.06	38.65	41.86	51.83	48.19	44.51	39.95	41.68
Gemini-2.0-Flash	43.49	16.38	29.93	28.67	33.66	48.27	33.95	40.76	54.27	40.12	46.13	40.73	37.02	38.62
Gemini-2.5-Flash	42.53	17.10	29.82	27.05	40.46	49.91	34.50	39.87	53.37	46.22	39.41	43.29	38.32	39.33
Gemini-2.5-Pro	51.15	19.60	35.37	29.71	37.65	50.96	37.44	39.29	56.50	43.29	47.35	45.12	43.62	41.81
Qwen-VL-Plus	37.77	10.38	24.07	24.68	21.75	32.98	33.91	28.45	33.55	33.78	30.95	28.60	4.39	26.77
Qwen-VL-Max	46.45	16.98	31.71	28.30	35.73	47.79	32.40	40.44	44.33	42.32	41.79	37.68	38.00	37.68
Open-Source Multi-Image MLLMs														
LLaVA-OneVision-0.5B	6.82	1.24	3.61	2.90	4.57	4.77	3.68	4.77	3.47	6.47	4.30	3.62	11.39	4.83
LLaVA-OneVision-7B	18.93	3.48	10.05	11.48	16.23	8.27	5.34	18.51	15.62	8.10	0.00	15.16	24.67	12.15
Qwen2.5-VL-7B-Ins	26.45	4.65	15.55	19.47	12.90	28.75	28.19	22.06	21.63	25.61	11.79	20.12	2.10	18.64
Qwen2.5-VL-72B-Ins	46.81	15.15	30.98	28.20	36.92	49.14	31.31	40.51	44.94	38.90	43.16	40.24	37.47	37.73
Embodied MLLMs														
RoboBrain-2.0-7B	36.93	8.19	22.51	15.46	25.32	32.72	31.81	19.85	30.85	23.24	31.51	23.89	24.53	25.35

above most models that range between 30–45, yet still below the human score of 74.30. In Table 6, For affordance prediction and failure analysis, Gemini-2.5-Pro reaches 65.21 and 45.14, respectively, while most other models remain in the 40–50 and 30–40 ranges. However, humans achieve 82.63 and 63.99 on these dimensions, highlighting the persistent gap between current MLLMs and robust human-level embodied intelligence.

Closed-source models outperform open-source models Closed-source models maintain a clear lead across four of five evaluation dimensions, with performance margins often reaching 10–15%. Only in perception reasoning do open-source models approach comparable results, while gaps in instruction comprehension, planning, affordance reasoning, and failure analysis remain substantial. Within the same model family, performance consistently improves with model size—for example, planning scores in Table 4 increase from 33.31 to 37.74 for GPT-4o-mini vs. GPT-4o, from 38.07 to 41.68 for Claude 3.5 vs. Claude 3.7, and from 39.33 to 41.81 for Gemini 2.5 Flash vs. Gemini 2.5 Pro.

Embodied MLLMs surpass general open-source counterparts The embodied MLLM RoboBrain-2.0-7B achieves higher scores than similarly sized general-purpose open-source models in *perception reasoning*, *generalized planning*, and *affordance prediction*, demonstrating the effectiveness of domain-specific embodied training in enhancing multimodal reasoning and planning performance.

Performance varies across dimensions: RoboBench results reveal large differences across task categories. Perception reasoning achieves relatively higher accuracy, while generalized planning

Table 5: Results on Planning Q2, Q3. Instr. Compre. stand for Instruction Comprehension.

Model	Instr. Compre.		Generalized Planning					
	Explicit Goal		Single Arm		Mater. Afford.		World Knowl.	
	Q2	Q3	Q2	Q3	Q2	Q3	Q2	Q3
Basic Reference								
Human Evaluation	45.28	74.32	27.52	71.35	43.62	71.2	43.89	69.83
GPT-4o-text-only	33.07	56.25	17.22	51.77	29.27	**51.83**	33.33	51.22
Closed-Source MLLMs								
GPT-4o-Mini	39.06	61.25	22.35	52.32	31.56	49.13	**44.27**	36.59
GPT-4o	41.41	62.50	21.64	57.49	31.92	51.87	42.19	46.34
Claude-3.5-Sonnet	33.86	55.00	21.56	57.49	29.60	52.54	39.60	46.34
Claude-3.7-Sonnet	38.54	61.25	22.88	59.13	30.98	57.88	32.29	60.98
Gemini-2.0-Flash	41.93	68.75	**26.27**	63.49	34.98	61.95	42.19	56.10
Gemini-2.5-Flash	40.89	71.25	24.28	**72.21**	33.46	69.39	41.15	60.98
Gemini-2.5-Pro	40.10	**75.00**	25.09	70.30	36.48	**69.67**	33.85	**65.85**
Qwen-VL-Plus	41.15	52.50	20.65	54.22	29.27	50.49	39.58	51.22
Qwen-VL-Max	**44.27**	66.25	23.31	65.94	31.72	62.76	35.94	53.66
Open-Source Multi-Image MLLMs								
LLaVA-OneVision-0.5B	9.42	42.50	15.91	45.23	10.77	49.84	5.92	41.46
LLaVA-OneVision-7B	40.26	58.75	53.10	52.04	44.54	51.95	43.42	60.98
Qwen2.5-VL-7B-Ins	22.92	55.00	12.87	53.13	17.30	49.41	18.75	51.22
Qwen2.5-VL-72B-Ins	42.45	61.25	26.19	60.76	33.56	52.08	32.81	48.78
Embodied MLLMs								
RoboBrain-2.0-7B	36.17	53.75	15.36	54.28	23.62	49.58	28.53	50.26

Table 6: Results on Affordance Prediction and Failure Analysis. Naviga. stand for Navigation.

Model	Affordance Prediction				Failure Analysis		
	Static	Dynamic	Naviga.	Avg	Execution	Planning	Avg
Basic Reference							
Human Evaluation	86.08	80.02	81.85	82.63	47.30	80.67	63.99
GPT-4o-text-only	44.89	40.70	38.19	39.88	25.17	37.93	31.55
Closed-Source MLLMs							
GPT-4o-Mini	50.64	42.88	42.30	46.39	17.66	44.60	31.13
GPT-4o	55.61	49.14	49.91	51.91	22.29	57.01	39.65
Claude-3.5-Sonnet	56.26	54.25	53.84	54.77	16.12	47.52	31.82
Claude-3.7-Sonnet	60.02	52.38	50.07	54.06	18.32	54.24	36.28
Gemini-2.0-Flash	61.65	61.76	66.89	63.37	28.48	59.80	44.14
Gemini-2.5-Flash	61.20	52.04	52.01	54.29	18.54	67.65	43.10
Gemini-2.5-Pro	70.54	62.03	63.96	65.21	15.96	**74.31**	**45.14**
Qwen-VL-Plus	51.74	37.42	47.97	48.18	13.91	40.00	26.96
Qwen-VL-Max	70.01	56.26	50.85	59.43	17.22	57.93	37.58
Open-Source Multi-Image MLLMs							
LLaVA-OneVision-0.5B	20.56	28.56	27.69	24.76	21.19	24.67	22.93
LLaVA-OneVision-7B	23.83	33.61	33.43	30.29	**29.14**	34.00	31.56
Qwen2.5-VL-7B-Ins	49.73	38.03	42.16	43.15	13.91	26.90	20.41
Qwen2.5-VL-72B-Ins	**71.54**	51.94	47.67	56.67	12.59	50.72	31.66
Embodied MLLMs							
RoboBrain-2.0-7B	51.87	54.63	41.61	49.37	7.95	42.00	41.24

remains the most challenging. These disparities reflect the varying cognitive demands of each dimension and the limits of current evaluation, highlighting where future progress is most needed.

5.2 FINE-GRAINED RESULTS ACROSS SUB-TASKS

We further break down results across sub-tasks, highlighting distinct weaknesses in instruction comprehension, perception, planning, and failure analysis.

Implicit vs. Explicit Instruction Understanding: Models perform significantly worse on implicit instructions compared to explicit ones, even when images and ground truth remain the same. In Table 4, for explicit instructions, model scores typically range between 40–50, whereas for implicit ones they drop to 10–20, reflecting an average decrease of about 30%. This indicates a clear weakness in grounding indirect human demands into actionable goals. A likely reason is that current MLLMs fail to fully integrate scene context when inferring task intent. Future improvements require models to jointly consider language, perception, and context to accurately infer and execute human intentions.

Perception Challenges: While models exhibit reasonable performance in object property analysis, they struggle with fundamental robotic perception and spatiotemporal reasoning. Common failures include misidentifying the robot type or viewpoint, and being unable to accurately localize events in time. In Table 3, the success rate on robotic-view recognition is only 43.48, and temporal grounding accuracy reaches merely 49.68. Moreover, causal reasoning across time remains weak, most models achieve only 30–40, with the Gemini series being the notable exception. These results suggest that future models should incorporate stronger embodiment-aware perception modules and explicit spatial–temporal reasoning mechanisms to achieve more robust embodied understanding.

Planning Limitations: Models exhibit limited capacity in complex planning scenarios. (1) Cross-Embodiment Planning: Current models often fail to coordinate dual-arm operations or perform mobile manipulation, resulting in suboptimal spatial search and movement decisions. We attribute this to the nature of existing planning data, which primarily focuses on concrete action execution without explicitly distinguishing between left–right arm assignments or positional movements. (2) Cross-Object Planning: Performance drops sharply when tasks involve uncommon objects, symbolic reasoning, or world knowledge, suggesting weak integration of heterogeneous information. In contrast, results on common, well-seen objects remain relatively acceptable. (3) Cross-View Planning: When the front-view camera is occluded, multi-view inputs effectively enhance model performance. For example, in rearranging letter blocks on a table, the front camera may fail to capture the letters facing upward, leading to incorrect left–right hand assignments. Incorporating wrist-camera views allows the model to make more accurate grasping decisions. This demonstrates the benefit of multi-view reasoning: In Table 4, using multi-image inputs improves planning scores from 33.66 to 38.51
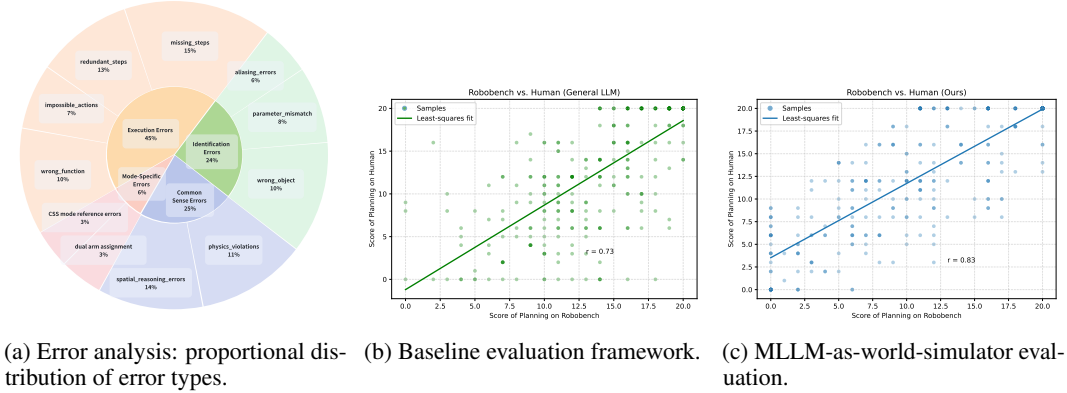
(a) Error analysis: proportional distribution of error types.

(b) Baseline evaluation framework.

(c) MLLM-as-world-simulator evaluation.

Figure 6: Left: planning error analysis; middle/right: alignment with human evaluation. A larger Pearson correlation r indicates stronger alignment.

for GPT-4o and from 44.51 to 48.19 for Claude-3.7-Sonnet, highlighting its potential to strengthen planning robustness in real-world environments.

Failure Analysis: Analyzing execution errors is considerably more challenging than diagnosing planning errors, as reflected by consistently lower scores — execution error scores typically range from 10–20, whereas planning error scores fall between 40–60, shown in Table 6. Execution-level failures often require fine-grained distinctions, such as differentiating positional errors from rotational ones (e.g., the correct location but an incorrect gripper angle), which demand expert-level embodied understanding. Even human performance remains relatively low in such cases — 47.30 compared to 80.67 in planning error — highlighting the intrinsic difficulty of this task. These findings suggest that enhancing fine-grained perceptual capabilities is essential for improving models' performance in execution-level error diagnosis.

5.3 PLANNING ERROR ANALYSIS

We further analyze planning failures observed in our benchmark to understand typical error patterns of embodied models. Errors are categorized into Execution, Identification, Common Sense, and Mode-Specific types, each reflecting a distinct limitation in embodied reasoning. **Execution Errors** account for the majority of failures, primarily caused by missing or redundant actions and incorrect function selection. These indicate insufficient procedural reasoning and weak grounding of task sequences. **Identification Errors** stem from visual aliasing, parameter mismatches, or object misreferencing, revealing the model's difficulty in precise object-action binding. **Common Sense Errors** reflect the lack of physical and spatial reasoning, with models often generating infeasible or directionally inconsistent actions. **Mode-Specific Errors** arise from formatting or symbolic reference violations, showing poor adherence to structured task requirements. Overall, these findings reveal that current models struggle to maintain logical consistency, object grounding, and physical feasibility in long-horizon planning. Improving these aspects is essential for achieving robust embodied intelligence. The proportional distribution of error type is shown in Figure 6a.

5.4 HUMAN EVALUATION OF THE DESIGNED METRIC

To assess the validity of our evaluation metric, we conducted a human study on planning data. We uniformly sampled 438 long-horizon tasks from the Q1 planning inference sets, covering diverse dimensions and model outputs. Expert annotators (with prior training) assigned holistic scores in the range 0–20, referencing the ground-truth planning list and considering multiple factors, such as *skill selection*, *object identification*, and *plan executability*.

We then compare the alignment between human scores and two automated approaches: (i) the *MLLM-as-world-simulator* method, and (ii) an *LLM-based pairwise trajectory scoring* baseline. Our method achieve a higher Pearson correlation coefficient with human judgments, shown in Figure 6b and 6c, indicating stronger consistency and validating the effectiveness of the proposed metric.

6 CONCLUSION

We present RoboBench, a benchmark to systematically evaluate MLLMs as the embodied brain for robotic manipulation. It traces the full execution pipeline across five dimensions—Instruction Comprehension, Perception Reasoning, Generalized Planning, Affordance Prediction, and Failure Analysis—covering 25 tasks with 6092 QA pairs. A core contribution is a MLLM world-simulator framework that evaluates not only symbolic plan fidelity but also embodied feasibility. Experiments expose major gaps in implicit instruction grounding, perception, long-horizon planning, affordance reasoning, and failure diagnosis. RoboBench thus provides a unified scaffold for measuring embodied cognition and guiding the development of more robust, generalizable embodied intelligence.

7 ACKNOWLEDGE

REFERENCES

Ruichuan An, Sihan Yang, Ming Lu, Renrui Zhang, Kai Zeng, Yulin Luo, Jiajun Cao, Hao Liang, Ying Chen, Qi She, et al. Mc-llava: Multi-concept personalized vision-language model. *arXiv preprint arXiv:2411.11706*, 2024.

Ruichuan An, Sihan Yang, Renrui Zhang, Zijun Shen, Ming Lu, Gaole Dai, Hao Liang, Ziyu Guo, Shilin Yan, Yulin Luo, et al. Unictokens: Boosting personalized understanding and generation via unified concept tokens. *arXiv preprint arXiv:2505.14671*, 2025.

Anthropic. Claude 3.5 models, 2024. URL `https://www.anthropic.com/news/claude-3-5-models`.

Anthropic. Claude 3.7 sonnet and claude code, 2025. URL `https://www.anthropic.com/news/claude-3-7-sonnet`.

Alisson Azzolini, Junjie Bai, Hannah Brandon, Jiaxin Cao, Prithvijit Chattopadhyay, Huayu Chen, Jinju Chu, Yin Cui, Jenna Diamond, Yifan Ding, et al. Cosmos-reason1: From physical common sense to embodied reasoning. *arXiv preprint arXiv:2503.15558*, 2025.

Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibo Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, et al. Qwen2. 5-vl technical report. *arXiv preprint arXiv:2502.13923*, 2025.

Johan Bjorck, Fernando Castañeda, Nikita Cherniadev, Xingye Da, Runyu Ding, Linxi Fan, Yu Fang, Dieter Fox, Fengyuan Hu, Spencer Huang, et al. Gr00t n1: An open foundation model for generalist humanoid robots. *arXiv preprint arXiv:2503.14734*, 2025.

Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Lucy Xiaoyang Shi, James Tanner, Quan Vuong, Anna Walling, Haohuan Wang, and Ury Zhilinsky. π_0: A vision-language-action flow model for general robot control, 2024. URL `https://arxiv.org/abs/2410.24164`.

Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess, Adnan Esmail, Michael Robert Equi, Chelsea Finn, Niccolo Fusai, Manuel Y Galliker, et al. π_0.5: a vision-language-action model with open-world generalization. In *9th Annual Conference on Robot Learning*, 2025.

Qingwen Bu, Jisong Cai, Li Chen, Xiuqi Cui, Yan Ding, Siyuan Feng, Shenyuan Gao, Xindong He, Xuan Hu, Xu Huang, et al. Agibot world colosseo: A large-scale manipulation platform for scalable and intelligent embodied systems. *arXiv preprint arXiv:2503.06669*, 2025.

Hao Chen, Jiaming Liu, Chenyang Gu, Zhuoyang Liu, Renrui Zhang, Xiaoqi Li, Xiao He, Yandong Guo, Chi-Wing Fu, Shanghang Zhang, et al. Fast-in-slow: A dual-system foundation model unifying fast manipulation within slow reasoning. *arXiv preprint arXiv:2506.01953*, 2025a.

Kaiyuan Chen, Shuangyu Xie, Zehan Ma, Pannag R Sanketi, and Ken Goldberg. Robo2vlm: Visual question answering from large-scale in-the-wild robot manipulation datasets. *arXiv preprint arXiv:2505.15517*, 2025b.

Shoubin Chen, Zehao Wu, Kai Zhang, Chunyu Li, Baiyang Zhang, Fei Ma, Fei Richard Yu, and Qingquan Li. Exploring embodied multimodal large models: Development, datasets, and future directions, 2025c. URL https://arxiv.org/abs/2502.15336.

Tianxing Chen, Zanxin Chen, Baijun Chen, Zijian Cai, Yibin Liu, Qiwei Liang, Zixuan Li, Xianliang Lin, Yiheng Ge, Zhenyu Gu, et al. Robotwin 2.0: A scalable data generator and benchmark with strong domain randomization for robust bimanual robotic manipulation. *arXiv preprint arXiv:2506.18088*, 2025d.

Yi Chen, Yuying Ge, Yixiao Ge, Mingyu Ding, Bohao Li, Rui Wang, Ruifeng Xu, Ying Shan, and Xihui Liu. Egoplan-bench: Benchmarking multimodal large language models for human-level planning, 2024a. URL https://arxiv.org/abs/2312.06722.

Zhe Chen, Jiannan Wu, Wenhai Wang, Weijie Su, Guo Chen, Sen Xing, Muyan Zhong, Qinglong Zhang, Xizhou Zhu, Lewei Lu, Bin Li, Ping Luo, Tong Lu, Yu Qiao, and Jifeng Dai. Internvl: Scaling up vision foundation models and aligning for generic visual-linguistic tasks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 24185–24198, June 2024b.

Sijie Cheng, Kechen Fang, Yangyang Yu, Sicheng Zhou, Bohao Li, Ye Tian, Tingguang Li, Lei Han, and Yang Liu. Videogothink: Assessing egocentric video understanding capabilities for embodied ai, 2024a. URL https://arxiv.org/abs/2410.11623.

Sijie Cheng, Zhicheng Guo, Jingwen Wu, Kechen Fang, Peng Li, Huaping Liu, and Yang Liu. Egothink: Evaluating first-person perspective thinking capability of vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 14291–14302, June 2024b.

Zhili Cheng, Yuge Tu, Ran Li, Shiqi Dai, Jinyi Hu, Shengding Hu, Jiahao Li, Yang Shi, Tianyu Yu, Weize Chen, Lei Shi, and Maosong Sun. Embodiedeval: Evaluate multimodal llms as embodied agents, 2025. URL https://arxiv.org/abs/2501.11858.

Xiaowei Chi, Hengyuan Zhang, Chun-Kai Fan, Xingqun Qi, Rongyu Zhang, Anthony Chen, Chi min Chan, Wei Xue, Wenhan Luo, Shanghang Zhang, and Yike Guo. Eva: An embodied world model for future video anticipation, 2024. URL https://arxiv.org/abs/2410.15461.

Xiaowei Chi, Peidong Jia, Chun-Kai Fan, Xiaozhu Ju, Weishi Mi, Kevin Zhang, Zhiyuan Qin, Wanxin Tian, Kuangzhi Ge, Hao Li, et al. Wow: Towards a world omniscient world model through embodied interaction. *arXiv preprint arXiv:2509.22642*, 2025.

Andy Clark. *Being there: Putting brain, body, and world together again*. MIT press, 1998.

Open X-Embodiment Collaboration et al. Open X-Embodiment: Robotic learning datasets and RT-X models. https://arxiv.org/abs/2310.08864, 2023.

Gheorghe Comanici, Eric Bieber, Mike Schaekermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blistein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025.

Danny Driess, Fei Xia, Mehdi SM Sajjadi, Corey Lynch, Aakanksha Chowdhery, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, Wenlong Huang, et al. Palm-e: An embodied multimodal language model. In *ICML*, 2023.

Jiafei Duan, Wilbert Pumacay, Nishanth Kumar, Yi Ru Wang, Shulin Tian, Wentao Yuan, Ranjay Krishna, Dieter Fox, Ajay Mandlekar, and Yijie Guo. Aha: A vision-language-model for detecting and reasoning over failures in robotic manipulation. *arXiv preprint arXiv:2410.00371*, 2024.

Hao-Shu Fang, Hongjie Fang, Zhenyu Tang, Jirong Liu, Chenxi Wang, Junbo Wang, Haoyi Zhu, and Cewu Lu. Rh20t: A comprehensive robotic dataset for learning diverse skills in one-shot. *arXiv preprint arXiv:2307.00595*, 2023.

Jensen Gao, Bidipta Sarkar, Fei Xia, Ted Xiao, Jiajun Wu, Brian Ichter, Anirudha Majumdar, and Dorsa Sadigh. Physically grounded vision-language models for robotic manipulation. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 12462–12469. IEEE, 2024.

Yi Han, Cheng Chi, Enshen Zhou, Shanyu Rong, Jingkun An, Pengwei Wang, Zhongyuan Wang, Lu Sheng, and Shanghang Zhang. Tiger: Tool-integrated geometric reasoning in vision-language models for robotics. *arXiv preprint arXiv:2510.07181*, 2025.

Chi-Pin Huang, Yueh-Hua Wu, Min-Hung Chen, Yu-Chiang Frank Wang, and Fu-En Yang. Thinkact: Vision-language-action reasoning via reinforced visual latent planning. *arXiv preprint arXiv:2507.16815*, 2025.

Wenlong Huang, Chen Wang, Ruohan Zhang, Yunzhu Li, Jiajun Wu, and Li Fei-Fei. Voxposer: Composable 3d value maps for robotic manipulation with language models, 2023. URL https://arxiv.org/abs/2307.05973.

Wenlong Huang, Chen Wang, Yunzhu Li, Ruohan Zhang, and Li Fei-Fei. Rekep: Spatio-temporal reasoning of relational keypoint constraints for robotic manipulation, 2024. URL https://arxiv.org/abs/2409.01652.

Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*, 2024.

Stephen James, Zicong Ma, David Rovick Arrojo, and Andrew J Davison. Rlbench: The robot learning benchmark & learning environment. *IEEE Robotics and Automation Letters*, 5(2):3019–3026, 2020.

Yuheng Ji, Huajie Tan, Jiayu Shi, Xiaoshuai Hao, Yuan Zhang, Hengyuan Zhang, Pengwei Wang, Mengdi Zhao, Yao Mu, Pengju An, et al. Robobrain: A unified brain model for robotic manipulation from abstract to concrete. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pp. 1724–1734, 2025.

Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan Foster, Grace Lam, Pannag Sanketi, Quan Vuong, Thomas Kollar, Benjamin Burchfiel, Russ Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang, and Chelsea Finn. Openvla: An open-source vision-language-action model, 2024. URL https://arxiv.org/abs/2406.09246.

Jason Lee, Jiafei Duan, Haoquan Fang, Yuquan Deng, Shuo Liu, Boyang Li, Bohan Fang, Jieyu Zhang, Yi Ru Wang, Sangho Lee, et al. Molmoact: Action reasoning models that can reason in space. *arXiv preprint arXiv:2508.07917*, 2025.

Bo Li, Yuanhan Zhang, Dong Guo, Renrui Zhang, Feng Li, Hao Zhang, Kaichen Zhang, Peiyuan Zhang, Yanwei Li, Ziwei Liu, and Chunyuan Li. Llava-onevision: Easy visual task transfer, 2024a. URL https://arxiv.org/abs/2408.03326.

Jinming Li, Yichen Zhu, Zhibin Tang, Junjie Wen, Minjie Zhu, Xiaoyu Liu, Chengmeng Li, Ran Cheng, Yaxin Peng, and Feifei Feng. Improving vision-language-action models via chain-of-affordance. *arXiv preprint arXiv:2412.20451*, 2024b.

Jinming Li, Yichen Zhu, Zhiyuan Xu, Jindong Gu, Minjie Zhu, Xin Liu, Ning Liu, Yaxin Peng, Feifei Feng, and Jian Tang. Mmro: Are multimodal llms eligible as the brain for in-home robotics?, 2024c. URL https://arxiv.org/abs/2406.19693.

Manling Li, Shiyu Zhao, Qineng Wang, Kangrui Wang, Yu Zhou, Sanjana Srivastava, Cem Gokmen, Tony Lee, Erran Li Li, Ruohan Zhang, Weiyu Liu, Percy S Liang, Fei-Fei Li, Jiayuan Mao, and Jiajun Wu. Embodied agent interface: Benchmarking llms for embodied decision making. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang (eds.), *Advances in Neural Information Processing Systems*, volume 37, pp. 100428–100534. Curran Associates, Inc., 2024d. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/b631da756d1573c24c9ba9c702fde5a9-Paper-Datasets_and_Benchmarks_Track.pdf.

Manling Li, Shiyu Zhao, Qineng Wang, Kangrui Wang, Yu Zhou, Sanjana Srivastava, Cem Gokmen, Tony Lee, Erran Li Li, Ruohan Zhang, et al. Embodied agent interface: Benchmarking llms for embodied decision making. *Advances in Neural Information Processing Systems*, 37:100428–100534, 2025.

Qixiu Li, Yaobo Liang, Zeyu Wang, Lin Luo, Xi Chen, Mozheng Liao, Fangyun Wei, Yu Deng, Sicheng Xu, Yizhong Zhang, et al. Cogact: A foundational vision-language-action model for synergizing cognition and action in robotic manipulation. *arXiv preprint arXiv:2411.19650*, 2024e.

Fanqi Lin, Yingdong Hu, Pingyue Sheng, Chuan Wen, Jiacheng You, and Yang Gao. Data scaling laws in imitation learning for robotic manipulation. *arXiv preprint arXiv:2410.18647*, 2024a.

Ji Lin, Hongxu Yin, Wei Ping, Yao Lu, Pavlo Molchanov, Andrew Tao, Huizi Mao, Jan Kautz, Mohammad Shoeybi, and Song Han. Vila: On pre-training for visual language models, 2023.

Weifeng Lin, Xinyu Wei, Ruichuan An, Peng Gao, Bocheng Zou, Yulin Luo, Siyuan Huang, Shanghang Zhang, and Hongsheng Li. Draw-and-understand: Leveraging visual prompts to enable mllms to comprehend what you want. *arXiv preprint arXiv:2403.20271*, 2024b.

Weifeng Lin, Xinyu Wei, Ruichuan An, Tianhe Ren, Tingwei Chen, Renrui Zhang, Ziyu Guo, Wentao Zhang, Lei Zhang, and Hongsheng Li. Perceive anything: Recognize, explain, caption, and segment anything in images and videos. *arXiv preprint arXiv:2506.05302*, 2025.

Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. Libero: Benchmarking knowledge transfer for lifelong robot learning. *Advances in Neural Information Processing Systems*, 36:44776–44791, 2023a.

Fangchen Liu, Kuan Fang, Pieter Abbeel, and Sergey Levine. Moka: Open-vocabulary robotic manipulation through mark-based visual prompting. In *First Workshop on Vision-Language Models for Navigation and Manipulation at ICRA 2024*, 2024a.

Jiaming Liu, Hao Chen, Pengju An, Zhuoyang Liu, Renrui Zhang, Chenyang Gu, Xiaoqi Li, Ziyu Guo, Sixiang Chen, Mengzhen Liu, et al. Hybridvla: Collaborative diffusion and autoregression in a unified vision-language-action model. *arXiv preprint arXiv:2503.10631*, 2025.

Songming Liu, Lingxuan Wu, Bangguo Li, Hengkai Tan, Huayu Chen, Zhengyi Wang, Ke Xu, Hang Su, and Jun Zhu. Rdt-1b: a diffusion foundation model for bimanual manipulation. *arXiv preprint arXiv:2410.07864*, 2024b.

Yang Liu, Weixing Chen, Yongjie Bai, Xiaodan Liang, Guanbin Li, Wen Gao, and Liang Lin. Aligning cyber space with physical world: A comprehensive survey on embodied ai, 2024c. URL https://arxiv.org/abs/2407.06886.

Yuan Liu, Haodong Duan, Yuanhan Zhang, Bo Li, Songyang Zhang, Wangbo Zhao, Yike Yuan, Jiaqi Wang, Conghui He, Ziwei Liu, Kai Chen, and Dahua Lin. Mmbench: Is your multi-modal model an all-around player?, 2024d. URL https://arxiv.org/abs/2307.06281.

Zeyi Liu, Arpit Bahety, and Shuran Song. Reflect: Summarizing robot experiences for failure explanation and correction. *arXiv preprint arXiv:2306.15724*, 2023b.

Yulin Luo, Ruichuan An, Bocheng Zou, Yiming Tang, Jiaming Liu, and Shanghang Zhang. Llm as dataset analyst: Subpopulation structure discovery with large language model. In *European Conference on Computer Vision*, pp. 235–252. Springer, 2024.

Arjun Majumdar, Anurag Ajay, Xiaohan Zhang, Pranav Putta, Sriram Yenamandra, Mikael Henaff, Sneha Silwal, Paul Mcvay, Oleksandr Maksymets, Sergio Arnaud, Karmesh Yadav, Qiyang Li, Ben Newman, Mohit Sharma, Vincent Berges, Shiqi Zhang, Pulkit Agrawal, Yonatan Bisk, Dhruv Batra, Mrinal Kalakrishnan, Franziska Meier, Chris Paxton, Alexander Sax, and Aravind Rajeswaran. Openeqa: Embodied question answering in the era of foundation models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 16488–16498, June 2024.

Yao Mu, Qinglong Zhang, Mengkang Hu, Wenhai Wang, Mingyu Ding, Jun Jin, Bin Wang, Jifeng Dai, Yu Qiao, and Ping Luo. Embodiedgpt: Vision-language pre-training via embodied chain of thought. *Advances in Neural Information Processing Systems*, 36:25081–25094, 2023.

Pierre Sermanet, Tianli Ding, Jeffrey Zhao, Fei Xia, Debidatta Dwibedi, Keerthana Gopalakrishnan, Christine Chan, Gabriel Dulac-Arnold, Sharath Maddineni, Nikhil J Joshi, et al. Robovqa: Multimodal long-horizon reasoning for robotics. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 645–652. IEEE, 2024.

Mohit Shridhar, Jesse Thomason, Daniel Gordon, Yonatan Bisk, Winson Han, Roozbeh Mottaghi, Luke Zettlemoyer, and Dieter Fox. Alfred: A benchmark for interpreting grounded instructions for everyday tasks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2020a.

Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. Alfworld: Aligning text and embodied environments for interactive learning. *arXiv preprint arXiv:2010.03768*, 2020b.

Chan Hee Song, Valts Blukis, Jonathan Tremblay, Stephen Tyree, Yu Su, and Stan Birchfield. Robospatial: Teaching spatial understanding to 2d and 3d vision-language models for robotics. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pp. 15768–15780, 2025.

Fuchun Sun, Runfa Chen, Tianying Ji, Yu Luo, Huaidong Zhou, and Huaping Liu. A comprehensive survey on embodied intelligence: Advancements, challenges, and future perspectives. *CAAI Artificial Intelligence Research*, 3:9150042, 2024. doi: 10.26599/AIR.2024.9150042. URL https://www.sciopen.com/article/10.26599/AIR.2024.9150042.

Huajie Tan, Xiaoshuai Hao, Cheng Chi, Minglan Lin, Yaoxu Lyu, Mingyu Cao, Dong Liang, Zhuo Chen, Mengsi Lyu, Cheng Peng, et al. Roboos: A hierarchical embodied framework for cross-embodiment and multi-agent collaboration. *arXiv preprint arXiv:2505.03673*, 2025.

BAAI RoboBrain Team. Robobrain 2.0 technical report. *arXiv preprint arXiv:2507.02029*, 2025.

Gemini Robotics Team, Saminda Abeyruwan, Joshua Ainslie, Jean-Baptiste Alayrac, Montserrat Gonzalez Arenas, Travis Armstrong, Ashwin Balakrishna, Robert Baruch, Maria Bauza, Michiel Blokzijl, et al. Gemini robotics: Bringing ai into the physical world. *arXiv preprint arXiv:2503.20020*, 2025.

Octo Model Team, Dibya Ghosh, Homer Walke, Karl Pertsch, Kevin Black, Oier Mees, Sudeep Dasari, Joey Hejna, Tobias Kreiman, Charles Xu, Jianlan Luo, You Liang Tan, Lawrence Yunliang Chen, Pannag Sanketi, Quan Vuong, Ted Xiao, Dorsa Sadigh, Chelsea Finn, and Sergey Levine. Octo: An open-source generalist robot policy, 2024. URL https://arxiv.org/abs/2405.12213.

Wanxin Tian, Shijie Zhang, Kevin Zhang, Xiaowei Chi, Yulin Luo, Junyu Lu, Chunkai Fan, Qiang Zhou, Yiming Zhao, Ning Liu Siyu Lin, et al. Seea-r1: Tree-structured reinforcement fine-tuning for self-evolving embodied agents. *arXiv preprint arXiv:2506.21669*, 2025.

Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Yang Fan, Kai Dang, Mengfei Du, Xuancheng Ren, Rui Men, Dayiheng Liu, Chang Zhou, Jingren Zhou, and Junyang Lin. Qwen2-vl: Enhancing vision-language model's perception of the world at any resolution, 2024. URL https://arxiv.org/abs/2409.12191.

Kun Wu, Chengkai Hou, Jiaming Liu, Zhengping Che, Xiaozhu Ju, Zhuqin Yang, Meng Li, Yinuo Zhao, Zhiyuan Xu, Guang Yang, et al. Robomind: Benchmark on multi-embodiment intelligence normative data for robot manipulation. *arXiv preprint arXiv:2412.13877*, 2024.

Jihan Yang, Shusheng Yang, Anjali W. Gupta, Rilyn Han, Li Fei-Fei, and Saining Xie. Thinking in space: How multimodal large language models see, remember, and recall spaces, 2024. URL https://arxiv.org/abs/2412.14171.

Rui Yang, Hanyang Chen, Junyu Zhang, Mark Zhao, Cheng Qian, Kangrui Wang, Qineng Wang, Teja Venkat Koripella, Marziyeh Movahedi, Manling Li, Heng Ji, Huan Zhang, and Tong Zhang. Embodiedbench: Comprehensive benchmarking multi-modal large language models for vision-driven embodied agents, 2025. URL https://arxiv.org/abs/2502.09560.

Wentao Yuan, Jiafei Duan, Valts Blukis, Wilbert Pumacay, Ranjay Krishna, Adithyavairavan Murali, Arsalan Mousavian, and Dieter Fox. Robopoint: A vision-language model for spatial affordance prediction for robotics. *arXiv preprint arXiv:2406.10721*, 2024.

Shiduo Zhang, Zhe Xu, Peiju Liu, Xiaopeng Yu, Yuan Li, Qinghui Gao, Zhaoye Fei, Zhangyue Yin, Zuxuan Wu, Yu-Gang Jiang, and Xipeng Qiu. Vlabench: A large-scale benchmark for language-conditioned robotics manipulation with long-horizon reasoning tasks, 2024a. URL https://arxiv.org/abs/2412.18194.

Yuan Zhang, Fei Xiao, Tao Huang, Chun-Kai Fan, Hongyuan Dong, Jiawen Li, Jiacong Wang, Kuan Cheng, Shanghang Zhang, and Haoyuan Guo. Unveiling the tapestry of consistency in large vision-language models. *Advances in Neural Information Processing Systems*, 37:118632–118653, 2024b.

Yuan Zhang, Chun-Kai Fan, Tao Huang, Ming Lu, Sicheng Yu, Junwen Pan, Kuan Cheng, Qi She, and Shanghang Zhang. Autov: Learning to retrieve visual prompt for large vision-language models. *arXiv preprint arXiv:2506.16112*, 2025.

Yuanhan Zhang, Jinming Wu, Wei Li, Bo Li, Zejun Ma, Ziwei Liu, and Chunyuan Li. Video instruction tuning with synthetic data, 2024c. URL https://arxiv.org/abs/2410.02713.

Enyu Zhao, Vedant Raval, Hejia Zhang, Jiageng Mao, Zeyu Shangguan, Stefanos Nikolaidis, Yue Wang, and Daniel Seita. Manipbench: Benchmarking vision-language models for low-level robot manipulation. *arXiv preprint arXiv:2505.09698*, 2025.

Enshen Zhou, Jingkun An, Cheng Chi, Yi Han, Shanyu Rong, Chi Zhang, Pengwei Wang, Zhongyuan Wang, Tiejun Huang, Lu Sheng, et al. Roborefer: Towards spatial referring with reasoning in vision-language models for robotics. *arXiv preprint arXiv:2506.04308*, 2025a.

Enshen Zhou, Qi Su, Cheng Chi, Zhizheng Zhang, Zhongyuan Wang, Tiejun Huang, Lu Sheng, and He Wang. Code-as-monitor: Constraint-aware visual programming for reactive and proactive robotic failure detection. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pp. 6919–6929, 2025b.

A APPENDIX OVERVIEW

This appendix provides additional materials that complement the main paper, covering related works, baselines, extended analyses, statistics, case studies, and prompts:

- **More Related Works (Appendix A.1)** Expanded survey of Embodied MLLM, including end-to-end and hierarchical VLAs, and planner-based agent systems.

- **Baseline Models (Appendix B)**: Configurations and comparative results of large-scale foundation models under adjusted evaluation settings.

- **Extended Experiment Analysis (Appendix C)**: In-depth error analyses that complement the main findings.

- **Statistics (Appendix D)**: Dataset source breakdowns, affordance prediction visualizations, and planning complexity distributions.

- **Case Study (Appendix E)**: Qualitative examples contrasting correct and incorrect model reasoning, with error type annotations and visual illustrations.

- **Prompts (Appendix F)**: Representative prompts used during data generation, model inference, and evaluation, enabling transparency and reproducibility.

A.1 MORE RELATED WORKS

Embodied Multimodal Large Language Model Recent advances in multimodal large language models (MLLMs) have delivered strong perception–reasoning–planning capabilities, leading to growing interest in their applications for embodied intelligence Hurst et al. (2024); Li et al. (2024a); Chen et al. (2024b); Lin et al. (2023); Zhang et al. (2024c); Wang et al. (2024); Bai et al. (2025); Chen et al. (2025c); Zhang et al. (2024b); Luo et al. (2024); An et al. (2024; 2025); Zhang et al. (2025); Lin et al. (2024b; 2025); Tian et al. (2025); Chi et al. (2025); Tan et al. (2025). One line of work develops vision–language–action (VLA) systems for language-conditioned manipulation Team et al. (2024); Kim et al. (2024); Black et al. (2024); Collaboration et al. (2023). Early end-to-end VLAs extend MLLMs with action heads (autoregressive Kim et al. (2024), diffusion Li et al. (2024e), flow matching Black et al. (2024), or hybrid Liu et al. (2025)). While promising, these models often struggle with long-horizon tasks and out-of-distribution generalization. To mitigate these issues, hierarchical VLAs let the MLLM emit explicit, semantically rich guidance for the controller—e.g., next-step predictions Black et al. (2025), static affordances Li et al. (2024b), or spatial trajectories Huang et al. (2025); Lee et al. (2025)—thereby externalizing task knowledge and strengthening transfer. This language-mediated structuring of state, goals, and affordances connects MLLM priors to embodied cognition, improving the robustness Black et al. (2025). A complementary direction treats the MLLM as a high-level planner that issues explicit constraints and sketches for low-level executors—such as spatial waypoints and temporal anchors Driess et al. (2023); Huang et al. (2023; 2024); Han et al. (2025) or manipulability regions and trajectories Huang et al. (2024); Ji et al. (2025); Team (2025); Zhou et al. (2025a). These systems demonstrate strong zero- and few-shot generalization to novel objects and tasks while maintaining robotic feasibility. Despite this progress, MLLMs' embodied cognition—the ability to plan, reason, and adapt under physical and task constraints—has not been systematically assessed. In this work, we articulate the core capabilities required of embodied brains and provide a systematic evaluation across them.

B BASELINE MODELS

To ensure fair and comprehensive evaluation, we benchmarked against a diverse set of baseline models spanning both open- and closed-source systems. Table 7 provides a systematic overview of these models, including their accessibility (open vs. closed source), creators, and corresponding API endpoints or model identifiers.

Table 7: The specific information of models in Experimental Results

Model Name	Type	Creator	API Name/Huggingface Model ID
GPT-4o Hurst et al. (2024)	Closed-Source	OpenAI	gpt-4o-2024-11-20
GPT-4o-mini Hurst et al. (2024)	Closed-Source	OpenAI	gpt-4o-mini-2024-07-18
Gemini-2.5-Pro Comanici et al. (2025)	Closed-Source	Google	gemini-2.5-pro
Gemini-2.5-Flash Comanici et al. (2025)	Closed-Source	Google	gemini-2.5-flash
Claude-3.7-Sonnet Anthropic (2025)	Closed-Source	Anthropic	claude-3.7-sonnet-20250219
Claude-3.5-Sonnet Anthropic (2024)	Closed-Source	Anthropic	claude-3.5-sonnet-20241022
Qwen-VL-Max Bai et al. (2025)	Closed-Source	Qwen	Qwen-VL-Max
Qwen-VL-Plus Bai et al. (2025)	Closed-Source	Qwen	Qwen-VL-Plus
Qwen2.5-VL-72B-Ins Bai et al. (2025)	Open-Source	Qwen	Qwen/Qwen2.5-VL-72B-Instruct
Qwen2.5-VL-7B-Ins Bai et al. (2025)	Open-Source	Qwen	Qwen/Qwen2.5-VL-7B-Instruct
LLaVA-OneVision-0.5B Li et al. (2024a)	Open-Source	NTU	lmms-lab/llava-onevision-qwen2-0.5b-ov
LLaVA-OneVision-7B Li et al. (2024a)	Open-Source	NTU	lmms-lab/llava-onevision-qwen2-7b-ov
RoboBrain-7B-2.0 Team (2025)	Open-Source	BAAI	BAAI/RoboBrain2.0-7B

C MORE EXPERIMENT ANALYSIS

C.1 ERROR DEFINITIONS AND ADDITIONAL ANALYSIS

C.1.1 ERROR TAXONOMY

In this section, we define the types of errors and sub-errors encountered in our benchmark. We categorize errors into four main types: *Execution Errors*, *Identification Errors*, *Common Sense Errors*, and *Mode-Specific Errors*. Each error type corresponds to a distinct failure mode of the embodied brain. For instance, execution errors arise when generating or executing action sequences, identification errors occur when associating objects or parameters, common sense errors reflect violations of physical or spatial priors, and mode-specific errors are tied to task-specific formatting requirements. A detailed taxonomy of error types and definitions is provided in Table 8.

C.1.2 ERROR ANALYSIS

Execution Errors. This category reflects failures in generating valid action sequences. We identify four common subtypes: (1) *Missing Steps*: the predicted sequence omits essential actions, such as failing to include a "pick up object" step, which prevents successful task completion. (2) *Impossible Actions*: the model produces functions not defined in the available action space, e.g., outputting an undefined "teleport" action. (3) *Redundant Steps*: the model generates more actions than required, often inserting unnecessary "move" operations that deviate from the ground truth sequence. (4) *Wrong Function*: the model selects the incorrect action for the intended goal, such as using "push" instead of "pull."

Identification Errors. These errors occur when the model fails to correctly associate actions with the appropriate objects or parameters. (1) *Aliasing Errors*: the model confuses visually similar objects, such as misidentifying a crumpled paper ball as popcorn. (2) *Parameter Mismatch*: although the correct function is chosen, its parameters are incorrect, e.g., attempting to grasp an object with an inappropriate force setting. (3) *Wrong Object*: the output involves manipulating an irrelevant object, such as operating on a cup when the instruction specifies a book.

Common Sense Errors. These errors reveal violations of physical feasibility or basic spatial reasoning. (1) *Physics Violations*: the predicted plan contains physically impossible steps, such as folding multiple pieces of clothing simultaneously. (2) *Spatial Reasoning Errors*: the model fails to infer correct spatial or directional relations, e.g., twisting a faucet in the wrong direction.

Mode-Specific Errors. Certain failures are tied to formatting or mode-specific requirements. (1) *CSS Mode Reference Errors*: instead of adhering to symbolic references, the model outputs natural language descriptions of objects, such as predicting "the red cup" rather than the required identifier "Object_3."

Overall, these errors highlight distinct limitations across execution, perception, reasoning, and adherence to task-specific conventions. Addressing them requires better grounding of action functions,

Error Type	Sub-error	Definition
Execution Errors	Missing Steps	The predicted action list omits necessary functions for task completion.
	Impossible Actions	The model outputs action functions that are not defined in the action space.
	Redundant Steps	The model outputs an excessive number of action functions beyond requirements.
	Wrong Function	The model selects an incorrect action function for the intended operation.
Identification Errors	Aliasing Errors	Failing to distinguish between visually similar objects (e.g., recognizing a crumpled paper ball as popcorn).
	Parameter Mismatch	Producing incorrect parameter values for otherwise valid functions.
	Wrong Object	Referring to objects completely unrelated to the intended target.
Common Sense Errors	Physics Violations	Generating physically impossible action sequences (e.g., folding multiple clothes simultaneously).
	Spatial Reasoning Errors	Failing to infer correct spatial operations (e.g., twisting a faucet in the wrong direction).
Mode-Specific Errors	CSS Mode Reference Errors	Referring to objects directly with natural language descriptions instead of the required object indices.
	Dual Arm Assignment	In dual arm tasks, the robot fails to use the arm that is closer to the target object, resulting in an unreasonable operation.

Table 8: Error taxonomy with definitions in RoboBench.

stronger disambiguation of objects, improved spatial reasoning, and stricter alignment with mode constraints.

D   MORE STATISTICS OF PLANNING TASKS

This section provides additional statistics for planning tasks in **RoboBench**, covering detailed distributions of skills, actions, objects, instructions, and task dimensions, as well as dataset source distributions.

D.1   DETAILED TASK STATISTICS

Skill Statistics   Figure 7 presents the distribution of skills. The horizontal axis shows skill names, and the vertical axis indicates counts.

Action Statistics   Figure 8 shows the distribution of action sequence lengths across tasks, while Figure 9 shows action name vs. frequency counts.

D.2   DATASET SOURCE DISTRIBUTION

Figure 10 presents the distribution of source datasets for planning tasks in RoboBench. The figure shows the proportion of tasks contributed by each source.

E   CASE STUDY

To provide further insights, we present illustrative case studies of our evaluation. Figure 11 shows representative error cases where the model fails to reason correctly, while Figure 12 highlights successful examples that demonstrate the model's robustness.
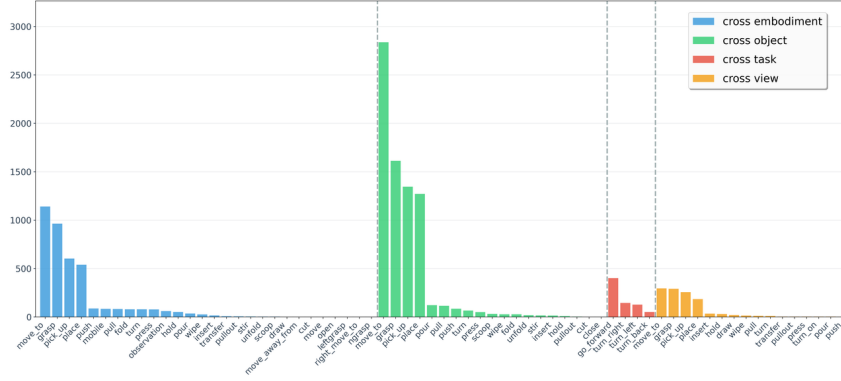
Figure 7: Skill distributions in RoboBench planning tasks. unique skill counts per data entry.
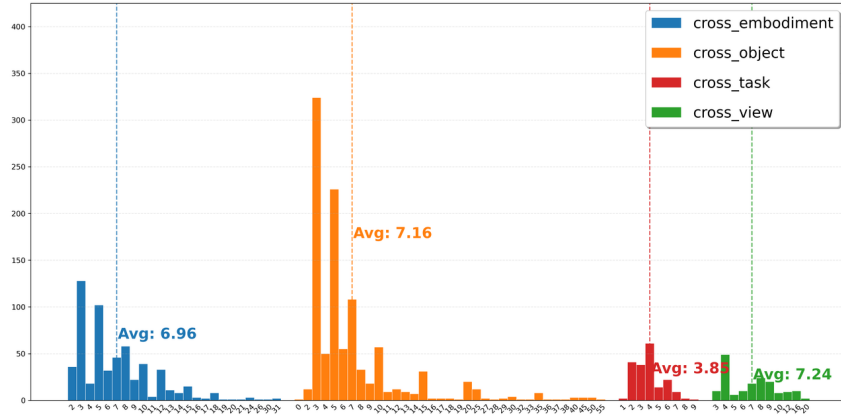

Figure 8: Distribution of action sequence lengths in RoboBench planning tasks.

F PROMPTS

This section presents the prompts employed for both **benchmark data construction** and **model evaluation** in RoboBench. We organize them by task type and processing stage, emphasizing reproducibility and precision. All prompt content placeholders are left empty for subsequent insertion.

F.1 BENCHMARK DATA CONSTRUCTION

F.1.1 PLANNING TASKS

1. Video to Natural Language Task Description   Robot videos are converted into natural language task descriptions, including sequential step descriptions. Three prompts are used to cover different robot types: single-arm, dual-arm, and mobile-manipulator robots, shown in Figure 13, 14 and 15.

2. NL Steps → Predefined Function Sequence   Natural language step descriptions are mapped to a sequence of predefined functions. Two function lists are used: Manipulation (Figure 16) and Navigation (Figure 17), followed by a conversion prompt (Figure 18) referencing these lists.

3. Function Instantiation   Instantiate function arguments with objects extracted from step descriptions. The prompt is presented in Figure 19.

4. Explicit → Implicit Instruction Conversion   To facilitate natural language task formulation for embodied robots, we convert explicit task instructions into implicit forms that imply the required action without directly naming the target object or task (Figure 20).
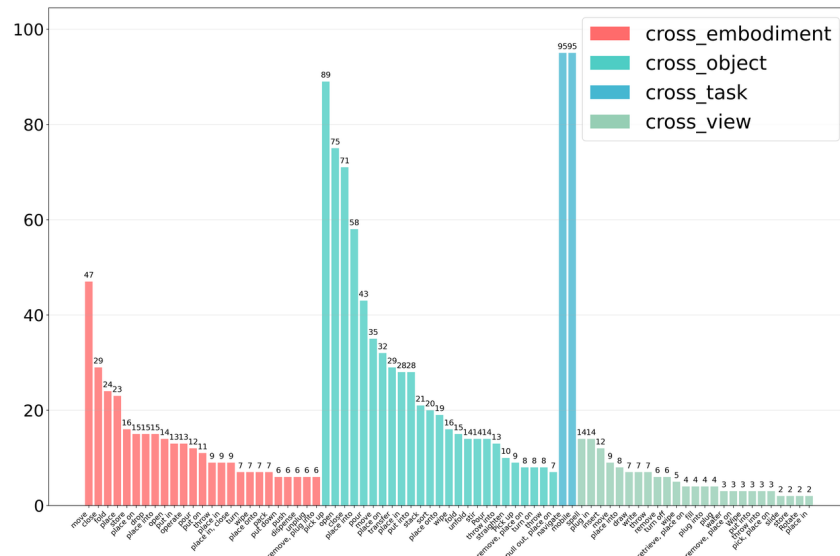
Figure 9: Frequency of each action name in RoboBench planning tasks.

Source Data Distribution
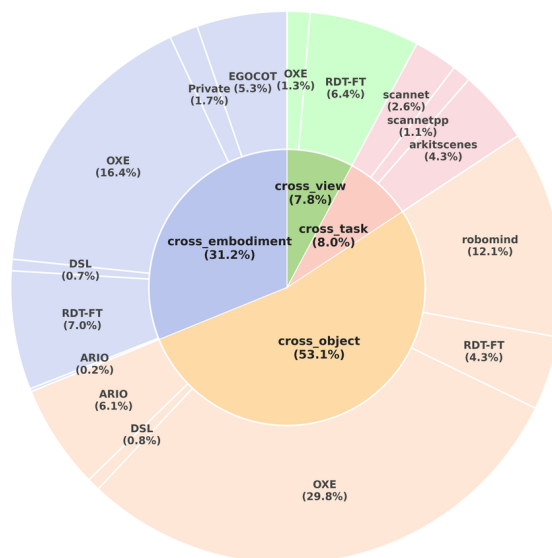
Figure 10: Source distribution of planning tasks in RoboBench dataset.

5. Fine-grained Attribute Extraction We design prompts to extract fine-grained, step-level information from video frames, including **objects** (Figure 21), **actions** (Figure 22), and **scene labels** (Figure 23). These structured annotations provide the foundation for evaluating compositional reasoning and downstream task performance.

F.1.2 PERCEPTION TASKS: FUNCTIONAL ATTRIBUTE SUBTASKS

To evaluate an AI system's understanding of object functionality in images, we generate multi-choice question-answer pairs based on objects highlighted with bounding boxes, focusing on detailed attribute analysis and functional reasoning (Figure 24).
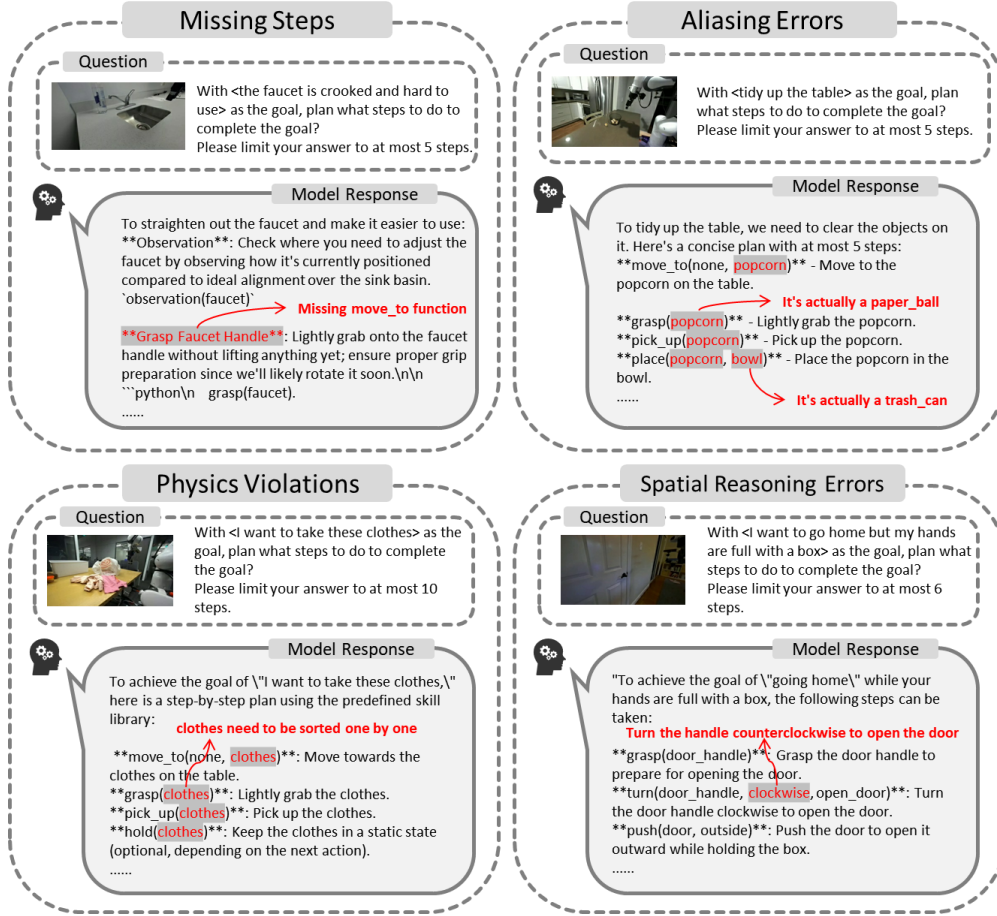
Figure 11: Representative error cases from our evaluation.

F.2 EVALUATION PROMPTS

F.2.1 PLANNING TASKS: Q1 EVALUATION

For Q1-type planning tasks, we provide a structured two-step evaluation: first, extracting the action sequence from model responses (Figure 25), and second, scoring the extracted plans against ground truth data using embodied task reasoning criteria (Figure 26).

F.2.2 PLANNING TASKS: Q2 EVALUATION

For Q2-type planning tasks, we follow a two-stage procedure: first, extracting the next actionable step from the model response in a structured robotic action format (Figure 27), and second, quantitatively scoring the extracted step against the ground truth using strict skill, object, and parameter evaluation criteria (Figure 28).

F.2.3 PLANNING TASKS: Q3 EVALUATION

For Q3-type planning tasks, the evaluation involves converting model responses into a strict binary decision, where only "yes" or "no" is returned without any additional text or formatting, as illustrated in Figure 29.

Figure 12: Successful evaluation examples illustrating robustness.

You will analyze a video (represented by image frames) of a robotic arm performing a specific task, where the task is described as: {desc}. Note that the referenced task summary might not be accurate or complete. Your task is to identify the primary task during the video with the help of the referenced description, summarize the task and rewrite the description, extract the necessary steps to complete it, and specify the frame range for each step. Follow these instructions:

1. **Task Identification**: First, identify the main task the robotic arm is performing. This task could be a clear goal or a series of related activities (e.g., assembling furniture, repairing equipment, preparing food, etc.). Briefly describe the primary task in one sentence.

2. **Step Extraction**: Once the task is identified, extract the key steps required to complete it, ensuring that each step is clearly described and logically ordered. Each step may include:
   - Specific actions (e.g., tightening screws, stirring mixtures, pressing buttons, etc.)
   - Frame window: Specify the start and end frame for each step (from 0 to {maxframeid}, since the video has {cnt} frames).

3. **Frame Range Constraints**:
   - **No Overlapping Frames**: Ensure that the frame ranges of each step do not overlap with each other. Each frame should be assigned to exactly one step.
   - **Full Frame Coverage**: Ensure that all {cnt} frames (from 0 to {maxframeid}) are included in the steps. No frames should be missed or duplicated.
   - **Must start from frame 0**

4. **Notes**:
   - Please annotate as finely as possible and try not to have more than ten frames of the same thing being done (unless it is difficult to distinguish).
   - A step has only one verb (unless two or more actions are strictly performed simultaneously).
   - Please ensure the accuracy of labeling and image matching.
   - For example, if the gripper changes from an open state to a closed state in a sequence of frames, the action is called `pick`; On the contrary, if it changes from a closed state to an open state, the action is called `release`.

5. **Failure Identification**: If the robotic arm attempts an action but does not succeed, clearly indicate this in the step description. For example, if the robotic arm tries to pick up a block but fails, the step description should be something like `Attempt to pick up a block but fails`.

6. **Output Format**: Provide the task description and steps in two parts, formatted as JSON:
   - **Task Summary**: A string summarizing the primary task in the video without mentioning the subjects - the robotic arm.
   - **Steps**: An array where each element represents a step, containing:
     - `step_description`: A concise description of the step with the action being performed in the format of verb phrases without mentioning the subjects - the robotic arm (e.g., "Add syrup in the glass").
     - `start_frame`: The start frame of the step (from 0 to {cnt-1}).
     - `end_frame`: The end frame of the step (from 0 to {cnt-1}).

Task Description: {desc}

Example Output Format 1: (This is an example with a total frame length of 30, and the specific situation depends on the actual frame length.)

```json
{{
  "task_summary": "Assembling an office desk.",
  "steps": [
    {{
      "step_description": "Remove all components and screws from the package.",
      "start_frame": 0,
      "end_frame": 4
    }},
    {{
      "step_description": "Use a screwdriver to attach the legs to the tabletop.",
      "start_frame": 5,
      "end_frame": 14
    }},
    ...
    {{
      "step_description": "Ensure all screws are tight and the desk is stable.",
      "start_frame": 25,
      "end_frame": 29
    }}
  ]
}}
```

Example Output Format 2: ...

Figure 13: Prompt for converting single-arm robot videos into task and step descriptions in natural language.

<div style="border:1px solid #000; padding:10px;">

Prompt: Video → NL Task & Steps (Dual-Arm)

You will analyze a video (represented by image frames) of a dual-arm robotic system performing a specific task, where the task is described as: {desc}. Note that the referenced task summary might not be accurate or complete. Your task is to identify the primary task during the video with the help of the referenced description, summarize the task and rewrite the description, extract the necessary steps to complete it, and specify the frame range for each step. Follow these instructions:

1. **Left-Right Hand Identification**: The video is recorded from a first-person view of the dual-arm robotic system. Your first task is to accurately identify which hand is the left arm ([left]) and which hand is the right arm ([right]). This is crucial as you proceed with the task analysis. Use visual cues such as the relative position of the hands, their orientation, and any distinguishing features to determine which hand is on the left and which is on the right.

2. **Task Identification**: Once the left and right arms are correctly identified, determine the main task the robotic system is performing. This task could be a clear goal or a series of related activities (e.g., assembling furniture, repairing equipment, preparing food, etc.). Briefly describe the primary task in one sentence.

3. **Step Extraction**: After identifying the task and distinguishing the left and right arms, extract the key steps required to complete the task, ensuring that each step is clearly described and logically ordered. **Make sure the following criteria are met**:

   - The **first step** must always start with start_frame equal to 0, and the **last step** must end with end_frame equal to {cnt-1}.
   - Every step must explicitly describe the actions of both the left arm ([left]) and right arm ([right]). If one arm is inactive during a step, mention its inactive state (e.g., [left] holds the object steady while [right] tightens the screw). If both arms are involved in the same action, use [both] to describe their joint activity, but still ensure to detail the left and right arm roles (e.g., [both] lift the object, with [left] supporting the base and [right] holding the top).
   - **Frame window**: Specify the start and end frame for each step (from 0 to {cnt-1}, since the video has {cnt} frames).
   - **Handling Camera Obstruction**: If at any point the view is obstructed by one or both of the robotic arms, add a special marker [block] at the beginning of the step description. Then, based on the context before and after the obstruction, infer the likely actions taking place and provide the most accurate analysis possible. The [block] token should be used only when the camera is blocked, and after the token, describe the inferred task as usual (e.g., [block] [left] holds the object steady, [right] tightens the bolt).

4. **Frame Range Constraints**:

   - **No Overlapping Frames**: Ensure that the frame ranges of each step do not overlap with each other. Each frame should be assigned to exactly one step.
   - **Full Frame Coverage**: Ensure that all {cnt} frames (from 0 to {cnt-1}) are included in the steps. No frames should be missed or duplicated.

5. **Notes**:

   - Please annotate as finely as possible and try not to have more than ten frames of the same thing being done (unless it is difficult to distinguish).
   - A step has only one verb (unless two or more actions are strictly performed simultaneously).
   - Please ensure the accuracy of labeling and image matching.
   - For example, if the gripper changes from an open state to a closed state in a sequence of frames, the action is called pick; On the contrary, if it changes from a closed state to an open state, the action is called release.

6. **Output Format**: Provide the task description and steps in two parts, formatted as JSON:

   - **Task Summary**: A string summarizing the primary task in the video without mentioning the subjects - the robotic arms.
   - **Steps**: An array where each element represents a step, containing:
     - step_description: A concise description of the step, specifying the actions of [left], [right], or [both] arms (e.g., [left] holds the frame, [right] tightens screws, [both] lift the object). Always describe [left] first and [right] second, when applicable.
     - start_frame: The start frame of the step (from 0 to {cnt-1}).
     - end_frame: The end frame of the step (from 0 to {cnt-1}).

Task Description: {desc}

Example Output Format 1: (This is an example with a total frame length of 30, and the specific situation depends on the actual frame length.)

```json
{
  "task_summary": "Assembling an office desk.",
  "steps": [
    {
      "step_description": "[left] removes all components from the package
              while [right] holds the package steady.",
      "start_frame": 0,
      "end_frame": 4
    },
    ...
    {
      "step_description": "[both] ensure all screws are tight and the desk is stable.",
      "start_frame": 29,
      "end_frame": 29
    }
  ]
}
```

**Example Output Format 2: ...

</div>

Figure 14: Prompt for converting dual-arm robot videos into task and step descriptions in natural language.

Prompt: Video → NL Task & Steps (Mobile Manipulator)

You will analyze a **first-person perspective (ego-centric) video** (represented by image frames) of a robot performing a specific task, where the task is described as: "{desc}". Note that the referenced task summary might not be accurate or complete.
Your task is to identify the primary task during the video with the help of the referenced description, summarize the task and rewrite the description, extract the necessary steps to complete it, and specify the frame range for each step. Each step will include a **state** and a corresponding **action description**. Please ensure that the division of atomic tasks is as detailed as possible, the task description is as clear as possible, and there are no vague descriptions.

1. **Task Identification**: First, identify the main task the robot is performing. This task could be a clear goal or a series of related activities (e.g., assembling furniture, repairing equipment, preparing food, etc.). Briefly describe the primary task in one sentence.

2. **Step Extraction**: Once the task is identified, extract the key steps required to complete it, ensuring that each step is clearly described and logically ordered. Each step consists of two parts:

   - **State**: The state of the robot during the step—whether it is moving (`[mobile]`), manipulating (`[manipulation]`), or observing (`[observation]`):

     – `[mobile]`: The robot changes position, but the camera view does not rotate. Describe the movement of the robot (e.g., "Move towards the table").

     – `[manipulation]`: The robot's position and camera view remain static, but the hands are performing actions. Use `[left]`, `[right]`, or `[both]` to describe the actions of the hands (e.g., "`[left]` holds the frame, `[right]` tightens the screws").

     – `[observation]`: This state **must** be triggered when the ego-centric camera's view or orientation changes, indicating the robot is observing or searching for an object. This step is required as a transition whenever the camera view changes, regardless of the robot's other actions. For each `[observation]` state, use the following format:

       * **[goal]**: Clearly specify the object or scene the robot is searching for, using `[find]` to indicate what the robot is looking for or detecting (e.g., "`[find]` the legs of the desk").
       * **[current object or scene]**: Describe what the robot currently sees (e.g., "A table and a chair are visible").
       * **[search result]**: Indicate whether the target object has been found (`yes`, `no`, `part`).

   - **Action Description**: A concise description of what happens during the step, based on the state.

3. **Output Format**: Provide the task description and steps in two parts, formatted as JSON:

   - **Task Summary**: A string summarizing the primary task in the video without mentioning the subjects - the robot's hands or the mobile base.

   - **Steps**: An array where each element represents a step, containing:

     – `state`: The current state of the robot (`[mobile]`, `[manipulation]`, or `[observation]`).

     – `action_description`: A detailed description of the action in the current state (e.g., "`[mobile]` Move towards the table", "`[manipulation]` `[left]` holds the frame, `[right]` tightens the screws", "`[observation]` `[goal]`:`[find]` the legs of the desk. `[current object or scene]`: A table is visible. `[search result]`: no"). Always describe `[left]` first and `[right]` second if the two hands have different functions, when applicable.

     – `start_frame`: The start frame of the step (from 0 to {cnt-1}).

     – `end_frame`: The end frame of the step (from 0 to {cnt-1}).

4. **Frame Range Constraints**:

   - **No Overlapping Frames**: Ensure that the frame ranges of each step do not overlap with each other. Each frame should be assigned to exactly one step.

   - **Full Frame Coverage**: Ensure that all {cnt} frames (from 0 to {cnt-1}) are included in the steps. No frames should be missed or duplicated.

5. **Important Notes**:

   - Please annotate as finely as possible and try not to have one step with more than ten frames doing the same thing (unless it is surely difficult to distinguish).

   - A step has only one verb (unless two or more actions are strictly performed simultaneously).

   - Please ensure the accuracy of labeling and image matching.

   - The video consists of exactly {cnt} frames. Therefore, ensure that the **end frame** for the last step is always {cnt-1}, and no frame should exceed this value.

   - Ensure that the **[observation]** state appears **at the beginning** of the task to search for the initial target object, and as needed thereafter.

   - For **[observation]** states, use the special tokens **[goal]**, **[current object or scene]**, and **[search result]**.

   - Always ensure **[observation]** concludes with `search result: yes` before transitioning to `[mobile]` or `[manipulation]`.

   - **Left and Right Hand Determination**: Use global context, spatial relationships, and object positions to distinguish between `[left]` and `[right]` hands.

Task Description: {desc}

Example Output Format (for a video with 60 frames):

```json
{
    "task_summary": "Assembling an office desk.",
    "steps": [
        {
            "state": "[observation]",
            "action_description": "[goal]: [find] the legs of the desk. [current object or scene]:
                A table is visible. [search result]: no.",
            "start_frame": 0,
            "end_frame": 1
        },
        ...
    ]
}
```

Figure 15: Prompt for converting mobile-manipulator robot videos into task and step descriptions in natural language.

Manipulation Function List

Functions for the actions of a gripper:

- `move_to(object, target_object)`: Move the gripper. First parameter is the object currently held (`none` if empty), second parameter is the target object. Example: `move_to(none, towel)`, `move_to(panda_toy, bowl)`.

- `hold(object)`: Keep an object in static hold. Not applicable if gripper is empty. Example: `hold(cup)`.

Functions for grabbing and releasing:

- `pick_up(object)`: Pick up a graspable object. Example: `pick_up(apple)`.

- `grasp(object)`: Lightly hold an object (pick-upable or not). Example: `grasp(door_handle)`.

- `place(object, target_object)`: Place an object at a location or relative position. Example: `place(apple, table)`, `place(apple, right_of_banana)`.

Functions for using a tool:

- `scoop(tool, contents, container)`: Use a tool to scoop contents. Use `unknown` if contents uncertain. Example: `scoop(spoon, water, bowl)`.

- `pour(container, contents, target_container)`: Pour contents into target. Example: `pour(bowl, water, pot)`.

- `wipe(tool, object, target_object)`: Wipe object using tool on target. Example: `wipe(towel, water, table)`.

- `stir(tool, contents, target_container)`: Stir contents with tool in container. Example: `stir(spoon, soup, pot)`.

- `draw(tool, character, target_object)`: Draw a character using a tool on target. Example: `draw(marker, 'A', whiteboard)`.

- `cut(tool, object, target_object)`: Cut an object with a tool on target. Example: `cut(knife, tomato, chopping_board)`.

Functions for interacting directly:

- `fold(object, target_position)`, `unfold(object, target_position)`: Fold or unfold object to target position.

- `turn(object, direction, state_of_target_object)`: Rotate object to target state. Directions: {clockwise, anticlockwise, up, down, forward, backward, left, right}. Example: `turn(faucet, clockwise, middle_of_sink)`.

- `press(tool, object)`: Press object using tool (`none` if no tool). Example: `press(none, red_button)`.

- `push(object, target_location)`, `pull(object, target_location)`: Push or pull object to target. Example: `push(chair, under_of_table)`, `pull(towel, right_side_of_table)`.

- `insert(object, target_object)`, `pullout(object, target_object)`: Insert or remove object from target. Example: `insert(plug, socket)`, `pullout(plug, socket)`.

Dual-arm specific:

- `transfer(left/right, right/left, object)`: Transfer object between hands. Example: `transfer(left, right, bottle)`.

Mobile-manipulation specific:

- `observation(object)`: Object not visible, needs to be found. Example: `observation(chair)`.

- `mobile(target_object)`: Object visible but distant, move robot to approach. Example: `mobile(table)`.

No Operation:

- `no_ops`: Stay still or maintain current state.

Figure 16: Predefined Manipulation function list used in NL-to-function conversion.

Navigation Function List

Navigation:

- `turn_left()`: Rotate the robot $90°$ to the left in place.

- `turn_right()`: Rotate the robot $90°$ to the right in place.

- `turn_back()`: Rotate the robot $180°$ in place (turn around).

- `go_forward(target_location)`: Move forward in the current facing direction until reaching the specified `target_location`. The parameter corresponds to the semantic name of the target location, e.g., `go_forward(dining_table)` or `go_forward(kitchen_door)`.
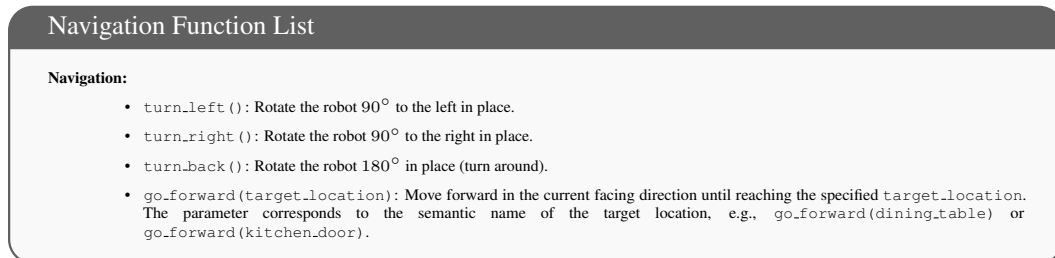
Figure 17: Predefined navigation function list used in natural-language-to-function conversion. The design ensures a compact yet expressive action space that supports systematic evaluation of compositional navigation instructions.

Prompt: NL Steps → Function Sequence

Task Description
You need to process a set of **natural language step descriptions** for a robotic manipulation task and extract **generalized functions**. All generated functions must be selected from a predefined list of functions, ensuring that the function names and parameters match the definitions in the list. If the input action cannot be matched to any function in the list, use `special_action()` and generate a suitable action.

Rules & Requirements

1. **Select from Predefined Functions**:
    - All generated functions must be selected from the following categories:
        - **Object Manipulation**
        - **Movement and Navigation**
        - **Opening and Closing**
        - **Special Operations**
        - **no_ops**
        - `special_action()`
    - If the input action cannot be matched to any function in the list, use `special_action()`.
    - If the input natural language action description indicates a left-handed action and a right-handed action respectively, you need to match a function for each of the left-handed and right-handed actions. Note that the description with `[both]` tag means both hands operate simultaneously and it is also suitable to match function to each of them.

2. **Function Selection Rules**:
    - Select the most appropriate function based on the semantics of the action.
    - Refer to the explanations provided for each function to ensure correct selection.

3. **Output Format**:
    - The output must be in a single-line JSON array format.
    - Each function should be enclosed in quotes and separated by commas.
    - The order of functions should match the input.
    - The number of output functions must be strictly equal to the number of input **natural language step descriptions**.
    - If `special_action()` is used, include the generated description in the JSON.

Predefined Functions List with Explanations
[Predefined Functions List with Explanations]...
Example List
Input:
pick up cup, move to shelf with block, push ball, place book on desk, hover over table, [left] approach glasses case, [right] approach lid, [left] hold case steady, [right] move lid, [both] close laptop lid
Output:

```
["pick_up(object)", "move_to(object, target_object)", "push(object, target_object)",
 "place(object, target_object)", "no_ops",
 "left:move_to(object, target_object), right:move_to(object, target_object)",
 "left:no_ops, right:move_to(object, target_object)",
 "left:close(object), right:close(object)"]
```

Task
Convert the following natural language step descriptions into structured generalized functions:

Figure 18: Prompt for converting natural language steps into a sequence of predefined functions.

Prompt: Function Instantiation

Task Description: Your task is to process a set of natural language action descriptions and generate function calls based on provided function templates. The function calls must accurately reflect the semantics of the input actions and use the correct objects or targets from the descriptions.

Rules & Requirements

- **Input:**
  - A list of natural language action descriptions (e.g., "reach for the juice bottle") and a list of function templates (e.g., "move_to(object)").
  - The number of function templates matches the number of action descriptions, and they correspond in order.
- **Output:**
  - A list of function calls that match the input actions and templates (e.g., "move_to(juice_bottle)").
  - The number of function calls must match the number of input descriptions.
  - If an object or target is composed of multiple words, separate them with underscores (_).
- **Function Call Rules:**
  - Replace the placeholder (e.g., object, target_object, goal, current_scene_or_object) in the function template with the appropriate object or target from the natural language description.
  - Ensure the function call accurately reflects the action described in the input.
- **Output Format:** The output must be a list of function calls in the correct format.

Function Templates contain three types of "observation", "mobile" and "manipulation". The specific templates and explanations are as follows:
[Predefined Functions List with Explanations]

Examples

Input:

- **Natural Language Descriptions:** "[goal]: [find] a cloth [current object or scene]: cabinets, countertop, and floor are visible [search result]: no", "move towards the cloth", "[right] places the oil bottle back on the counter", "[right] reaches for the bowl of shrimp", "[goal]: [find] spatula [current object or scene]: pan with shrimp and spatula are visible [search result]: yes", "reach for the juice bottle", "grasp the juice bottle", "lift the juice bottle", "move the juice bottle towards the blue sticky note", "place the juice bottle on the blue sticky note", "release the juice bottle".

- **Function Templates:** "observation(goal, current_scene_or_object, result)", "mobile(target_object)", "left:no_ops, right:place(object, target_object)", "left:no_ops, right:move_to(object, target_object)", "move_to(object)", "pick_up(object)", "move_to(object, target_object)", "place(object, target_object)".

Output:

```
["observation(a_cloth, ['cabinets','countertop','floor'],no)", "mobile(cloth)",
"left:no_ops, right:place(oil_bottle, counter)",
"left:no_ops, right:move_to(none, bowl_with_shrimp)",
"observation(spatula, ['pan_with_shrimp_and_spatula'], yes)",
"move_to(juice_bottle)", "pick_up(juice_bottle)",
"pick_up(juice_bottle)", "move_to(juice_bottle, blue_sticky_note)",
"place(juice_bottle, blue_sticky_note)",
"place(juice_bottle, None)"]
```

Task Convert the following natural language step descriptions and function templates to the special function calls. Language Descriptions: Function Templates:

Figure 19: Prompt for instantiating predefined functions with specific object references.

Prompt: Explicit → Implicit

Task: Generate implicit task instructions for embodied robots based on a provided explicit task and (if available) an image. The goal is to produce natural, everyday expressions that imply the need for the target task without explicitly mentioning the target object or directly stating the task.

Key Requirements:

1. **Target Entity Focus:**
   - Identify the unique characteristic of the target entity from the explicit task and ensure that the implicit instruction reflects this characteristic.
   - The implicit instruction must have a direct association with the target entity and task, not abstract references.

2. **Everyday Scenarios and Language:**
   - Use casual, real-life scenarios to imply the need for the task.
   - Avoid technical terms or abstract expressions. The instruction should feel like a natural request from a human in daily life.

3. **Image Integration (if provided):**
   - Analyze the provided image directly to extract relevant information about the scene.
   - The image may contain multiple objects, including distracting household items unrelated to the task. Focus on the target object from the explicit task while ensuring the implicit instruction remains relevant and natural.
   - Do not rely on textual descriptions of the image; all visual information must come from analyzing the provided image itself.

4. **No Direct Mentions:**
   - Do not directly mention the target object or the explicit task.
   - The need for the task should be implied through observations, feelings, or everyday needs (e.g., "I feel parched," instead of "pour water into the cup").

5. **Output Format:**
   - Provide 5 implicit task instruction suggestions for each input.
   - The output must be in list format, with each instruction as a separate list item to ensure consistency and ease of post-processing.

Examples:
Target object: pouring liquid from a bottle into a cup **Provided Image:** A table with a bottle, an empty cup, and other unrelated items such as soap, a photo frame, and a gift bag.
Instruction Output:
   - "Looking at the empty cup on the table makes me realize how thirsty I am. Could you help me with that?"
   - …

Target object: watering a plant **Provided Image:** A windowsill with several potted plants, one with drooping leaves.
Instruction Output:
   - "This plant's leaves look a bit droopy today. Could you help bring it back to life?"
   - …

Target object: selecting a Batman toy **Provided Image:** A shelf filled with various superhero toys, including a Batman figure.
Instruction Output:
   - "I'm a big fan of DC series, please help me choose a suitable toy."
   - …

Input Format:
   - **Target object:** {task_summary}
   - **Provided Image:** {image_path}

Output Format:
   - **Instruction Output:**
       - "Instruction 1"
       - "Instruction 2"
       - "Instruction 3"
       - "Instruction 4"
       - "Instruction 5"

Figure 20: Prompt for converting explicit task instructions into implicit form.

Prompt: Step-level Object Extraction

The following is a long-term robot task instruction text. Please analyze the objects being manipulated and other objects mentioned contained in it and generate the output part in a strict format. All words are in the singular form. Object properties (if there are properties other than quantity) are placed in parentheses. You also need to analyze, what kind of object is this manipulation for, choosing from rigid (like bowl), articulated (like microwaves), deformable (like cloth), special (like liquid).
For example,

- task instruction: picking grapes and placing them on a plate; output: [manipulated: grape=rigid][other: plate]
- task instruction: picking up a bowl; output: [manipulated: bowl=rigid][other: none]
- task instruction: pick up a small coca-cola can and place it on a blue paper; output: [manipulated: can(small, coca-cola)=rigid][other: paper(blue)]
- task instruction: pouring dice from a white cup into a pink cup; output: [manipulated: dice=special, cup(white)=rigid][other: cup(pink)]
- task instruction: pouring water from one cup to another; output: [manipulated: water=special, cup=rigid][other: cup]
- task instruction: transferring a usb flash drive between two devices; output: [manipulated: usb flash drive=rigid][other: device]
- task instruction: folding a cloth; output: [manipulated: cloth=deformable][other: none]
- task instruction: opening a laptop; output: [manipulated: laptop=articulated][other: none]
- task instruction: closing a cabinet door; output: [manipulated: door(cabinet)=articulated][other: none]
- task instruction: wiping a table; output: [manipulated: cloth=deformable][other: table]
- task instruction: writing the number 14 on a whiteboard; output: [manipulated: mark pen=rigid][other: whiteboard]

Figure 21: Prompt used to extract **objects** at the step level from video frames.

Prompt: Step-level Action Extraction

Below are task instructions for long-term robotic tasks. Your goal is to extract the primary action from each task instruction.
For instance:

- Task instruction: dragging a strainer backwards, should return: `drag`
- Task instruction: Spelling "THU" with blocks, should return: `spell`
- Task instruction: inserting a three-pronged object into its matching slot, should return: `insert`
- Task instruction: filling a bottle with water from a dispenser, should return: `fill`
- Task instruction: pick apple and place in the drawer, should return: `pick, place in`

These examples illustrate the type of concise action extraction required. The task instruction provides a high-level overview of the task, and your job is to distill it into its core action(s). This prompt is designed for the Gemini model, and it is crucial to focus on identifying the main action verbs that define the task.

Figure 22: Prompt used to extract **actions** at the step level from video frames.

Prompt: Step-level Scene Label Extraction

Please analyze the following image scene and provide a set of scene tags. The specific example of the tags is as follows:

- {'Primary Tag': 'Lab Scene', 'Secondary Tag': 'Robotics Testing Area', 'Tertiary Tag': 'Shape-Sorter Toy', 'Perspective Tag': 'Third-person View'}
- {'Primary Tag': 'Lab Scene', 'Secondary Tag': 'Robot control room', 'Tertiary Tag': 'Workbench', 'Perspective Tag': 'Third-person View'}
- {'Primary Tag': 'Family Scene', 'Secondary Tag': 'Kitchen', 'Tertiary Tag': 'Chopping Board', 'Perspective Tag': 'First-person View'}

Among them, the tertiary tag represents the specific background of these images (for example, all operations in the experiment are carried out on the chopping board); the secondary tag is a collection of tertiary tags, such as stoves, chopping boards, these tertiary scene tags should be clustered into the kitchen, this secondary scene tag, basket frames, and basketball hoops, these tertiary scene tags should be clustered into the basketball court, this secondary scene tag; the primary tag is the clustering of secondary tags, such as kitchens, living rooms, and bedrooms, these secondary scene tags should be clustered into the family scene, this primary tag; basketball courts, volleyball courts, and ping-pong tables, these secondary scene tags should be clustered into the gym, this primary tag; the perspective tag indicates whether the image is shot from a first-person or third-person perspective. The first-person perspective should have the entire scene moving with the operation of the robotic arm, while in the third-person perspective, the overall scene should remain stationary, with only the robotic hand moving.

Note:

1. If you don't know what scene tags to fill in, you can fill in ¡unknown¿ as the answer.
2. The Primary Tag, Secondary Tag, and Tertiary Tag must fill in the background information of the several frames of images I gave you. Please do not fill in other information!!!

Again, your result must use the following format:

- Before providing your tag answers, please explain the reasoning behind the labels you have given.

```
[
    {
        "Primary Tag": "the primary scene tag",
        "Secondary Tag": "the secondary scene tag",
        'Tertiary Tag': "the tertiary scene tag",
        "Perspective Tag": "the perspective type of images"
    }
]
```

Figure 23: Prompt used to extract **scene labels** at the step level from video frames.

Prompt: Functional Attribute QA Generation

As an AI visual assistant, you are tasked with analyzing an image that includes a single marked bounding box (colored green) and generate a multi-choice question-answer pair. The bounding box format is [x, y, width, height], x and y typically represent the coordinates of the top-left corner of the bounding box, while width and height represent the width and height of the box. The image bounding box: {bounding_box}. Initially, you must acquire a comprehensive and intricate understanding of the image, ideally down to each pixel area. Your primary focus should be on deciphering the specific objects or contexts associated with the marks I have placed within the image, as this is of paramount importance. Your analysis will encompass two principal functions:

Bounding Box Description: For each marked bounding box in the image, provide a thorough description using natural language. Detail attributes such as the category, type, color, functionality, and other characteristics of the object, including its location, condition, and any additional pertinent attributes. Envision yourself observing directly and convey your observations as thoroughly and promptly as possible.

– *Template for Role 1*:
'*Bounding Box 1': [Comprehensive Description]*
Continue in this manner for marked bounding box.

Multi-choice Question-Answer Pair Generation: Using the functionality that you describe in Role 1, generate a question about the functionality or use of the object inside the Bounding Box. Provide multiple-choice options and the correct answer. Ensure the question and options are relevant to the object's functionality. The choice in the question must be a little bit difficult to answer, make sure is a good question to distinguish where the model are smart or not.

– *Template for Role 2*:
"*Question*": *What functionality does the object inside the Green Bounding Box have? (A) ... (B) ... (C) ... (D) ...*
"Answer": (...)
–*Example*:
"*Question*": *What functionality does the object inside the Green Bounding Box have? (A) Used to store water (B) Used to store food (C) Used for drinking hot beverages (D) Used to plant flowers*
"*Answer*": *(C)*

Figure 24: Prompt for generating question-answer pairs describing object functionalities.

Prompt: Q1 Evaluation - Action List Extraction

Task: You are given an input dataset containing a robotic manipulation task goal, a previously executed step, and a response describing the remaining steps. Your task is to extract structured action plans in a specific function format.

Instructions:

1. **Extract Key Information:** Identify the task goal from the `prompt` field and assign it to the `"task_summary"` field. Extract action functions from the `previous_step` and `response` fields to construct the sequence of necessary steps in `"plan_step"`.

2. **Strict Action Function Format:** Use only the predefined action functions listed below. Do not modify function names or introduce new ones. Ensure all function names match exactly. Arguments (`object`, `target_object`, `carry_object`, `direction`) should be generalized but faithful to the task.

3. **Maintain Execution Order:** The `"plan_step"` list should follow the correct execution order.

4. **Reasoning Explanation:** Provide a `"reason"` field explaining how the `"task_summary"` and `"plan_step"` were derived, including how you determined the format (single-arm vs dual-arm).

Predefined Action Functions: `""" + PREDEFINED_ACTIONS + """`

Output Format (JSON):
For Single-Arm Tasks:

```
{
  "task_summary": "<task goal>",
  "plan_step": ["<action_function_1>", "<action_function_2>", ...],
  "reason": "<your reasoning>"
}
```

For Dual-Arm Tasks:

```
{
  "task_summary": "<task goal>",
  "plan_step": [
    "<action_function_1>",
    "left:<action_function_2>, right:<action_function_3>",
    ...
  ],
  "reason": "<your reasoning>"
}
```

The data provided is as follows: {data}. Please output your results as required.

Figure 25: Prompt for extracting structured action lists from model outputs for Q1-type questions.

Prompt: Q1 Evaluation - Scoring

Task: Act as a judge of embodied multimodal task planning. Evaluate and score the planning capabilities of multimodal models in embodied intelligence scenarios.

Input & Data Format: You receive:

- **Image**: Current scene with visual constraints (obstacles, spatial layouts, object positions)
- **GT action list**: Ground Truth actions from Gemini video segmentation
- **GT DAG**: Task dependencies (parallel/sequential relationships)
- **Model plan**: Action list output by the evaluated model

Format: {{'GT action list': [...], 'GT dag': [...], 'model plan action list': [...]}}

Mode Detection:

- **Standard Mode**: Textual object names (e.g., `pick_up(apple)`)
- **CSS Mode**: Numeric object IDs (e.g., `pick_up(3)`)

Predefined Action Functions: `""" + PREDEFINED_ACTIONS + """`

Evaluation Criteria (0–10 points each):

1. **Node Correctness**: Match skill-object-parameter nodes between GT and model outputs. *Scoring:* floor((correct/total) * 10). Flexible matching in Standard Mode; strict in CSS Mode.

2. **Task Completion Degree**: Count achieved critical object state changes (turn, place, pick_up, push/pull). Ignore robot motions (e.g., move_to, grasp). *Scoring:* floor((achieved/total) * 10).

Flexible Equivalence (Standard Mode Only): Functional equivalence (e.g., drawer ~ drawer_handle, lamp_switch ~ power_button) is accepted.

Output Format (JSON):

```
{
    "node_correctness": {
        "reason": "... GT vs Model nodes, matching table, formulas ...",
        "result": x
    },
    "task_completion_degree": {
        "reason": "... visual constraints, rollout simulation ...",
        "result": y
    },
    "planning_issue_analysis": {
        "issue_types": [...],
        "detailed_analysis": "..."
    },
    "comprehensive_evaluation": "..."
}
```

Final Guidelines: Enforce mode detection, sequential rollout with DAG constraints, strict vs flexible matching rules, state-tracking for critical object changes, and consistent scoring formulas. Include required artifacts in reasoning: object matching tables, rollout trajectories, and issue categorization.

Figure 26: Prompt for scoring model outputs for Q1-type questions based on ground truth.

Prompt: Q2 Evaluation - Action Extraction

Task: Extract the next step from the given response and represent it in a structured robotic action format.

Action Format: The extracted step must follow the format: `skill(element1, element2, ...)`

Examples:

- `grasp(microwave_handle)`
- `push(microwave_handle, close)`
- `move_to(none, drawer)`

Input: Response: {response}

Output Instruction: Return **only** the extracted step in the correct format. Do not include any explanations, extra text, or additional symbols.

Figure 27: Prompt for extracting structured actions from model outputs for Q2-type questions.

Prompt: Q2 Evaluation - Prompt-based Scoring

Task: You are given two robot action steps (an *extracted* step from a model and a *ground-truth* step). Evaluate their similarity using the prescribed, quantitative criteria below and provide a concise justification for each sub-score.

Input Data Format:

```
{
    "extracted_step": "The step extracted from model response",
    "gt_step": "The ground truth step"
}
```

Evaluation Criteria:

1. **Skill usage accuracy (0 or 1 point).** Consider only the skill/action token (e.g., `grasp`, `push`, `move_to`) in both steps. Award 1 iff the skills are exactly identical (strict match after normalization); otherwise award 0.
   Normalization rule: lowercase, strip extra spaces/underscores for comparison (e.g., `Pick_Up` → `pickup` for matching purposes).

2. **Operation object reasonableness (0, 0.5, or 1 point).** Compare the object argument(s) referenced by the operations:
   - `1.0`: Identical or clearly the same referent (alias/orthographic variants accepted, e.g., `door_handle` vs `door handle`).
   - `0.5`: Related or part/whole/category-level match (e.g., `table` vs `table_leg`, `cup` vs `mug`).
   - `0.0`: Unrelated or incompatible objects.

   Evaluate object similarity using semantic equivalence and task context.

3. **Parameter accuracy (0, 0.5, or 1 point).** Evaluate additional parameters (e.g., target positions, directions, contents). **Important:** if Skill score = 0 or Object score = 0, Parameter score = 0. Otherwise:
   - `1.0`: Parameters fully correct and precise for execution.
   - `0.5`: Parameters partially correct or imprecise but salvageable.
   - `0.0`: Parameters incorrect or irrelevant.

Evaluation Guidelines:

- Enforce **strict** skill matching; be **flexible** on objects in Standard mode (use context), but adhere to strict ID matching in CSS/ID mode if applicable.
- Provide brief, explicit reasons for each sub-score (one-line justification).
- Consider execution precision required by robotic control when judging parameters (e.g., `push(drawer, close)` vs `push(drawer, slight_nudge)`).
- Examples:
  - Extracted: `push(drawer, close)` — GT: `push(drawer_handle, open)` → Skill=1, Object=0.5 (part-whole), Parameter=0.5 (partial).
  - Extracted: `move_to(none, table)` — GT: `move_to(none, table)` → Skill=1, Object=1, Parameter=1.

Output Format (JSON): Return a single JSON object strictly matching the structure below. Each '"reason"' entry should be a concise justification for the assigned score.

```
{
    "skill_usage_accuracy": {"result": x, "reason": "brief explanation of evaluation"},
    "operation_object_reasonableness": {"result": y, "reason": "brief explanation of evaluation"},
    "parameter_accuracy": {"result": z, "reason": "brief explanation of evaluation"}
}
```

The data provided is as follows:
Extracted step: extracted_step Ground truth step: gt_step
Please output your results exactly in the JSON format above.

Figure 28: Prompt for scoring Q2-type questions based on model outputs and prompt instructions.

Prompt: Q3 Evaluation - Yes/No Conversion

Task: From a given response, extract only the binary decision: `"yes"` or `"no"`. No explanations, additional words, or formatting are allowed.

Input Format:

```
Response: {response}
```

Output Requirement:

- Return strictly one token: `"yes"` or `"no"`.
- Output must be lowercase.
- No punctuation, whitespace, or extra text is permitted.

Examples:

- Response: "Yes, that is correct." → Output: `yes`
- Response: "No, it does not work." → Output: `no`

Figure 29: Prompt for converting model outputs into yes/no answers for Q3-type questions.