Reasoning Language Model Inference Serving Unveiled: An Empirical Study

Qi Li[1,2*], **Junpan Wu**[4*], **Xiang Liu**[1*], **Yuxin Wang**[3], **Zeyu Li**[1],
Zhenheng Tang[6], **Yuhan Chen**[1], **Shaohuai Shi**[5], **Xiaowen Chu**[1†]

[1] The Hong Kong University of Science and Technology (Guangzhou)
[2] Shenzhen International Graduate School, Tsinghua University   [3] HKBU
[4] University of Wisconsin-Madison   [5] Harbin Institute of Technology, Shenzhen
[6] The Hong Kong University of Science and Technology
`lqinfdim@163.com`

Abstract

The reasoning large language model (RLLM) has been proven competitive in solving complex reasoning tasks such as mathematics, coding, compared to general LLM. However, the serving performance and behavior of RLLM remains *unexplored*, which may undermine the deployment and utilization of RLLM in real-world scenario. To close this gap, in this paper, we conduct a comprehensive study of RLLM service. We first perform a pilot study on comparing the serving performance between RLLM and traditional LLM and reveal that there are several distinct differences regarding serving behavior: (1) *significant memory usage and fluctuations*; (2) *straggler requests*; (3) *adaptive running time*; (4) *domain preference*. Then we further investigate whether existing inference optimization techniques are valid for RLLM. Our main takeaways are that model quantization methods and speculative decoding can improve service system efficiency with small compromise to RLLM accuracy, while prefix caching, KV cache quantization may even degrade accuracy or serving performance for small RLLM. Lastly, we conduct evaluation under real world workload modeled by Gamma distribution to verify our findings. Empirical results for real world workload evaluation across different dataset are *aligned* with our main findings regarding RLLM serving. We hope our work can provide the research community and industry with insights to advance RLLM inference serving. The reproduction details can be found in §F.

1 Introduction

Large language models (LLM) such as GPT [1], Claude [2, 3], Gemini [4], Llama [5] have emeraged as powerful knowledge bases through pre-training. These models, trained on vast Internet-crawled corpora such as C4 [6], PILE [7] and guided by scaling law [8, 9], have accumulated large-scale knowledge, and exhibited remarkable performance on various knowledge extensive tasks. Despite these advancements, LLMs are criticized for their unsatisfactory capabilities on complex reasoning tasks, e.g., challenging mathematics, and programming tasks.

Recently, reasoning large language models (RLLM) like OpenAI o1 [10], DeepSeek R1 [11], Qwen-3 [12] have sparked a growing body of research into *test time scaling* [13, 14] via *long chain-of-thought reasoning* [15], significantly improving their mathematical reasoning, coding tasks and knowledge reasoning capabilities, e.g., even a 1.5B open source RLLM can surpass giant cutting-edge LLMs

---

[*]Equal Contribution.
[†]Corresponding Author.

like GPT-4o on math tasks [11]. Such achievements make it possible to deploy a small to medium RLLM as a powerful assistant to light the burden of workload for the staff of small entities or even for person, democratizing the use of cutting-edge RLLMs. Hence, it is desirable for small entity with *limited GPU resources* to *efficiently* deploy RLLM with inference engine privately for internal use.

Nevertheless, current LLM serving engine, e.g. vLLM [16], LMDeploy [17], Tensor-RT [18], are initially designed for traditional LLM , other than for RLLM. Though optimization techniques for LLM serving (§2) have been extensively studied, it remains largely *unexplored* whether RLLM exhibits distinct serving characteristics from LLM. If so, directly applying existing LLM serving techniques to RLLM may leave sub-optimal serving performance. Thus, it is natural to ask the following critical research question:

Is there any distinct difference in serving behaviors between LLM and RLLM?

To answer the above question, we perform systematic study of efficient RLLM serving. We first establish the **ASU assessment framework** (§3.2) for assessing RLLM serving. To justify whether there exists a distinct difference in serving behavior between RLLM and LLM, we design a benchmark suite named *ASU-Perf* and conduct a pilot investigation with it on different scale LLM and RLLM (§4). We found that when requests arrive in batches, the serving behavior of RLLMs *differ significantly* from LLMs, and the *main findings* can be primarily summarized in the following aspects: (1) RLLM exhibits significant KV Cache fluctuations and usage; (2) long tail distribution of requests running time caused by slow requests; (3) RLLM solves different difficulty level problems with adaptive running time; (4) RLLM excels LLM on math reasoning while on-par on knowledge intensive tasks.

To understanding RLLM serving further, we first conduct extensive evaluations with various optimization techniques across diverse benchmarks (§5). We find that the model quantization and speculative decoding integrated in serving engine can improve serving efficiency and performance with only small compromising on accuracy of RLLM. However, prefix caching, and KV cache quantization do not always improve serving efficiency. They degrade the accuracy or serving performance for small RLLM, e.g., 7B model. Lastly, we conduct evaluation (§6) under real world workload modeled by Gamma distribution to verify our findings with different scale language models across different domain. Empirical results of real world workload evaluation indicate that the serving behaviors of RLLM are distinct from the LLM and are *aligned* with our main findings regarding RLLM serving.

We hope our work can provide the research community and industry with insightful perspectives to help advance studies in efficient RLLM serving. To the best of our knowledge, we are the *first* to dissect the RLLM serving performance. The main contributions of this paper are the following.

- Conceptually, we propose *ASU*, a framework to assess RLLM serving, which considers accuracy of response, RLLM service-provider side metric, and user side performance metrics together (§3).
- Technically, we introduce *ASU-Perf*, a benchmarking suite for evaluating RLLM serving (§3).
- Empirically, we reveal key differences of serving behaviors between RLLM and LLM: *Significant Memory Fluctuations and Usage*, *Straggler Requests*, and *Adaptive Running Time* (§4).
- We conduct extensive experiments on some RLLM serving optimization techniques (§5).
- We empirically validate our findings in real-world workload and verify their generalization (§6).

2   Preliminaries

In this section, we provide preliminaries of RLLM, LLM serving and its metric. For comprehensive introduction of LLM serving optimization and recent advancement, please refer to Appendix E.

RLLM and LLM. LLMs have demonstrated remarkable capabilities across various natural language processing tasks. However, standard LLMs often encounter difficulties when faced with complex problems that require multi-step reasoning, planning, and deeper cognitive processes, sometimes referred as "System-2 tasks" [19]. To address these limitations, RLLMs have emerged, specifically engineered to enhance these deliberative reasoning abilities. A key technique employed by RLLMs is the "long Chain of Thought" (long CoT) prompting strategy [20]. This approach encourages the model to generate extended, explicit step-by-step reasoning pathways, breaking down complex problems into more manageable parts. Unlike standard LLMs that might provide more direct or less detailed answers, RLLMs utilizing long CoT can better navigate the intricacies of tasks, leading to

more accurate and justifiable solutions by methodically thinking through the problem. This distinction allows RLLMs to tackle challenges in domains like advanced mathematics, intricate logical puzzles, and long-horizon planning more effectively than their conventional counterparts.

LLM Serving. To exploit LLM in real-world scenarios, current practice generally delegates the inference procedure as an individual serving service. The design goal of such serving systems is to accommodate inference output to client users with *low latency* and *high throughput* and full use of GPU memories. Unlike the encoder-based language model [21] like BERT [22], LLM first processes input prompts with intensive computation at the *prefill stage* and then generates output tokens one by one within each iteration at *decoding stage*, which limited by the memory capacity of the hardware. Traditional serving systems process prompts batch by batch, resulting in ineffective memory utilization. Orca [23] introduces *continuous batching* schedule at granularity of each token generation iteration to improve throughput of serving system. To handle as much input requests, the memory space for serving system should be efficient yet elaborated managed. Since decoding phase needs to re-use KV values of their prompt tokens which are stored in GPU, vLLM [16], a high performance serving engine, introduces PagedAttention with paged memory fragmentation and sharing mechanism , which alleviates memory fragmentation and enables allocation in demand. Considering the prefill is compute-intensive task, while the decode is memory-intensive task, for further improvement, DistServe [24] disaggregates the prefill and decode phase by assign computation of these two stages to different GPUs, which co-optimizes the resource allocation and parallelism tailored for each phase.

Serving Performance Metrics. To measure the performance of serving system, there are multiple metrics can be chosen: (1) Time to first token (TTFT) is the time it takes to process the prompt until generate the first token. It measures how long a user must wait before seeing the model's output; (2) End-to-end request latency (E2E latency) indicates the time it takes from submitting the first token of a request to receiving the last token of response, including the time for queueing and batching and network latencies in real-world scenario; (3) Time between tokens (TBT, a.k.a Intertoken latency, ITL) is the average time between the generation of consecutive tokens in a sequence; (4) Tokens per second (TPS) of system represents the mean of total output tokens number per second , accounting for all the requests happening simultaneously; (5) Requests per second (RPS) is the average number of requests that can be successfully completed by the system in a 1-second period. For More details of LLM benchmarking metrics, please refer to §E.2 and related resource [25, 26].

3 Experimental Settings

In this section, we present experimental setups (§3.1) and the ASU assessment framework (§3.2).

3.1 Setups

Here, we list necessary experimental setups. For implementation details, please refer to Appendix G.

Language Models. We employ 4 different scale models to assess their serving performance and serving behavior. General LLM : Qwen-2.5-Math 7B [27], Qwen-2.5-14B , Qwen-2.5-32B [28], and meta-llama/Llama-3.3-70B-Instruct [5] and their long-cot tuned counterparts RLLM: DeepSeek-R1-Distill-Qwen-7B, DeepSeek-R1-Distill-Qwen-14B, DeepSeek-R1-Distill-Qwen-32B , and DeepSeek-R1-Distill-Llama-70B for fair comparison.

Evaluating Datasets. We adopt four different widely used datasets to evaluate the performance of RLLM. Since RLLMs are particularly trained for system-2 reasoning tasks [15], we mainly perform benchmarking with mathematical problems. We adopt three different difficulty level math reasoning datasets: GSM8K [29] as easy level, MATH-500 [30, 31] as medium level, AIME-2024 [32] as the hardest level. To further distinguish are there any differences of serving performance and behaviors for RLLM in reasoning math problem or knowledge-based problem , we also used GPQA [33] dataset for knowledge reasoning. More details of these datasets are introduced in §G.1.

LLM Inference Engine. We employ 2 most adopted open source LLM inference engines, vLLM and SGLang [34] in evaluation. We use OpenAI compatible API of these engines.

Evaluation Suite. We employ *ASU-Perf*, an benchmark suite proposed *by us* for evaluating LLM and RLLM serving performance with different inference engine. We leverage it in all of evaluation.

3.2 The ASU Assessment Framework

The adoption of RLLM hinges on whether their are capable of generating value that outweighs their inference costs [35]. Assessing this tradeoff requires metrics that account for both performance and serving costs for both service provider and users. For RLLM service providers and users, the performance metrics they care about differ: providers seek to maximize system throughput, while users expect rapid model responses. In addition, it is essential to ensure response accuracy while optimizing RLLM serving system performance as much as possible. Thus, we propose *ASU* (*Accuracy, Service-end, User-end*), a trinity framework for assessing RLLM serving performance by together considering response accuracy, RLLM service provider end and user end. For accuracy metric, we employ evaluation *own metric* for each dataset. For service provider side metrics, we use throughput metric TPS (token per second) . For user-side metrics, we use TTFVT (time to first visible token) , a variant of TTFT , since we assume reasoning tokens of RLLM are invisible to users like commercial RLLM like OpenAI o1, and E2E requests running time as metrics.

In the next section, we will dive into the characteristic of RLLM serving via detailed experiments.

4 Pilot Investigations: Serving LLM v.s. RLLM

In this section, we perform an comprehensive investigation to RLLM and LLM inference serving.

Experiments. We involve eight prevailing models in evaluations. For fair comparison, RLLM model we employed is the tuned counterpart of evaluated LLM, e.g., Qwen-2.5-Math-7B and its tuned RLLM counterpart DeepSeek-R1-Distill-Qwen-7B. We conduct evaluation with 7B, 14B, 32B, 70B language models on different inference engines. For comprehensively assessment, we perform evaluation with different token budget and batch size. We use all the datasets described in §3.1.

Figure 1: Results of token budget variation across different datasets for 14B and 32B RLLM .

Main Results. 1) *Results with Different Token Budget*: Unlike traditional LLMs, RLLMs engage in deliberate reasoning by generating lengthy chains of thought prior to answer, which significantly increases token consumption. However, as existing LLM services are priced based on token usage, this results in substantially higher costs. To justify the impact of token budget for RLLM serving, we conduct evaluation with varying token budget from $0.5K$ to $20K$ across benchmarks. The results are presented in Figure 1. We found that, for the majority of datasets, a token budget of 4096 to 8192 can achieved sufficiently good performance. It is worth noting that, as the token budget increases, the performance of RLLMs on the GPQA and AIME24 datasets declined, which may indicate the *overthinking* problem [36] of RLLM. Please refer to §I.1 for full results.

2) *Results with Different Batch Size*. We also explore the impact of different batch sizes on RLLM serving performance with the same experimental setting. We find that increasing the batch size does not affect model accuracy on various datasets. Nevertheless, it reduces the time required for RLLMs to process the same number of requests, and improves throughput metric TPS, but at the cost of increased average TTFVT. Please refer to §I.2 for full results with different batch size.

Serving Performance and Behaviors. To investigate RLLM serving behaviors, we analyzed the running logs of the inference serving engine and conducted a visualization of the running traces, as shown in Figure 2. As illustrated, RLLMs achieve much higher accuracy on math datasets than same scale LLM, but a on-par performance on knowledge reasoning such as GPQA. The full results are presented in §I.3. To dissect the difference of serving behavior, we present the running trace in §I.4.

Main Findings for RLLM Serving Characteristics. Given the above results in pilot studies, we have the following findings in comparison of RLLM and LLM serving behaviors :
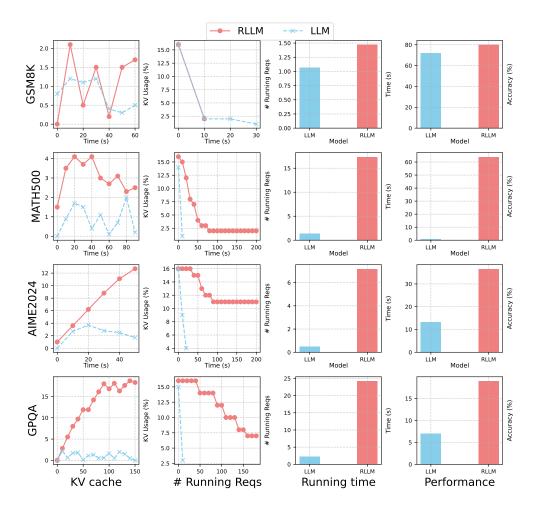
Figure 2: The serving performance and behavior comparison of a batch requests between 7B RLLM and LLM. We can read from this figure that (1) RLLM exhibits significant KV Cache fluctuations than LLM; (2) long tail distribution of requests running time caused by straggler requests; (3) adaptive running time of RLLM; (4) domain preference on math. Please refer to §I.3 for more results.

- *Significant Memory Usage and Fluctuations:* We observed significant fluctuations in memory utilization of inference engine when serving RLLM . In extreme cases, the usage varied dramatically between 3% and 70%, whereas traditional LLMs typically maintain KV cache usage below 3%. We attribute these fluctuations to the excessive length of the reasoning chains generated by RLLMs, which result in high memory consumption. During inference, the engine must retain KV caches for the reasoning chains until the requests are completed, after which they are discarded.

- *Straggler Requests:* When requests arrive at the inference engine in batches, or an RLLM receives multiple requests simultaneously, significant disparities in request difficulty can lead to some requests taking much longer time to complete than others. We denote these slow requests as *straggler requests*. These straggler requests ends either reaching the token budget or finishing the reasoning process. During this time, only a small number of requests remain running in inference engine, resulting in a noticeable drop in system throughput and hardware utilization. In contrast, LLMs exhibit much smaller variations in execution time for requests within the same batch.

- *Adaptive Running Time of RLLM:* We found that, given the same number of samples with same batch size, the runtime of RLLMs varies significantly across different datasets and is strongly correlated with the difficulty of the tasks. In contrast, traditional LLMs exhibit much smaller runtime differences across datasets, with little sensitivity to task difficulty. When the number of
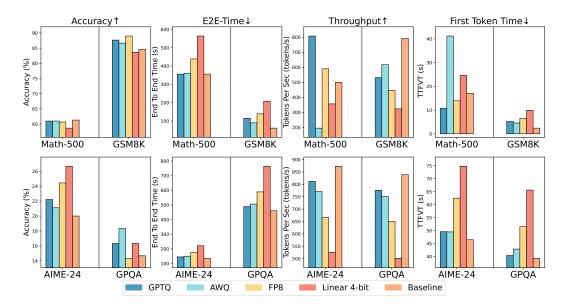
5

Figure 3: Empirical results of current LLM quantization methods on 7B RLLM. current methods maintain or improve all serving-related metrics with less memory footprint while keep accuracy.

samples varies, the runtime of LLMs on each dataset scales approximately linearly with the dataset size, even when there are substantial differences in task difficulty.

- *Domain Preference:* RLLMs and LLMs exhibit significant performance differences on the mathematical reasoning , while on-par on knowledge tasks, which align with existing works.

Discussion and Analysis for Findings. Based on the working mechanisms of inference engines we employed in benchmarking, we discuss the reason of why the above revealed phenomena occur.

1) *Straggler requests*: In some mathematical datasets, e.g., Math-500, the difficulty of individual problems varies. We assume that requests arrive in batches, and new requests are sent only after all requests in a batch are completed. As easier problems are answered shortly, the few remaining difficult requests continue running in the engine, leading to the entire batch's runtime being extended by these straggler requests. This situation results in reduced system throughput.

2) *Memory fluctuation and usage*: This issue is caused by the KV Cache management strategy of existing inference engines. Since RLLMs generate more tokens than traditional LLMs, the KV Cache utilization for RLLM is much higher under the same scale model with same precision in the same inference engine. This leads to a rapid increase in KV Cache usage, and since current inference engines discard the KV Cache once completing requests, it results in a sharp drop in cache usage.

3) *Adaptive running time*: RLLM generates varying reasoning chain lengths depending on problem difficulty—more difficulty lead to longer chains and running time. Hence, RLLM's runtime is typically correlated with problem difficulty, while LLMs generally may not be affected by difficulty.

Our findings indicate notable differences in serving RLLMs and LLMs. To enable more effective deployment of RLLMs, we explore some optimization techniques for inference in the next section.

5 Observations on RLLM Serving Optimization

In this section, we take a closer look at the techniques that may optimize RLLM serving performance. The prerequisite for assessing these optimization techniques is that they must *preserve* the RLLM's accuracy as much as possible. It holds throughout this section. More results are presented in §J.

5.1 Is model weight quantization methods effective in boosting RLLM serving?

Model weight quantization (MWQ) refers to the techniques that reduce number of bits for model parameters with the minimal loss in performance. Current LLM quantization methods are mainly
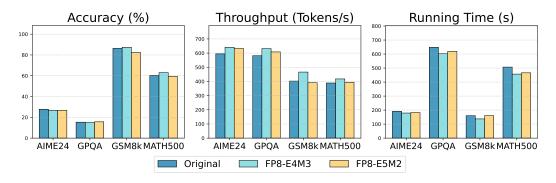
Figure 4: Empirical results for KV cache quantization on 14B model across different datasets.

fallen into the post-training quantization approaches. For more comprehensive introduction of LLM quantization, please refer to [37] and [38].To investigate the impact of model weight quantization, we employ 4 most adopted (also supported by current open source LLM serving engine) quantization methods for LLM: GPTQ [39] (Int4), AWQ [40] (4-bit), FP8 [41], and Linear 4-bit [42](L4) with BitsAndBytes [43]. We conduct experiments on 7B, 14B RLLM.

Main results. The evaluation results of quantized 7B RLLM using different quantization methods are presented in Figure 3. GPTQ-IN4 and FP8 quantization preserve the original model performance on most datasets, incurring only a minor degradation of approximately 3% or even perform better, while maintaining or improving all serving-related metrics with less memory footprint. However, GPTQ exhibits a substantial performance drop of around 15–25% on more challenging mathematical tasks such as AIME24. In contrast, AWQ and L4 maintain performance across all datasets but result in a marked reduction in inference efficiency, nearly doubling E2E time and halving throughput. These highlight the limitations of these approaches. The comprehensive results are presented in §J.1.

Observation 5.1. MWQ methods exert differing impacts on various metrics of RLLM inference .

5.2 Could KV Cache Quantization Lead to Better RLLM Serving Performance?

As illustrated in [16], to serve traditional LLM, at least 30% of GPU memory is perserved to store KV cache in the generation process. For RLLM, the demand for KV cache storage would be paramount since its much longer output length (including chain of thought reasoning), which makes it evitable for efficient management of memory. KV cache quantization emerges as an appealing approach to this end. We employ two KV cache quantization methods natively supported by vLLM: FP8-E5M2, and FP8-E4M3 [44] for inference serving evaluation.

Main results. The results of KV Cache quantization for 14B RLLM are presented in Figure 4. We found that using KV cache quantization effectively accelerates the operation of RLLMs while maintaining performance comparable to the original. Surprisingly, while the 14B or 32B RLLM maintained performance with minimal degradation after KV cache quantization, the 7B RLLM experienced almost complete performance deterioration, as shown in §J.2. Furthermore, we observed that KV cache quantization can also improve other metrics such as TTFVT and TPS.

Observation 5.2. KV Cache quantization can improve running efficiency for sufficient large RLLM.

5.3 Is Prefix Caching Useful for Contributing Efficient RLLM Serving?

Prefix Cache (PC) is a cache optimization policy that reuse computed KV values for prefill stage. By using this technique, new prompts that share same prefixes (exactly, same prefix tokens) with previous prompts processed by serving systems can reuse these KV cache. This technique is very useful such as long document query or multi-round conversation where requires multiple recomputation of same text. Empirical studies show that the prefix cache can provide a huge performance benefit in such scenarios. To evaluate the utility of prefix cache in RLLM serving, we compare the performance of 4 different RLLMs across all datasets with or without prefix caching enabling in vLLM and SGLang.

Main results. The results of PC evaluation on different datasets are shown in Figure 5. We find that for sufficiently large RLLMs (14B and above), prefix caching significantly improves runtime speed
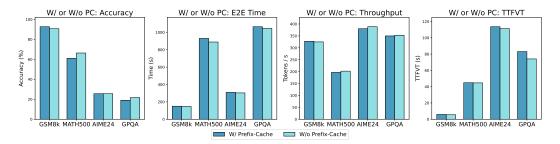
Figure 5: Empirical results of comparison for enable or disable prefix caching on 32B RLLM.

and serving metrics without compromising performance. However, for 7B models, prefix caching negatively impacts efficiency, leading to increased latency. Detailed results are in §J.3.

Observation 5.3. PC can accelerate larger RLLMs (14B and above) without performance degrade.

5.4 Does Speculative Decoding Help to Improve RLLM Serving Performance?

Speculative decoding (SD) refers to a bunch of approaches that improves inter-token latency in memory-bound LLM inference. The initial speculating sampling usually employs a faster homogeneous LLM as draft model to generate a multiple tokens draft, and then the larger LLM can decide to accept or reject this draft by scoring. The results in [45] show that the overhead of draft model is much smaller than larger LLM forwarding, which makes it feasible to be utilized in real world scenario. Recently, many works in speculative decoding [46] like n-gram matching [47], MLP speculators [48], and Eagle algorithm [49, 50] are proposed. Despite these advancement, current support and compatibility of speculating decoding for RLLM in serving framework is poor. Given this situation, we only assess n-gram matching algorithm for 7B, 14B and 32B RLLM serving with vLLM iframework. The other experimental settings is keeping the same as in §5.1 for fair comparison.

Main results. The main results for speculative decoding evaluation of 7B RLLM are listed in Figure 11. See §J.4 for full results. We find that speculative decoding improves the inference serving running time of RLLM across all scales, without degrading model performance on benchmarks. However, speculative decoding significantly reduces throughput and degrades the TTFVT metric.

Observation 5.4. SD improves the running time of RLLMs and deteriorates metrics like TPS.

Summary. This section suggests that many existing LLM inference optimization techniques can be directly applied to RLLMs seamless. However, surprisingly, some of these techniques have the opposite effect on smaller RLLMs, e.g., 7B. We leave the investigation of this phenomenon to future.

6 Applying to Real World Workload

In previous section (§4), we have shown that the serving behaviors of RLLM is significantly different from the LLM. However, we assumed that the serving engine receives requests simultaneously in batches, with each new batch arriving only after the system has completed processing the previous one. This assumption may be overly idealized and not fully consistent with real-world conditions. Prior works [51, 52, 53] have shown that, in real-world applications, the burstiness of requests received by the serving engine is typically modeled using the *Gamma distribution*. To validate our insights regarding RLLM serving in §4 under real-world scenarios, we implement a workload generator like BurstGPT-Perf [53] that is capable of producing requests following a Gamma distribution in our proposed Serve-Pref suite, enabling the generation of streaming, stochastic, and bursty workloads. We then perform empirical studies with it on various scale language models (7B, 14B, 32B) across different datasets to validate our findings.

Main Results. As shown in Figure 6, the average KV cache usage rate of RLLM is much higher than LLM. More surprisingly, for RLLM, the utilization of the serving engine's KV cache can remain close to 100% for long periods, forcing some new requests to wait in the waiting queue before running. This may significantly prolong request turnaround time in the serving engine, severely degrading user experience. We attribute the persistently high KV cache utilization to the accumulation of numerous stragglers in the system. The running requests in the engine are also much higher when serving
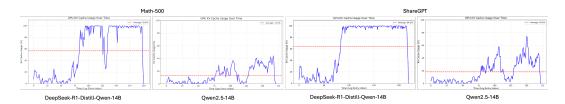
Figure 6: KV cache usage of 14B models under real-world workload across different datasets.
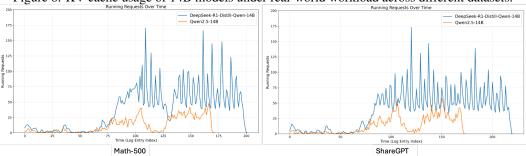

Figure 7: Num of running requests in the inference engine for 14B models under real-world workload.

RLLM compared to LLM, as shown in Figure 7. The above phenomena hold consistently across different datasets, demonstrating the generalizability of our findings. These results demonstrate our findings in §4 remain valid under real-world workloads. Please refer to Appendix K for more results.

7 Related Work

We introduce necessary related work in this section. More related work can be found in Appendix D.

Reasoning Large Language Models. Recent advancement in RLLM , such as OpenAI o1 [10] have demonstrated significant improvement in system-2 tasks such as mathematics and programming via test time scaling , which generates long chain of thought (CoT) reasoning text before answer the question. Compared with chain-of-thought in traditional LLM, the reasoning process of RLLM have the following characteristics: (1) much longer reasoning process; (2) extensive exploration to unreached logic node; (3) backtrack and reflection; (4) aha moment. Recent cutting edge RLLMs such as QwQ [54], Kimi K1.5 [55], Gemini-2.5-flash [56], Seed-think-v.15 [57], Qwen3 [12] have continually improve the performance on complex reasoning dataset.

LLM Inference and Serving.

Due to the large scale of LLM, they present considerable challenges in efficient serving, undermining the real world utilities of these models. Numerous works have been proposed to alleviate these problems from 6 different views: (1) model parameter memory optimization: model weight quantization like gptq [39], awq [40], FP8 [41], model pruning, model parallelism, CPU offloading ; (2) request scheduling: inter-request scheduling, and intra-request scheduling (3) dynamic memory optimization: KV cache quantization [44], KV cache reuse and dropping [58, 59]; (4) efficient decoding: speculating decoding [45] [49] [50], flash decoding [60] ;(5) system optimization: prefill-decoding disaggregation like [24] [61] [62]; (6) model and algorithm optimization: hard-aware algorithm like flash attention [60], linear attention, mixture of experts.

8 Conclusion

In this work, we systematically investigate the serving performance and behavior of RLLM. We reveal that RLLMs have several different serving behavior compared with traditional LLM, which makes current LLM serving engines struggle to unleash the power of RLLM and fall to reach the optimal performance. Additionally, we further investigate whether existing inference optimization techniques are valid for RLLM. Lastly, we conduct evaluation under real world workload modeled by Gamma distribution, and the results are *aligned* with our main findings regarding RLLM serving.

9 Reproducibility Statement

Below we summarize some critical aspects to facilitate reproducible results:

- Datasets. The datasets we used are all publicly accessible, which is introduced in G.1. The website for download these data are listed in F.
- Models. We provide the details about our adopted model and hyperparameters in F.
- Environment. All experiments are conducted with multiple runs on NVIDIA Tesla RTX4090-24GB GPUs, RTX A6000-48GB GPUs and NVIDIA A100-PCIE-40GB GPUs with Python 3.11 and PyTorch 2.5.

References

[1] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.

[2] Anthropic. The claude 3 model family: Opus, sonnet, haiku, 2024.

[3] Anthropic. Introducing claude 4. *online*, 2025.

[4] Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, et al. Gemini: a family of highly capable multimodal models, 2023.

[5] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, and Archie Sravankumar. The llama 3 herd of models, 2024.

[6] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.

[7] Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, et al. The pile: An 800gb dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027*, 2020.

[8] Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.

[9] Jack W Rae, Sebastian Borgeaud, Trevor Cai, Katie Millican, Jordan Hoffmann, Francis Song, John Aslanides, Sarah Henderson, Roman Ring, Susannah Young, et al. Scaling language models: Methods, analysis & insights from training gopher. *arXiv preprint arXiv:2112.11446*, 2021.

[10] Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry, Alex Beutel, Alex Carney, et al. Openai o1 system card. *arXiv preprint arXiv:2412.16720*, 2024.

[11] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

[12] Qwen team. Qwen3: Think deeper, act faster, 2025.

[13] Charlie Victor Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM test-time compute optimally can be more effective than scaling parameters for reasoning. In *The Thirteenth International Conference on Learning Representations*, 2025.

[14] Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori Hashimoto. s1: Simple test-time scaling, 2025.

[15] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.

[16] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention, 2023.

[17] LMDeploy Contributors. Lmdeploy: A toolkit for compressing, deploying, and serving llm. https://github.com/InternLM/lmdeploy, 2023.

[18] NVIDIA. TensorRT. https://github.com/NVIDIA/TensorRT-LLM, 2023.

[19] Zhong-Zhi Li, Duzhen Zhang, Ming-Liang Zhang, Jiaxin Zhang, Zengyan Liu, Yuxuan Yao, Haotian Xu, Junhao Zheng, Pei-Jie Wang, Xiuyi Chen, et al. From system 1 to system 2: A survey of reasoning large language models. *arXiv preprint arXiv:2502.17419*, 2025.

[20] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models, 2024.

[21] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.

[22] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 conference of the North American chapter of the association for computational linguistics: human language technologies, volume 1 (long and short papers)*, pages 4171–4186, 2019.

[23] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for {Transformer-Based} generative models. In *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, pages 521–538, 2022.

[24] Yinmin Zhong, Shengyu Liu, Junda Chen, Jianbo Hu, Yibo Zhu, Xuanzhe Liu, Xin Jin, and Hao Zhang. {DistServe}: Disaggregating prefill and decoding for goodput-optimized large language model serving. In *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24)*, pages 193–210, 2024.

[25] Nguyen Vinh, Gao Wenwen, Apsey Emily, Kudleppanavar Ganesh, Shah Neelay, and Bermudez Elias. Llm benchmarking: Fundamental concepts, 2025.

[26] Nvidia inc. A comprehensive guide to nim llm latency-throughput benchmarking, 2024.

[27] An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu, Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu, Xingzhang Ren, and Zhenru Zhang. Qwen2.5-math technical report: Toward mathematical expert model via self-improvement, 2024.

[28] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

[29] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

[30] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 2)*, 2021.

[31] Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let's verify step by step. In *The Twelfth International Conference on Learning Representations*, 2023.

[32] MAA Committees. Aime problems and solutions, 2024.

[33] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. In *First Conference on Language Modeling*, 2024.

[34] Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. Sglang: Efficient execution of structured language model programs, 2024.

[35] Mehmet Hamza Erol, Batu El, Mirac Suzgun, Mert Yuksekgonul, and James Zou. Cost-of-pass: An economic framework for evaluating language models, 2025.

[36] Xiaoye Qu, Yafu Li, Zhaochen Su, Weigao Sun, Jianhao Yan, Dongrui Liu, Ganqu Cui, Daizong Liu, Shuxian Liang, Junxian He, Peng Li, Wei Wei, Jing Shao, Chaochao Lu, Yue Zhang, Xian-Sheng Hua, Bowen Zhou, and Yu Cheng. A survey of efficient reasoning for large reasoning models: Language, multimodality, and beyond, 2025.

[37] Xunyu Zhu, Jian Li, Yong Liu, Can Ma, and Weiping Wang. A survey on model compression for large language models. *Transactions of the Association for Computational Linguistics*, 12:1556–1577, 2024.

[38] Ruihao Gong, Yifu Ding, Zining Wang, Chengtao Lv, Xingyu Zheng, Jinyang Du, Haotong Qin, Jinyang Guo, Michele Magno, and Xianglong Liu. A survey of low-bit large language models: Basics, systems, and algorithms. *arXiv preprint arXiv:2409.16694*, 2024.

[39] Elias Frantar, Saleh Ashkboos, Torsten Hoefler, and Dan Alistarh. OPTQ: Accurate quantization for generative pre-trained transformers. In *The Eleventh International Conference on Learning Representations*, 2023.

[40] Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. Awq: Activation-aware weight quantization for on-device llm compression and acceleration. *Proceedings of Machine Learning and Systems*, 6:87–100, 2024.

[41] Andrey Kuzmin, Mart Van Baalen, Yuwei Ren, Markus Nagel, Jorn Peters, and Tijmen Blankevoort. Fp8 quantization: The power of the exponent. *Advances in Neural Information Processing Systems*, 35:14651–14662, 2022.

[42] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms. *Advances in neural information processing systems*, 36:10088–10115, 2023.

[43] bitsandbytes foundation. bitsandbytes, 2022.

[44] vllm project. Llm compressor, 2024.

[45] Charlie Chen, Sebastian Borgeaud, Geoffrey Irving, Jean-Baptiste Lespiau, Laurent Sifre, and John Jumper. Accelerating large language model decoding with speculative sampling. *arXiv preprint arXiv:2302.01318*, 2023.

[46] Heming Xia, Zhe Yang, Qingxiu Dong, Peiyi Wang, Yongqi Li, Tao Ge, Tianyu Liu, Wenjie Li, and Zhifang Sui. Unlocking efficiency in large language model inference: A comprehensive survey of speculative decoding. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics: ACL 2024*, pages 7655–7671, Bangkok, Thailand, August 2024. Association for Computational Linguistics.

[47] vLLM Team. Speculating by matching n-grams in the prompt, 2024.

[48] Davis Wertheimer, Joshua Rosenkranz, Thomas Parnell, Sahil Suneja, Pavithra Ranganathan, Raghu Ganti, and Mudhakar Srivatsa. Accelerating production llms with combined token/embedding speculators. *arXiv preprint arXiv:2404.19124*, 2024.

[49] Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. Eagle: Speculative sampling requires rethinking feature uncertainty. In *Forty-first International Conference on Machine Learning*, 2024.

[50] Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. Eagle-3: Scaling up inference acceleration of large language models via training-time test. *arXiv preprint arXiv:2503.01840*, 2025.

[51] Bingyang Wu, Yinmin Zhong, Zili Zhang, Shengyu Liu, Fangyue Liu, Yuanhang Sun, Gang Huang, Xuanzhe Liu, and Xin Jin. Fast distributed inference serving for large language models. *arXiv preprint arXiv:2305.05920*, 2023.

[52] Zhuohan Li, Lianmin Zheng, Yinmin Zhong, Vincent Liu, Ying Sheng, Xin Jin, Yanping Huang, Zhifeng Chen, Hao Zhang, Joseph E Gonzalez, et al. {AlpaServe}: Statistical multiplexing with model parallelism for deep learning serving. In *17th USENIX Symposium on Operating Systems Design and Implementation (OSDI 23)*, pages 663–679, 2023.

[53] Yuxin Wang, Yuhan Chen, Zeyu Li, Xueze Kang, Yuchu Fang, Yeju Zhou, Yang Zheng, Zhenheng Tang, Xin He, Rui Guo, et al. Burstgpt: A real-world workload dataset to optimize llm serving systems. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*, pages 5831–5841, 2025.

[54] Qwen Team. Qwq-32b: Embracing the power of reinforcement learning, 2025.

[55] Kimi Team, Angang Du, Bofei Gao, Bowei Xing, Changjiu Jiang, Cheng Chen, Cheng Li, Chenjun Xiao, Chenzhuang Du, Chonghua Liao, et al. Kimi k1. 5: Scaling reinforcement learning with llms. *arXiv preprint arXiv:2501.12599*, 2025.

[56] Google DeepMind. Gemini 2.5: Our most intelligent ai model, 2025.

[57] ByteDance Seed, Yufeng Yuan, Yu Yue, Mingxuan Wang, Xiaochen Zuo, Jiaze Chen, Lin Yan, Wenyuan Xu, Chi Zhang, Xin Liu, et al. Seed-thinking-v1. 5: Advancing superb reasoning models with reinforcement learning. *arXiv preprint arXiv:2504.13914*, 2025.

[58] Xiang Liu, Zhenheng Tang, Peijie Dong, Zeyu Li, Yue Liu, Bo Li, Xuming Hu, and Xiaowen Chu. Chunkkv: Semantic-preserving kv cache compression for efficient long-context llm inference, 2025.

[59] Xiang Liu, Zhenheng Tang, Hong Chen, Peijie Dong, Zeyu Li, Xiuze Zhou, Bo Li, Xuming Hu, and Xiaowen Chu. Can llms maintain fundamental abilities under kv cache compression? *arXiv preprint arXiv:2502.01941*, 2025.

[60] Dao Tri, Haziza Daniel, o Massa Francisc, and Sizov Grigory. Flash-decoding for long-context inference. *online*, 2023.

[61] Cunchen Hu, Heyang Huang, Junhao Hu, Jiang Xu, Xusheng Chen, Tao Xie, Chenxi Wang, Sa Wang, Yungang Bao, Ninghui Sun, et al. Memserve: Context caching for disaggregated llm serving with elastic memory pool. *arXiv preprint arXiv:2406.17565*, 2024.

[62] Ruoyu Qin, Zheming Li, Weiran He, Jialei Cui, Feng Ren, Mingxing Zhang, Yongwei Wu, Weimin Zheng, and Xinran Xu. Mooncake: Trading more storage for less computation—a kvcache-centric architecture for serving llm chatbot. In *23rd USENIX Conference on File and Storage Technologies (FAST 25)*, pages 155–170, 2025.

[63] OpenAI. Openai o3 and o4-mini system card. *Online*, 2025.

[64] Yixin Ye, Zhen Huang, Yang Xiao, Ethan Chern, Shijie Xia, and Pengfei Liu. Limo: Less is more for reasoning, 2025.

[65] Hugging Face. Open r1: A fully open reproduction of deepseek-r1, January 2025.

[66] Noam Shazeer. Fast transformer decoding: One write-head is all you need. *arXiv preprint arXiv:1911.02150*, 2019.

[67] Baolin Li, Yankai Jiang, Vijay Gadepally, and Devesh Tiwari. Llm inference serving: Survey of recent advances and opportunities. *arXiv preprint arXiv:2407.12391*, 2024.

[68] Sehoon Kim, Coleman Hooper, Thanakul Wattanawong, Minwoo Kang, Ruohan Yan, Hasan Genc, Grace Dinh, Qijing Huang, Kurt Keutzer, Michael W. Mahoney, Yakun Sophia Shao, and Amir Gholami. Full stack optimization of transformer inference: a survey, 2023.

[69] Ranran Zhen, Juntao Li, Yixin Ji, Zhenlin Yang, Tong Liu, Qingrong Xia, Xinyu Duan, Zhefeng Wang, Baoxing Huai, and Min Zhang. Taming the titans: A survey of efficient llm inference serving. *arXiv preprint arXiv:2504.19720*, 2025.

[70] Malgorzata Lazuka, Andreea Anghel, and Thomas Parnell. Llm-pilot: Characterize and optimize performance of your llm inference services. In *SC24: International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 1–18. IEEE, 2024.

[71] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. Holistic evaluation of language models. *Transactions on Machine Learning Research*, 2023.

[72] Renren Jin, Jiangcun Du, Wuwei Huang, Wei Liu, Jian Luan, Bin Wang, and Deyi Xiong. A comprehensive evaluation of quantization strategies for large language models. In *Findings of the Association for Computational Linguistics ACL 2024*, pages 12186–12215, 2024.

[73] Shiyao Li, Xuefei Ning, Luning Wang, Tengxuan Liu, Xiangsheng Shi, Shengen Yan, Guo-hao Dai, Huazhong Yang, and Yu Wang. Evaluating quantized large language models. In *Proceedings of the 41st International Conference on Machine Learning*, pages 28480–28524, 2024.

[74] Qi Li, Xiang Liu, Zhenheng Tang, Peijie Dong, Zeyu Li, Xinglin Pan, and Xiaowen Chu. Should we really edit language models? on the evaluation of edited language models. In *Proceedings of the Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2025.

[75] Amey Agrawal, Anmol Agarwal, Nitin Kedia, Jayashree Mohan, Souvik Kundu, Nipun Kwa-tra, Ramachandran Ramjee, and Alexey Tumanov. Etalon: holistic performance evaluation framework for llm inference systems. *arXiv preprint arXiv:2407.07000*, 2024.

[76] Jeff Rasley, Samyam Rajbhandari, Olatunji Ruwase, and Yuxiong He. Deepspeed: System optimizations enable training deep learning models with over 100 billion parameters. In *Proceedings of the 26th ACM SIGKDD international conference on knowledge discovery & data mining*, pages 3505–3506, 2020.

[77] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-lm: Training multi-billion parameter language models using model parallelism. *arXiv preprint arXiv:1909.08053*, 2019.

[78] Jian Hu, Xibin Wu, Zilin Zhu, Weixun Wang, Dehao Zhang, Yu Cao, et al. Openrlhf: An easy-to-use, scalable and high-performance rlhf framework. *arXiv preprint arXiv:2405.11143*, 2024.

[79] Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. *Advances in neural information processing systems*, 35:16344–16359, 2022.

[80] Zihao Ye, Lequn Chen, Ruihang Lai, Wuwei Lin, Yineng Zhang, Stephanie Wang, Tianqi Chen, Baris Kasikci, Vinod Grover, Arvind Krishnamurthy, et al. Flashinfer: Efficient and customizable attention engine for llm inference serving. *arXiv preprint arXiv:2501.01005*, 2025.

[81] Samyam Rajbhandari, Jeff Rasley, Olatunji Ruwase, and Yuxiong He. Zero: Memory optimiza-tions toward training trillion parameter models. In *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 1–16. IEEE, 2020.

[82] Yuhan Liu, Hanchen Li, Yihua Cheng, Siddhant Ray, Yuyang Huang, Qizheng Zhang, Kuntai Du, Jiayi Yao, Shan Lu, Ganesh Ananthanarayanan, et al. Cachegen: Kv cache compression and streaming for fast large language model serving. In *Proceedings of the ACM SIGCOMM 2024 Conference*, pages 38–56, 2024.

[83] Yihua Cheng, Kuntai Du, Jiayi Yao, and Junchen Jiang. Do large language models need a content delivery network? *arXiv preprint arXiv:2409.13761*, 2024.

[84] Jiayi Yao, Hanchen Li, Yuhan Liu, Siddhant Ray, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. Cacheblend: Fast large language model serving with cached knowledge fusion. *arXiv preprint arXiv:2405.16444*, 2024.

[85] Marah Abdin, Sahaj Agarwal, Ahmed Awadallah, Vidhisha Balachandran, Harkirat Behl, Lingjiao Chen, Gustavo de Rosa, Suriya Gunasekar, Mojan Javaheripi, Neel Joshi, et al. Phi-4-reasoning technical report. *arXiv preprint arXiv:2504.21318*, 2025.

[86] Akhiad Bercovich, Levy Itay, Izik Golan, Mohammad Dabbah, Ran El-Yaniv, Omri Puny, Ido Galil, Zach Moshe, Tomer Ronen, et al. Llama-nemotron: Efficient reasoning models, 2025.

[87] Cunchen Hu, Heyang Huang, Liangliang Xu, Xusheng Chen, Jiang Xu, Shuang Chen, Hao Feng, Chenxi Wang, Sa Wang, Yungang Bao, et al. Inference without interference: Disaggregate llm inference for mixed downstream workloads. *arXiv preprint arXiv:2401.11181*, 2024.

[88] Pratyush Patel, Esha Choukse, Chaojie Zhang, Aashaka Shah, Íñigo Goiri, Saeed Maleki, and Ricardo Bianchini. Splitwise: Efficient generative llm inference using phase splitting. In *2024 ACM/IEEE 51st Annual International Symposium on Computer Architecture (ISCA)*, pages 118–132. IEEE, 2024.

Appendix and Supplementary Material

A Use of LLMs Statement

We solemnly declare that the originality of ideas, writing, overall methodology, experiments, and other core contributions in this paper are entirely the work of the authors, with no involvement of any LLMs in the research process. LLMs were used solely for grammar checking and language polishing after drafting this submission.

B Limitation

In this work, we systematically investigate the serving performance of RLLM. Despite our comprehensive and thorough experiments, the evaluation of RLLM serving is limited in some extet due to limited support from the current ecosystem. We hope that future improvements in serving engines will enable broader and more comprehensive evaluations. Additionally, our hardware resources were limited, and we aim to extend our evaluations to a wider range of hardware platforms in the future.

C Boarder Impact

In this paper, we systematically investigate the serving performance of RLLM. We hope our work can provide the research community and industry with insightful perspectives to help advance studies in efficient RLLM serving, help to democratize the use of cutting-edge RLLMs for social good.

D Extended Related Work

D.1 Reasoning Large Language Models

Recent advancement in RLLM , such as OpenAI o1 [10] have demonstrated significant improvement in system-2 tasks such as mathematics and programming via test time scaling, which generates long chain of thought (CoT) reasoning text before answer the question. Compared with chain-of-thought in traditional LLM, the reasoning process of RLLM has the following characteristics: (1) much longer reasoning process; (2) extensive exploration to unreached logic node; (3) backtrack and reflection; (4) aha moment. Since OpenAI's o1 and o3 [63] are proprietary models, the research community has attempted to replicate their performance. s1 [14] try to achieve test time scaling with only $1k$ post-training samples. LIMO [64] exploits only 817 curated training samples, improving scores from 6.5% to 57.1% on AIME dataset. DeepSeek R1 [11] is the first open-source RLLM and achieves on-par performance with OpenAI o1. Followed by [65], which aims to fully reproduce R1 by the collaboration of open-source community. Recent cutting edge RLLMs such as QwQ [54], Kimi K1.5 [55], Gemini-2.5-flash [56], Seed-think-v.15 [57], Qwen3 [12] have continually improve the performance on complex reasoning dataset.

D.2 LLM Inference and Serving

LLM has become a cornerstone of deep learning in recent years, reshaping the landscape of AI research. Due to the large scale of LLM, they present considerable challenges in efficient serving, undermining the real-world utilities of these models. Numerous works have been proposed to alleviate these problems from 6 different views: (1) model parameter memory optimization: model weight quantization like gptq [39], awq [40], FP8 [41], model pruning, model parallelism, CPU offloading ; (2) request scheduling: inter-request scheduling, and intra-request scheduling (3) dynamic memory optimization: KV cache quantization [44], KV cache reuse and dropping; (4) efficient decoding: speculating decoding [45] [49] [50], flash decoding [60] ;(5) system optimization: prefill-decoding disaggregation architecture like [24] [61] [62]; (6) model and algorithm optimization: hard-aware algorithm like flash attention [60], linear attention, mixture of expert.

Recent advances in LLM inference have yielded a variety of specialized frameworks and serving engines that maximize GPU utilization through optimized kernels and memory strategies. High-performance libraries such as NVIDIA's FasterTransformer [66] and TensorRT-LLM [18], alongside open-source systems like vLLM [16] and SGLang [34], employ different techniques with continuous batching[23], speculative decoding[45], prefill-decode disaggregation[24] and many other methods, ensuring the GPU pipeline remains saturated. Complementing these efforts are dynamic scheduling

and memory management schemes that break large KV caches into reusable blocks and selectively merge or preempt operations, allowing much larger batch sizes with minimal overhead. Equally important are multi-way parallelism and algorithmic innovations that further boost throughput and reduce latency. Large models are commonly deployed across GPUs using tensor parallelism (splitting each layer's computation), pipeline parallelism (partitioning the model into sequential stages), and data parallel replication. Mixture-of-Experts (MoE) architectures extend this by routing tokens to different expert shards via expert parallelism, with communication optimizations to balance load. On the algorithmic side, parameter-efficient methods such as prompt and prefix tuning adapt frozen models via small "soft" prompts, speculative decoding [45] uses a lightweight draft model to accelerate token generation, and Simple Test-Time Scaling[14] applies budget-forcing at inference to improve reasoning quality.

Together, these system-level designs and algorithm-level approaches form a cohesive ecosystem that drives state-of-the-art performance in efficient LLM serving. Please see survey papers [67, 68, 69] for comprehensive introduction [70].

D.3   LLM Evaluation

Recently, with the rapid development of LLM, there is a growing interest in evaluating LLM from different aspects and topics. A holistic evaluation framework of language models is proposed [71]. Generally, the technical reports like [28, 12, 11] of LLM provides pre-relase comprehensive evaluation results. The quantization methods for LLM are evaluated in [72] and [73]. In [74], it evaluates the general abilities of post-edit LLM to assess the utility of existing knowledge editing methods. Work [70] and [75] evaluate LLM serving from a new perspective.

D.4   Ecosystem Support for RLLM Serving.

The development of LLMs has greatly benefited from the research community and the open-source ecosystem, including open platforms such as Hugging Face, Github, and Modelscope; open-source LLMs like Llama [5], Qwen [28], and Deepseek R1; open-source LLM infrastructure such as Deepspeed [76], Megatron-LM [77], vLLM [16], OpenRLHF [78], and SGLang [34]; various optimization techniques like Flash-Attention [79], FlashInfer [80], ZeRO [81], and LMCache [82, 83, 84]. The advancement of RLLMs continues this trend. With the open-sourcing of Deepseek R1 [11], a large number of open-source RLLMs like Phi-4 reasoning [85], and Llama-Nemotron [86] have emerged, further promoting the democratization of cutting-edge RLLM technology. Although existing LLM serving systems like vLLM, and SGLang provide some level of support for RLLMs, current support and optimization techniques remain significantly limited. Some techniques do not support RLLMs at all, for instance, Eagle speculative decoding currently lacks compatibility with RLLMs, while others fail to offer targeted optimizations and improvements specific to RLLM characteristics. As RLLMs continue to advance rapidly, we call on the research community and industry to collaborate in addressing the issues revealed in this paper.

E An Introduction to LLM Serving

The highly increased development of LLMs' application arise the demand of effectively using LLM serving systems. In this part, we introduce some optimization methods for serving systems and introduce more serving metrics. For more comprehensive introduction, please refer to [69] and [67].

E.1 Serving Performance

Recently, there are a lot of researches focus on optimizing the performance of serving system based on LLM architecture's characteristics and system-level tricks. Current LLMs are mostly using decode-only architecture, making the KV values of former tokens becomes key information for the next token. Hence, the first useful methods is storing all of KV value in memories(particularly in GPUs), this method significantly improve the efficiency of prefill stage. However, this method had already deployed for language models. For LLM, the most important method proposed first is continuous batching[23]. Continuous batching is processing requests in serving systems in iteration level, compared with former systems process requests in request-level. By using this technique, serving systems don't need to wait until the last request finishes its decoding, but replace requests with new requests once it ends decoding. This method enhance GPU's utilization, reducing waiting time for high-throughput serving systems. Next, considering the difference of prefill and decode that prefill is compute-intense stage which needs more GPU computing resources, while decode is memory-intense stage which needs more GPU memories compared with prefill, Prefill-Decode disaggregation[24] proposed a method that process prefill and decode in different GPUs, fully utilizing GPU resources based on the characteristics of the two phases. Despite this, GPU resources are still not fully utilized because the GPU pre-allocates a portion of GPU space for requests when storing previous KV cache. However, much of this space isn't effectively used, resulting in significant waste (for example, if a request occupies 8 tokens, the GPU allocates 2080 token spaces for decoding this request, but actually only produces 80 tokens, wasting space for 2000 tokens). At the same time, since the GPU allocates and reserves space for requests sequentially, this can lead to memory fragmentation and inefficient resource utilization when requests complete at different times. Paged attention borrows the concept of CPU paging, and in their serving system (vLLM) creates a mapping between virtual addresses and actual GPU addresses through virtual pages[16].

E.2 Serving Metric

With the high demand of deploying customized LLMs for practical utilization, there is a need to measure the cost efficiency of different LLM serving solutions. The cost of serving RLLM depends on how many requests it can handle per second while being responsive to client users and supporting an acceptable level of answer accuracy. To measure the performance of LLM serving system, there are multiple metrics can be chosen: (1) Time to first token (TTFT) is the time it takes to process the prompt until generate the first token. It measures how long a user must wait before seeing the model's output; (2) End-to-end request latency (E2E latency) indicates the time it takes from submitting the first token of a request to receiving the last token of response, including the time for queueing and batching and network latencies in real-world scenario; (3) Time between tokens (TBT, a.k.a Intertoken latency, ITL) is the average time between the generation of consecutive tokens in a sequence; (4) Tokens per second (TPS) of system represents the mean of total output tokens number per second , accounting for all the requests happening simultaneously; (5) Requests per second (RPS) is the average number of requests that can be successfully completed by the system in a 1-second period. In LLM serving systems, there are many metrics evaluating the performance, In this paper, we use metrics for reference that companies and personal users care most while using RLLM. I'll introduce them here for clear understanding.

- **Throughput:** Number of processed requests per second. This is the key metric for users since it directly determines overall system performance.
- **Time to First Token (TTFT):** Time from receiving a request until the first token is generated (i.e., the prefill stage is completed). This reflects how quickly the serving system handles the prefill stage. Techniques such as continuous batching [23] and paged attention [16] were proposed to optimize this metric.
- **Time to First Visible Token (TTFVT):** The time from receiving a request until the first token is actually displayed to the user. This metric is specific to RLLMs because some

20

inference systems hide the internal "thinking" steps and only reveal output once thinking is complete. Since RLLMs often perform a prolonged reasoning chain before producing any visible token, TTFVT is typically much larger than TTFT.

- **Time Between Tokens (TBT):** Average time between generation of consecutive tokens. For RLLMs, both the thinking stage and decoding stage share this metric. Recent algorithm-level optimizations such as S1 [14] target TBT. In this paper, TBT reflects the real-time per-token responsiveness of the model during interactive generation, capturing both computational and scheduling overhead.

- **KV Cache Utilization:** Proportion of total memory occupied by the KV cache during model execution. High utilization enables reuse of KV values by subsequent requests, reducing prefill time. However, excessive utilization triggers frequent evictions, degrading performance. Section 4 analyzes KV cache utilization and its impact on overall performance for RLLMs across datasets of varying difficulty.

- **Tokens per Second (TPS):** Total number of tokens generated per second across all active sessions. This combines throughput and per-token speed into one measure of generation capacity.

- **Requests per Second (RPS):** Total number of full-request pipelines completed per second. Unlike throughput (which counts raw requests), RPS tracks end-to-end request handling.

- **Model Initialization Latency:** Total time from service startup—including loading model weights, constructing computation graphs, allocating GPU memory, initializing optimizers, and any warm-up steps—until the system is ready to handle its first request. For MoE models (such as the DeepSeek model used in this paper) with Tensor Parallelism (TP) and Pipeline Parallelism (PP), this also involves partitioning and distributing parameters across multiple GPUs. This metric helps compare how different serving systems optimize model loading and initialization.

- **End-to-End Latency (E2E Latency):** Time from user request submission until receipt of the final token. This metric significantly influences user experience; for enterprises, improving RLLM end-to-end latency is also a critical concern.

F   Implementation and Reproduction Details

In this section, we would like to provide details for reproducing our experimental results.

F.1   Code Base

Our code and the *ASU-Perf* suite will be available once this paper accepted.

F.2   Models

Here, we list all of the model checkpoints used in our experiments.

RLLM checkpoints:

- deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B `https://hf-mirror.com/deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B`
- deepseek-ai/DeepSeek-R1-Distill-Qwen-7B `https://hf-mirror.com/deepseek-ai/DeepSeek-R1-Distill-Qwen-7B`
- deepseek-ai/DeepSeek-R1-Distill-Qwen-14B `https://hf-mirror.com/deepseek-ai/DeepSeek-R1-Distill-Qwen-14B`
- deepseek-ai/DeepSeek-R1-Distill-Qwen-32B `https://hf-mirror.com/deepseek-ai/DeepSeek-R1-Distill-Qwen-32B`
- deepseek-ai/DeepSeek-R1-Distill-Llama-70B `https://hf-mirror.com/deepseek-ai/DeepSeek-R1-Distill-Llama-70B`

LLM checkpoints:

- Qwen/Qwen2.5-Math-1.5B `https://hf-mirror.com/Qwen/Qwen2.5-Math-1.5B`
- Qwen/Qwen2.5-Math-7B `https://hf-mirror.com/Qwen/Qwen2.5-Math-7B`
- Qwen/Qwen2.5-14B `https://hf-mirror.com/Qwen/Qwen2.5-14B`
- Qwen/Qwen2.5-32B `https://hf-mirror.com/Qwen/Qwen2.5-32B`
- meta-llama/Llama-3.3-70B-Instruct `https://hf-mirror.com/meta-llama/Llama-3.3-70B-Instruct`

F.3   Datasets

Here, we list all of the benchmarking datasets used in our experiments.

- GSM8K `https://hf-mirror.com/datasets/openai/gsm8k`
- MATH-500 `https://hf-mirror.com/datasets/HuggingFaceH4/MATH-500`
- AIME-24 `https://hf-mirror.com/datasets/HuggingFaceH4/aime_2024`
- GPQA `https://hf-mirror.com/datasets/Idavidrein/gpqa`

F.4   Hyperparameters Settings for RLLM

The hyperparameters settings for RLLM we employed are as follows:

Batch Size: 8, 16, 32

Dataset Capacity: 100, (AIME24 30)

Temperature: 0.6, Top-p: 0.95, Top-k: 20, Request Timeout: 1200 sec

Experiments Repeat Time: 3

Performance Only Mode: False

Reasoning LLM Mode: True

CoT Visible (for TTFT): False

F.5 Hyperparameters Settings for LLM

The hyperparameters settings for LLM we employed are as follows:

Batch Size: 8, 16, 32

Dataset Capacity: 100, (AIME24 30)

Temperature: 0.7, Top-p: 0.8, Top-k: 20, Request Timeout: 1200 sec

Experiments Repeat Time: 3

Performance Only Mode: False

Reasoning LLM Mode: False

CoT Visible (for TTFT): False

G   Experiments Details

G.1   Details Evaluation Datasets

We use 4 different datasets in this paper, they are GSM8K, MATH500, AIME24, and GPQA. The details of these datasets are following.

- **GSM8K** [29]: The GSM8K dataset is a large collection of mathematical problem-solving tasks designed for training and evaluating AI models in the context of elementary school-level math. It primarily focuses on grade school math word problems that require multiple steps of reasoning and calculations to solve.

- **MATH500** [31]: a challenging dataset consisting of problems from high school math competitions across seven subjects (e.g., Prealgebra, Algebra, Number Theory) and difficulty levels based on AoPS (ranging from 1 to 5). Problems in these competitions range from level 1, the easiest, often found in AMC 8 exams, to level 5, like those in AIME.

- **AIME24** [32]:a dataset from the American Invitational Mathematics Examination, which tests math problem solving across multiple areas (e.g. algebra, counting, geometry, number theory, and probability). Because AIME 2024 contains only 30 examples, we don't considered examples of AIME from other years.

- **GPQA** [33]: a graduate-level dataset consisting of multiple-choice questions in subdomains of physics, chemistry, and biology. For our experiment, we select the highest quality subset, known as GPQA Diamond (composed of 198 questions).

G.2   Running Device

All of our experiments are running on three devices: a server with 8 RTX A6000 GPUs with 48GB VRAM, a server equipped with 8 RTX 4090 GPUs with 24GB VRAM, and a server with 8 NVIDIA A100-PCIE-40GB GPUs.

G.3   Inference Engine

We use vLLM [16] version 0.8.1 and SGLang [34] version 0.4.6.post1. For evaluation, we use OpenAI compatible API /v1/chat/completions.

Table 1: Performance Metrics with PD-disaggregated with 7B, 14B model

Model	Method	Dataset	Accuracy	Running Time	Token Per Sec	TTFVT	Output Tokens
7B	w/o PD-disa	GSM8K	82	0:58	713.3	0.21	41789
	w/ PD-disa	GSM8K	82	1:12	516.9	0.3	42052
	w/o PD-disa	MATH500	65	5:21	492.7	0.23	158431
	w/ PD-disa	MATH500	64	7:57	323.2	0.4	154318
	w/o PD-disa	AIME2024	266	2:01	935.3	0.27	112492
	w/ PD-disa	AIME2024	23.3	2:40	711.7	0.4	113718
	w/o PD-disa	GPQA	11	6:52	894.2	0.36	368377
	w/ PD-disa	GPQA	20	9:19	662.0	0.6	369821
14B	w/o PD-disa	GSM8K	87	6:03	163.0	0.31	59219
	w/ PD-disa	GSM8K	86	9:04	109.7	0.6	59605
	w/o PD-disa	MATH500	65	12:33	233.4	0.37	175808
	w/ PD-disa	MATH500	58	22:43	134.0	0.7	182705
	w/o PD-disa	AIME2024	23.3	4:14	449.9	0.42	114586
	w/ PD-disa	AIME2024	26.6	7:22	258.2	0.8	114048
	w/o PD-disa	GPQA	20	14:37	403.6	0.55	354155
	w/ PD-disa	GPQA	19	25:21	235.6	1.2	358227

H Extend Observation

H.1 Can Disaggregated Prefilling Improve RLLM Serving Performance ?

As discussed in Section 2 and paper [24], the process of LLM generates responds to a input prompt can be divided into two different phases. The LLM first processes input prompt in the *prefill phase*, which is computation intensive, to generate the first token of response within one iteration. After it , in the memory bounded *decoding phase*, LLM generates token one by one in each iteration until reaching the end token. These two phases have distinct different significance. However, many existing serving system co-locate the prefill and decoding at the same device, which may leads to sub-optimal performance and inter-phase interference as revealed in [24]. The disaggregated prefilling architecture was proposed to address this problem. It is first introduced in [24], followed by lines of recent works like [87], [88], [61], [62], notably improving the TTFT and throughput of system. However, current support for disaggregated prefilling is experimental and only available in vLLM. What's more, the only disaggregated prefilling feature support in vLLM is 1P1D scheme (1 prefilling worker and 1 decoding worker) currently. Hence, we merely perform evaluation with 1P1D on 7B (on two RTX-4090 GPU) and 14B (on two A6000 GPU) models across 4 evaluation datasets.

Main results. The results of PD-disaggregation are shown in Table 1. We found that under the 1P1D setup, PD-disaggregation does not improve the serving performance of RLLMs. On the contrary, it leads to a decline in system performance metrics. We find that the performance bottleneck of 1P1D serving for RLLMs lies in decoding, while the devices used for pre-filling are largely idle, which leads to suboptimal performance. Additionally, PD-disaggregation requires KV cache transfer between GPUs, and the communication overhead negatively impacts the serving of RLLMs.

Observation 6. PD-disaggregation (1P1D) deteriorates RLLM serving metrics compared to mixed PD. Since nearly half of the computing resources are idle.

I Detailed Empirical Results

I.1 Token Budget for Pilot Study

Full Figures of token budget exploration are listed in Figure 8.

I.2 Main Results for Pilot Study

We provide full results of RLLM and LLM serving comparison.

- 7B. RLLM in Table 2, LLM in Table 3
- 14B. RLLM in Table 4, LLM in Table 5
- 32B. RLLM in Table 6, LLM in Table 7
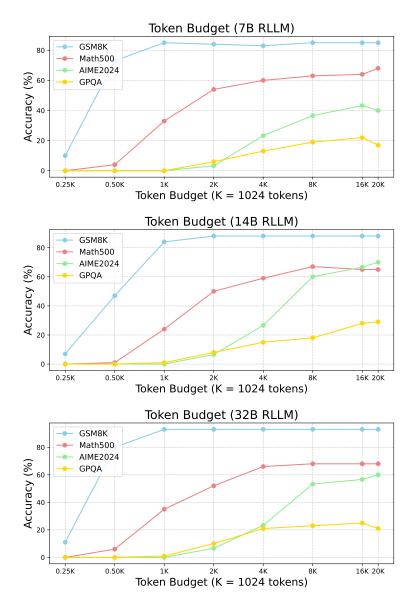
Figure 8: Results of token budget variation across different datasets for different scale RLLM .

- 70B. RLLM in Table 8, LLM in Table 9

I.3 Serving Behaviors for Pilot Study

For better presentation, we provide illustration about 14B and 32B model serving visualization in Figure 9 and 10.

Table 2: Serving Results of RLLM-7B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
RLLM-7B	8	4096	GSM8k	80.67	1m53s	426.87	2.4200	129629
			MATH500	61.33	9m34	302.47	16.5500	502833
			AIME24	23.33	4m7s	470.37	40.8200	340730
			GPQA	15.00	13m29s	477.14	35.3600	1124673
		8192	GSM8k	82.33	1m46s	443.95	2.3400	126155
			MATH500	64.33	14m09	240.55	20.2900	601161
			AIME24	38.89	8m7s	409.92	64.9600	592434
			GPQA	26.67	26m42s	410.33	65.9200	1941071
	16	4096	GSM8k	84.30	1m1s	791.40	2.3900	126451
			MATH500	59.30	5m56s	499.45	17.1000	505796
			AIME24	20.00	2m15s	872.91	46.4700	346986
			GPQA	14.67	7m40s	838.80	39.3300	1124379
		8192	GSM8k	85.00	1m1s	783.22	2.3900	125992
			MATH500	62.30	10m29s	351.14	22.2400	623861
			AIME24	37.78	4m35s	736.40	73.4100	590741
			GPQA	25.33	15m21s	706.20	72.5800	1921405
	64	4096	GSM8k	80.67	28s	1700.96	3.7200	128915
			MATH500	60.67	2m16s	1303.67	25.1500	506862
			AIME24	18.89	1m22s	1413.73	52.9100	345118
			GPQA	14.67	3m22s	1929.84	62.5700	1131700
		8192	GSM8k	84.67	25s	1872.10	3.5800	125931
			MATH500	61.33	4m1s	897.99	30.3300	624385
			AIME24	35.56	2m46s	1222.58	82.1100	600925
			GPQA	28.00	6m57s	1567.52	102.1000	1936615

Table 3: Serving Results of LLM-7B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
LLM-7B	8	4096	GSM8k	69.67	1m47s	394.13	0.0600	107378
			MATH500	3.30	3m19s	343.65	0.0713	178459
			AIME24	15.56	1m34s	392.93	0.0776	101905
			GPQA	3.00	3m45s	324.91	0.1366	183061
	16	4096	GSM8k	70.00	1m33s	477.32	0.0991	115613
			MATH500	1.67	2m13s	513.34	0.1258	178452
			AIME24	18.89	50s	699.36	0.1208	95181
			GPQA	0.04	2m42s	495.68	0.2114	204317
	32	4096	GSM8k	67.67	57s	762.61	0.1698	111292
			MATH500	1.67	1m31s	748.09	0.2006	176704
			AIME24	16.67	33s	1063.23	0.1971	94296
			GPQA	6.00	1m34s	754.67	0.3807	175708

Table 4: Serving Results of RLLM-14B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
RLLM-14B	8	4096	GSM8k	87.00	3m50s	280.83	5.9731	177347
			MATH500	55.00	11m49s	277.63	20.0256	566267
			AIME24	27.78	4m42s	411.29	50.6600	341998
			GPQA	16.33	15m25s	406.14	39.0104	1090910
		8192	GSM8k	87.67	3m30s	303.34	6.0476	175210
			MATH500	62.33	17m17s	218.77	25.5350	656437
			AIME24	47.78	9m18s	338.41	72.4221	558966
			GPQA	21.00	30m34s	341.19	67.0687	1842315
	16	4096	GSM8k	86.33	2m40s	402.48	7.6400	173204
			MATH500	60.33	8m27s	388.92	27.2100	563198
			AIME24	27.78	3m11s	594.52	61.7600	335082
			GPQA	15.33	10m48s	581.23	52.2400	1098423
		8192	GSM8k	86.33	3m01s	382.72	8.1400	180684
			MATH500	63.67	13m25s	291.14	33.2600	669369
			AIME24	47.78	6m28s	498.79	109.3600	568311
			GPQA	22.00	21m16s	481.30	95.8800	1820128
	32	4096	GSM8k	85.33	1m38s	619.91	10.1400	169539
			MATH500	57.33	5m38s	568.13	36.8600	554485
			AIME24	22.22	2m20s	839.99	92.3600	344929
			GPQA	15.00	8m18s	770.56	78.5500	1104010
		8192	GSM8k	86.33	2m13s	486.36	11.0500	182233
			MATH500	66.00	9m09s	414.53	42.8600	654295
			AIME24	50.00	4m22s	722.93	140.4300	563372
			GPQA	26.67	15m36s	658.74	136.5100	1818245

Table 5: Serving Results of LLM-14B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
LLM-14B	8	4096	GSM8k	74.17	1m18s	317.73	0.0603	60286
			MATH500	44.00	2m29s	286.86	0.0626	100818
			AIME24	2.22	1m47s	240.82	0.0626	69874
			GPQA	29.00	2m15s	305.50	0.1222	87836
		8192	GSM8k	74.17	1m16s	337.06	0.0164	59314
			MATH500	44.67	2m23s	282.50	0.0635	99443
			AIME24	3.33	1m54s	252.58	0.0619	77607
			GPQA	29.33	2m09s	228.90	0.1217	90576
	16	4096	GSM8k	77.33	44s	529.30	0.1536	55856
			MATH500	46.00	1m57s	365.86	0.0913	106980
			AIME24	3.33	1m05s	388.44	0.0913	66109
			GPQA	25.33	1m40s	424.31	0.2250	93954
		8192	GSM8k	77.33	45s	554.18	0.0850	57030
			MATH500	47.00	1m53s	385.42	0.0920	109010
			AIME24	4.44	58s	409.73	0.9020	62213
			GPQA	24.67	2m05s	381.19	0.2547	97573
	32	4096	GSM8k	74.09	44s	615.89	0.1350	60318
			MATH500	47.67	1m04s	597.29	0.1393	94561
			AIME24	5.56	41s	616.16	0.1354	68309
			GPQA	28.00	1m17s	403.77	0.4789	95110
		8192	GSM8k	74.59	52s	526.30	0.1320	61833
			MATH500	46.67	1m20s	523.91	0.1385	102168
			AIME24	2.22	41s	620.70	0.1445	69811
			GPQA	28.67	1m19s	541.19	0.4215	93025
	64	4096	GSM8k	84.00	2m15s	463.17	7.3815	169635
			MATH500	59.33	7m57s	406.74	25.6479	560880
			AIME24	23.33	3m3s	630.49	58.5514	340172
			GPQA	17.00	10m22s	607.50	50.8592	1098715
		8192	GSM8k	88.00	2m39s	403.58	7.2181	174708
			MATH500	68.00	12m53s	299.83	32.5928	667886
			AIME24	53.33	6m7s	532.62	112.0436	577277
			GPQA	25.33	20m43s	504.45	86.7966	1846931

Table 6: Serving Results of RLLM-32B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
RLLM-32B	8	4096	GSM8k	91.00	4m16s	192.12	5.2715	130170
			MATH500	64.00	24m58s	125.59	42.9319	537799
			AIME24	21.11	9m12s	215.58	104.0663	349099
			GPQA	20.00	5m0s	206.07	73.0046	1071843
		8192	GSM8k	90.33	4m11s	194.77	5.2037	129401
			MATH500	70.67	36m41s	97.83	51.2373	621128
			AIME24	45.56	18m36s	174.28	150.4045	575207
			GPQA	24.33	60m31s	166.67	129.2077	1779384
	16	4096	GSM8k	90.67	2m27s	324.74	5.4900	128546
			MATH500	66.33	14m49	201.98	44.6300	517659
			AIME24	25.56	5m03s	388.52	111.2300	346308
			GPQA	21.67	17m25s	352.72	74.0400	1070143
		8192	GSM8k	92.67	2m30s	324.04	5.5700	128060
			MATH500	68.33	25m38s	134.29	53.4000	597129
			AIME24	48.89	10m22s	309.03	170.1700	568936
			GPQA	28.67	35m30s	283.98	137.9600	1778995
	32	4096	GSM8k	91.33	1m39s	490.92	6.7103	129986
			MATH500	66.67	9m29s	309.67	47.9789	504391
			AIME24	28.89	3m01s	631.69	119.5658	335582
			GPQA	18.00	11m25s	541.78	94.9504	1077209
		8192	GSM8k	92.33	1m39s	494.92	6.4570	129510
			MATH500	68.67	16m6s	213.37	60.4897	593685
			AIME24	50.00	6m06s	502.69	186.3563	547959
			GPQA	23.00	23m13s	436.05	171.4893	1787777

Table 7: Serving Results of LLM-32B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
LLM-32B	8	4096	GSM8k	60.67	7m25s	87.34	0.1271	100271
			MATH500	46.67	8m24s	103.76	0.1339	131828
			AIME24	6.67	2m43s	131.04	0.1414	57729
			GPQA	29.00	7m6s	141.18	0.2291	144033
		8192	GSM8k	60.67	7m10s	89.53	0.0867	98287
			MATH500	44.00	7m42s	109.18	0.0874	127011
			AIME24	4.44	3m24s	127.82	0.0899	71654
			GPQA	29.00	6m18s	152.47	0.1141	136446
	16	4096	GSM8k	66.33	4m03s	130.10	0.1083	83076
			MATH500	49.00	4m32s	170.74	0.1104	115631
			AIME24	12.22	2m13s	185.90	0.1092	67231
			GPQA	31.00	4m23s	210.62	0.1544	139381
		8192	GSM8k	62.67	4m04s	126.56	0.1107	92586
			MATH500	46.03	5m01s	156.52	0.1115	122183
			AIME24	6.67	1m53s	200.03	0.1179	60391
			GPQA	23.33	3m46s	259.97	0.1659	136660
	32	4096	GSM8k	63.00	5m24s	125.09	0.1902	96723
			MATH500	45.00	5m44	148.90	0.2097	123313
			AIME24	7.78	1m39s	211.36	0.2206	57192
			GPQA	24.33	4m06s	231.08	0.3894	141199
		8192	GSM8k	62.15	3m59s	129.75	0.1094	87157
			MATH500	44.37	5m31s	148.39	0.1109	128664
			AIME24	7.78	1m52s	200.63	0.1113	63495
			GPQA	25.33	4m31s	246.55	0.1872	136403

Table 8: Serving Results of RLLM-70B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
RLLM-70B	8	4096	GSM8k	90.00	5m22s	146.65	6.6450	125391
			MATH500	54.67	30m29s	102.48	48.4212	538966
			AIME24	32.22	11m46s	162.46	108.0681	336881
			GPQA	22.67	38m11s	158.29	93.9947	1052080
		8192	GSM8k	90.00	5m30s	143.90	6.5438	125874
			MATH500	62.00	49m48s	78.31	61.8861	677161
			AIME24	54.44	23m18s	135.50	192.8731	561002
			GPQA	29.67	75m38s	128.03	159.9508	1702698
	16	4096	GSM8k	88.33	3m24s	230.39	7.4600	123185
			MATH500	57.00	19m45s	158.60	52.3600	539666
			AIME24	26.67	7m01s	277.08	140.8100	342355
			GPQA	23.00	23m35s	253.66	98.0500	1042191
		8192	GSM8k	88.00	3m32s	228.88	7.7200	125291
			MATH500	60.67	32m27s	112.85	69.6400	644132
			AIME24	55.56	13m46s	221.77	197.0300	539132
			GPQA	30.67	46m01s	204.78	181.9000	1659750
	32	4096	GSM8k	88.67	2m11s	352.22	9.7962	122786
			MATH500	56.67	13m0s	238.82	66.2581	534918
			AIME24	25.56	4m26s	438.58	184.7638	343514
			GPQA	23.00	15m57s	378.78	136.1667	1052451
		8192	GSM8k	89.00	2m12s	352.59	9.4598	123720
			MATH500	62.33	21m36s	166.26	83.2369	621237
			AIME24	51.11	8m24s	360.36	246.0471	537908
			GPQA	32.00	30m31s	307.04	228.6145	1650647

Table 9: Serving Results of LLM-70B

Model	BS	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
LLM-70B	8	4096	GSM8k	93.00	3m03s	144.56	0.2030	62746
			MATH500	59.33	10m38s	90.67	0.2283	148624
			AIME24	30.00	5m35s	107.39	0.2461	99826
			GPQA	53.33	12m18s	128.53	0.4749	243123
		8192	GSM8k	92.67	2m58s	144.42	0.1143	60838
			MATH500	59.00	13m12s	78.52	0.1197	162129
			AIME24	23.33	6m1s	99.04	0.1232	98349
			GPQA	51.67	11m24s	131.86	0.1847	233596
	16	4096	GSM8k	91.33	1m44s	243.39	0.1525	60890
			MATH500	60.00	6m35s	139.78	0.1730	144651
			AIME24	27.78	3m06s	171.31	0.1678	89789
			GPQA	47.67	7m28s	201.44	0.2694	231586
		8192	GSM8k	93.67	1m52s	227.03	0.1536	61434
			MATH500	59.00	6m38s	139.61	0.1704	142362
			AIME24	27.78	5m13s	124.39	0.1715	103021
			GPQA	49.00	10m26s	155.10	0.2551	247291
	32	4096	GSM8k	92.33	1m17s	346.24	0.6077	61987
			MATH500	60.33	5m29s	176.59	0.6520	149612
			AIME24	26.56	2m37s	243.93	0.7012	106146
			GPQA	51.00	5m40s	268.96	1.5823	237672
		8192	GSM8k	93.00	1m14s	357.73	0.2325	60831
			MATH500	61.00	7m7s	148.03	0.2520	164985
			AIME24	27.78	3m23s	186.51	0.2550	105595
			GPQA	53.33	4m54s	306.15	0.5346	233070
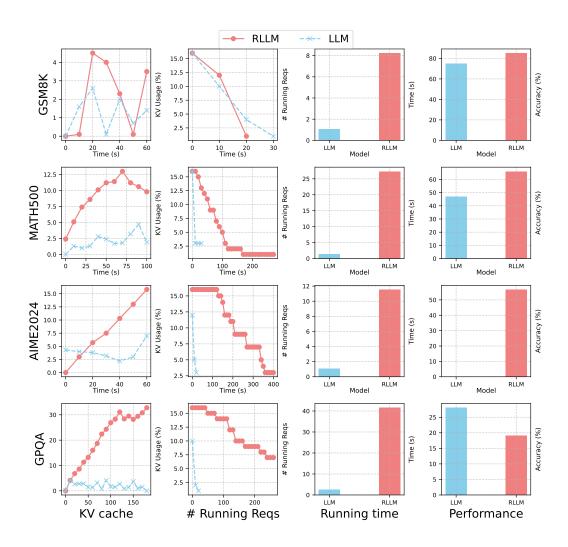
Figure 9: Results of RLLM vs LLM for 14B model size .
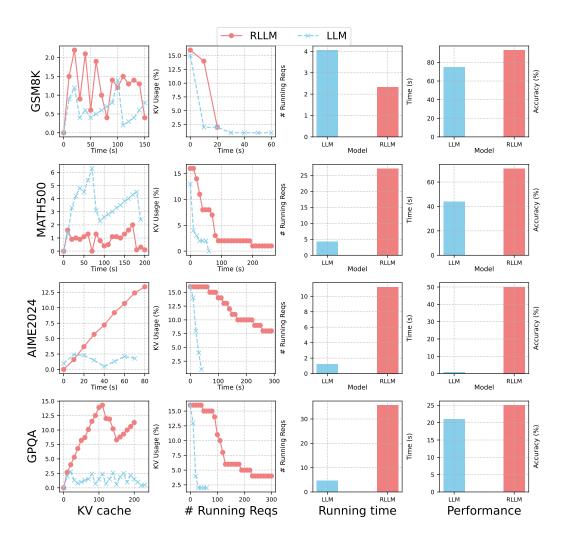
Figure 10: Results of RLLM vs LLM for 32B model size .

I.4 Running Traces Demo

INFO: 127.0.0.1:53458 – "POST /v1/chat/completions HTTP/1.1"
    200 OK

INFO 05–10 13:01:49 [loggers.py:80] Avg prompt throughput: 223.0
    tokens/s, Avg generation throughput: 164.5 tokens/s, Running:
    16 reqs, Waiting: 0 reqs, GPU KV cache usage: 1.0\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:01:59 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 650.0 tokens/s, Running:
    14 reqs, Waiting: 0 reqs, GPU KV cache usage: 2.8\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:02:09 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 495.5 tokens/s, Running:
    11 reqs, Waiting: 0 reqs, GPU KV cache usage: 3.9\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:02:19 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 406.1 tokens/s, Running:
    8 reqs, Waiting: 0 reqs, GPU KV cache usage: 4.0\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:02:29 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 320.8 tokens/s, Running:
    8 reqs, Waiting: 0 reqs, GPU KV cache usage: 5.1\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:02:39 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 320.7 tokens/s, Running:
    6 reqs, Waiting: 0 reqs, GPU KV cache usage: 4.9\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:02:49 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 271.3 tokens/s, Running:
    5 reqs, Waiting: 0 reqs, GPU KV cache usage: 4.9\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:02:59 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 180.7 tokens/s, Running:
    4 reqs, Waiting: 0 reqs, GPU KV cache usage: 4.5\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:03:09 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 164.8 tokens/s, Running:
    4 reqs, Waiting: 0 reqs, GPU KV cache usage: 5.0\%, Prefix
    cache hit rate: 11.6\%

INFO 05–10 13:03:19 [loggers.py:80] Avg prompt throughput: 0.0
    tokens/s, Avg generation throughput: 174.4 tokens/s, Running:
    4 reqs, Waiting: 0 reqs, GPU KV cache usage: 5.7\%, Prefix
    cache hit rate: 11.6\%

J Detailed Empirical Results for RLLM Serving Optimization

J.1 Model Weight Quantization

Full results of model weight quantization with different models are listed in Table 10 and 11.

Table 10: Results of RLLM-7B with Different Quantization Methods

Model	Method	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
				Budget-4096			
RLLM-7B	GPTQ	GSM8k	81.67	36s	1258.52	1.5015	125477
		MATH500	61.00	3m34s	807.99	10.8316	488868
		AIME24	16.67	1m25s	1383.83	28.3623	346988
		GPQA	16.00	4m48s	1342.53	23.6366	1127392
	AWQ	GSM8k	82.67	2m36s	306.10	5.7903	128626
		MATH500	59.33	14m55s	195.01	41.2395	501037
		AIME24	21.11	5m02s	390.68	104.9071	347634
		GPQA	14.33	17m22s	371.37	89.8281	1124314
	FP8	GSM8k	82.67	58s	805.44	1.8947	128016
		MATH500	64.00	4m44s	589.49	13.9700	479989
		AIME24	25.56	1m49s	1062.28	38.5411	341103
		GPQA	15.00	6m13s	1029.07	29.9774	1110873
	L4	GSM8k	82.67	1m21s	560.40	3.5042	119344
		MATH500	61.00	7m58s	356.31	24.6086	486082
		AIME24	20.00	3m05s	630.92	60.3252	343397
		GPQA	15.33	10m36s	599.34	51.2468	1109222
				Budget-8192			
RLLM-7B	GPTQ	GSM8k	80.33	57s	955.97	1.5007	135978
		MATH500	64.00	5m54s	591.44	12.3735	600650
		AIME24	31.11	2m59s	1156.19	44.9320	618796
		GPQA	25.33	9m47s	1131.34	44.2137	1959609
	AWQ	GSM8k	80.00	2m35s	306.42	5.6545	126262
		MATH500	66.67	25m03s	141.57	54.4308	618240
		AIME24	32.22	10m17s	333.68	162.2038	610667
		GPQA	24.33	34m54s	314.88	156.5370	1942820
	FP8	GSM8k	83.00	1m03s	774.97	2.0614	129654
		MATH500	63.00	7m38s	455.87	17.6713	24375
		AIME24	40.00	3m42s	880.38	61.7136	582089
		GPQA	27.67	12m40s	856.35	57.6786	1915156
	L4	GSM8k	80.33	1m24s	562.97	3.6269	123986
		MATH500	63.67	13m01s	264.82	29.4854	597768
		AIME24	36.67	6m34s	499.62	98.4701	586687
		GPQA	31.33	23m07s	461.52	94.4508	1879225

J.2 KV Cache Quantization

Full results of KV Cache quantization with different models are listed in Table 12, 13.

J.3 Prefix Caching

Full results of KV Cache quantization evaluation with different models are listed in Table 14, 15, 16, 17.

J.4 Speculative Decoding

The visualizatio for 7B model SD is in Figure 11. Full results of speculative decoding evaluation with different models are listed in this subsection. For RLLM, results are presneted in Table 18, 19, 20, 21.

Table 11: Results of RLLM-14B with Different Quantization Methods

| Model	Method	Dataset	Accuracy	Running Time	TPS	TTFT	Output Tokens
			Budget-4096				
RLLM-14B	GPTQ	GSM8k	87.67	1m55s	531.24	5.2007	167860
		MATH500	61.00	5m56s	535.47	18.478	547482
		AIME24	22.22	2m25s	811.91	49.699	346219
		GPQA	16.33	8m09s	775.78	40.463	1102007
	AWQ	GSM8k	86.67	1m30s	617.47	4.6092	151523
		MATH500	61.00	6m01s	534.74	19.725	551164
		AIME24	21.11	2m30s	771.16	49.545	342303
		GPQA	18.33	8m26s	750.95	42.983	1106144
	FP8	GSM8k	89.00	2m19s	446.80	6.5690	170719
		MATH500	60.67	7m16s	447.51	24.214	560177
		AIME24	24.44	2m55s	665.67	62.497	343575
		GPQA	14.33	9m49s	650.24	51.583	1113927
	L4	GSM8k	83.67	3m27s	323.02	9.9777	187673
		MATH500	58.67	9m24s	326.66	30.626	528533
		AIME24	26.67	3m41s	525.41	74.833	341339
		GPQA	16.33	12m44s	501.28	65.669	1113251
			Budget-8192				
RLLM-14B	GPTQ	GSM8k	84.67	1m48s	569.94	5.2287	168441
		MATH500	65.33	8m49s	418.08	22.0500	638458
		AIME24	40.00	5m03s	666.62	76.884	596659
		GPQA	25.00	16m14s	638.98	74.9600	1831969
	AWQ	GSM8k	86.67	1m48s	585.55	4.7892	155465
		MATH500	65.33	9m22s	399.52	23.658	647840
		AIME24	47.78	5m04s	640.79	84.494	575835
		GPQA	26.33	17m02s	621.02	77.938	1866185
	FP8	GSM8k	86.00	2m14s	467.49	6.5689	171349
		MATH500	63.33	11m40s	328.62	28.621	667889
		AIME24	51.11	5m43s	535.96	90.392	544733
		GPQA	27.33	19m27s	529.1	84.797	1817404
	L4	GSM8k	83.33	2m47s	376.01	9.3994	172748
		MATH500	63.00	14m44s	248.57	36.372	635662
		AIME24	47.78	7m58s	401.73	123.95	572091
		GPQA	21.33	28m55s	368.01	115.88	1879105
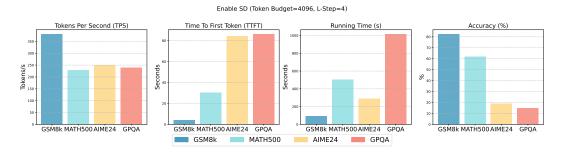
Figure 11: Results of 7B RLLM with SD enabled .

Table 12: Results of RLLM-7B with Different KV Cache Quantization Methods

Model	Method	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
			Budget-4096				
RLLM-7B	FP8-E4M3	GSM8k	9.33	7m5s	678.93	3.8153	843936
		MATH500	4.33	7m26s	849.23	17.1335	1113415
		AIME24	0.00	2m13s	936.07	23.4623	732374
		GPQA	0.33	16m27s	831.28	73.5032	2428386
	FP8-E5M2	GSM8k	2.67	9m8s	624.13	6.5660	1013649
		MATH500	0.33	9m28s	731.99	22.0233	1223241
		AIME24	0.00	2m50s	733.58	71.2345	367838
		GPQA	0.00	9m36s	739.58	36.1577	1225641
			Budget-8192				
RLLM-7B	FP8-E4M3	GSM8k	8.67	14m23s	634.50	5.0929	1631364
		MATH500	4.00	16m01s	796.17	36.4224	2271199
		AIME24	0.00	4m18s	855.67	50.5708	732374
		GPQA	0.67	7m39s	914.28	33.3658	1225852
	FP8-E5M2	GSM8k	2.33	19m04s	563.56	26.7276	1919839
		MATH500	0.33	20m02s	681.93	94.8742	2438321
		AIME24	0.00	5m59s	682.40	91.3258	730016
		GPQA	0.00	20m21s	678.99	102.0499	2450781

Table 13: Results of RLLM-14B with Different KV Cache Quantization Methods

Model	Method	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
			Budget-4096				
RLLM-14B	FP8-E4M3	GSM8k	87.33	2m17s	466.06	7.5905	169107
		MATH500	63.00	7m37s	417.60	25.1716	547992
		AIME24	26.67	2m59s	640.63	63.6228	338613
		GPQA	15.00	10m03s	631.30	50.1774	1106576
	FP8-E5M2	GSM8k	82.33	2m40s	390.53	7.6730	171445
		MATH500	59.33	7m46s	394.45	25.4353	527039
		AIME24	26.67	3m02s	632.26	62.9118	339940
		GPQA	15.67	10m19s	608.30	49.3198	1094547
			Budget-8192				
RLLM-14B	FP8-E4M3	GSM8k	83.67	3m03s	357.72	7.7833	180474
		MATH500	66.33	12m18s	305.61	30.8091	653017
		AIME24	52.22	5m53s	520.69	94.5251	545376
		GPQA	24.00	20m23s	511.37	89.8416	1838882
	FP8-E5M2	GSM8k	85.33	2m39s	401.41	7.8384	175761
		MATH500	62.67	12m48s	289.84	29.0644	633207
		AIME24	48.89	6m04s	513.09	92.6786	553741
		GPQA	26.67	20m51s	497.93	90.6068	1833865

Table 14: Results of RLLM-7B without Prefix Cache

Model	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
			Budget-4096			
RLLM-7B	GSM8k	79.00	1m26s	564.90	3.5931	130372
	MATH500	60.67	6m54s	420.14	22.4207	498018
	AIME24	18.89	2m46s	708.96	58.7487	349035
	GPQA	14.67	9m24s	686.08	47.7385	1124613
			Budget-8192			
RLLM-7B	GSM8k	81.33	1m14s	622.53	3.4274	124776
	MATH500	60.33	11m02s	335.91	26.4863	639141
	AIME24	38.89	5m26s	599.82	96.2117	588055
	GPQA	27.33	18m42s	583.69	89.0785	1932823

Table 15: Results of RLLM-14B without Prefix Cache

Model	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
			Budget-4096			
RLLM-14B	GSM8k	88.00	2m31s	420.23	8.2604	178175
	MATH500	57.67	8m17s	404.26	28.1407	576686
	AIME24	23.33	3m12s	604.58	69.2730	342347
	GPQA	13.67	10m52s	587.18	56.3880	1114920
			Budget-8192			
RLLM-14B	GSM8k	87.67	2m34s	412.78	8.0207	171501
	MATH500	62.33	12m31s	292.60	31.9481	655351
	AIME24	48.89	6m22s	502.43	107.6285	569489
	GPQA	23.33	21m20s	481.04	91.3230	1817908

Table 16: Results of RLLM-32B without Prefix Cache

Model	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
			Budget-4096			
RLLM-32B	GSM8k	92.67	2m30s	327.11	5.8756	129761
	MATH500	61.00	15m30s	196.65	44.8259	529743
	AIME24	25.56	5m09s	379.99	113.5958	344566
	GPQA	19.00	17m45s	349.79	83.0689	1082837
			Budget-8192			
RLLM-32B	GSM8k	92.33	2m32s	314.79	5.8947	130188
	MATH500	70.00	23m14s	152.31	58.6775	615489
	AIME24	56.67	10m24s	302.55	174.7540	559504
	GPQA	27.67	36m01s	275.68	139.8732	1752688

Table 17: Results of RLLM-70B without Prefix Cache

Model	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
			Budget-4096			
RLLM-70B	GSM8k	90.00	3m25s	228.12	8.2671	123955
	MATH500	54.00	20m22s	153.16	53.0075	535884
	AIME24	31.11	7m05s	269.74	136.8779	336872
	GPQA	19.33	24m06s	251.73	110.7695	1055915
			Budget-8192			
RLLM-70B	GSM8k	90.00	3m39s	218.50	8.3275	126163
	MATH500	61.33	31m57s	115.36	73.7041	634981
	AIME24	54.44	13m52s	221.07	224.2071	544412
	GPQA	31.33	46m37s	200.15	177.5976	1644534

Table 18: Results of RLLM-7B with Different Speculative Decoding Methods

Model	Budget	Dataset	Acc.	Running Time	TPS	TTFT	Output Tokens
				L-Step: 2			
RLLM-7B	4096	GSM8k	83.33	1m21s	411.47	3.8268	84044
		MATH500	63.67	8m26s	241.67	28.9989	342214
		AIME24	18.89	4m52s	266.43	91.1752	226340
		GPQA	14.33	16m04s	261.39	73.0445	720304
	8192	GSM8k	82.67	1m21s	412.96	3.7715	83798
		MATH500	61.33	13m50s	176.53	34.7626	415561
		AIME24	37.78	11m13s	192.58	156.9481	383225
		GPQA	26.67	37m16s	187.68	161.3714	1223433
				L-Step: 4			
RLLM-7B	4096	GSM8k	82.33	1m32s	380.14	3.9015	83812
		MATH500	62.00	8m26s	228.30	30.2225	325135
		AIME24	18.89	4m49s	250.19	84.2456	213773
		GPQA	15.00	16m54s	240.27	86.0878	699813
	8192	GSM8k	83.67	1m17s	418.62	4.0124	83857
		MATH500	65.00	13m05s	178.65	36.4870	400672
		AIME24	37.78	11m13s	180.60	151.0329	356258
		GPQA	28.33	37m11s	172.07	152.6671	1119974
				L-Step: 8			
RLLM-7B	4096	GSM8k	85.67	1m28s	378.44	4.0057	84219
		MATH500	61.00	8m35s	228.01	30.6229	328151
		AIME24	21.11	4m57s	248.74	92.2684	215501
		GPQA	14.67	16m25s	244.13	78.9886	685904
	8192	GSM8k	82.33	1m26s	388.10	4.0512	83643
		MATH500	60.33	13m15s	174.36	36.0669	392872
		AIME24	38.89	10m36s	184.35	146.6946	344680
		GPQA	23.33	36m48s	180.27	171.9770	1158190

Table 19: Results of RLLM-14B with Speculative Decoding Method

Model	Budget	L-Step	Dataset	Accuracy	Running Time	TPS	TTFT	Output Tokens
RLLM-14B	4096	4	GSM8k	88.00	3m3s	238.27	11.1347	113595
			MATH500	59.33	11m52s	187.51	45.6259	373475
			AIME24	24.44	6m09s	199.02	126.3789	217784
			GPQA	15.67	20m42s	196.47	105.7676	696919
	8192	4	GSM8k	85.00	2m47s	255.10	11.5077	112893
			MATH500	62.67	17m22s	148.10	52.2801	433701
			AIME24	52.22	12m17s	151.99	194.4668	337745
			GPQA	23.33	43m14s	146.90	182.6785	1099344

Table 20: Results of RLLM-32B with Speculative Decoding Method

Model	Budget	L-Step	Dataset	Accuracy	Running Time	TPS	TTFT	Output Tokens
RLLM-32B	4096	4	GSM8k	90.33	2m36s	218.41	7.2955	84809
			MATH500	62.33	15m45s	125.10	57.7919	333363
			AIME24	25.56	8m12s	148.29	157.7257	211331
			GPQA	16.67	26m29s	147.23	117.7640	665233
	8192	4	GSM8k	92.00	2m32s	222.34	7.3433	84595
			MATH500	68.67	26m6s	84.75	77.4161	376796
			AIME24	47.78	18m56s	99.98	240.3189	331180
			GPQA	24.00	56m02s	105.96	216.8525	1050802

Table 21: Results of RLLM-70B with Speculative Decoding Method

Model	Budget	L-Step	Dataset	Accuracy	Running Time	TPS	TTFT	Output Tokens
RLLM-70B	4096	4	GSM8k	88.67	3m30s	157.68	10.0978	84118
			MATH500	56.67	22m01s	98.28	72.7944	367031
			AIME24	32.22	10m25s	121.87	195.0349	221532
			GPQA	23.33	33m40s	117.38	159.7265	678418
	8192	4	GSM8k	90.33	3m33s	158.49	10.2136	84802
			MATH500	59.00	31m32s	77.02	86.9171	413304
			AIME24	51.11	21m15s	94.05	298.7859	349228
			GPQA	34.33	66m23s	90.79	278.9286	1049010

For LLM, results are presented in Table 22 (7B), 23 (32B).

Table 22: Results of LLM-7B with Different Speculative Decoding Methods

Model	Budget	L-Step	Dataset	Accuracy	Running Time	TPS	TTFT	Output Tokens
LLM-7B	4096	2	GSM8k	68.67	2m07s	262.59	0.2300	83610
			MATH500	2.33	2m40s	299.47	0.3243	119825
			AIME24	15.56	1m10s	321.47	0.3449	60913
			GPQA	6.33	2m29s	282.74	0.5113	90423
		4	GSM8k	66.67	1m47s	304.82	0.2263	80208
			MATH500	0.67	2m32s	309.60	0.3160	115604
			AIME24	18.89	1m02s	342.77	0.3234	54837
			GPQA	3.67	2m24s	297.60	0.5061	36045
		8	GSM8k	69.33	1m59s	268.42	0.2257	79618
			MATH500	2.00	2m54s	277.62	0.3213	121082
			AIME24	21.11	1m06s	313.01	0.3505	54877
			GPQA	3.00	2m27s	275.03	0.5067	85956

Table 23: Results of LLM-32B with Speculative Decoding Method

Model	Budget	L-Step	Dataset	Accuracy	Running Time	TPS	TTFT	Output Tokens
LLM-32B	4096	4	GSM8k	59.67	4m07s	104.92	0.3605	61457
			MATH500	45.33	3m57s	142.52	0.5117	77136
			AIME24	6.67	1m19s	170.42	0.5746	32802
			GPQA	23.67	3m12s	199.91	0.9070	79295
	8192	4	GSM8k	63.00	3m33s	108.95	0.3581	53146
			MATH500	45.67	3m33s	146.57	0.5238	69419
			AIME24	3.33	1m14s	182.46	0.5354	31939
			GPQA	24.67	3m04s	204.56	0.8977	77214

K Extended Results for Real World Benchmarking

Figure 12: KV cache usage of 7B models under real-world workload across different datasets.
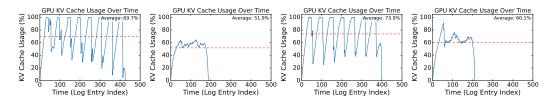

Figure 13: KV cache usage of 32B models under real-world workload across different datasets.
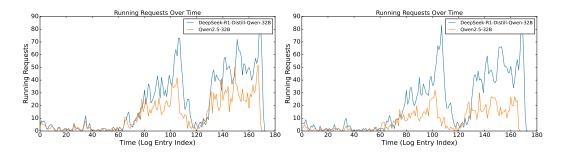
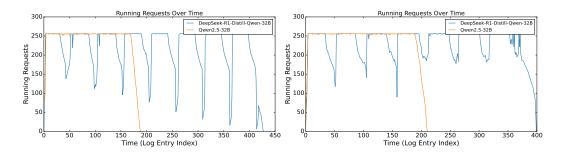Figure 14: Num of running requests in the inference engine for 7B models under real-world workload.

Figure 15: Num of running requests in the inference engine for 32B models under real-world workload.